

# ***Introduction to GPU Architecture***

Ofer Rosenberg,  
PMTS SW, OpenCL Dev. Team  
AMD

Based on  
“From Shader Code to a Teraflop: How GPU Shader Cores Work”,  
By Kayvon Fatahalian, Stanford University

# Content

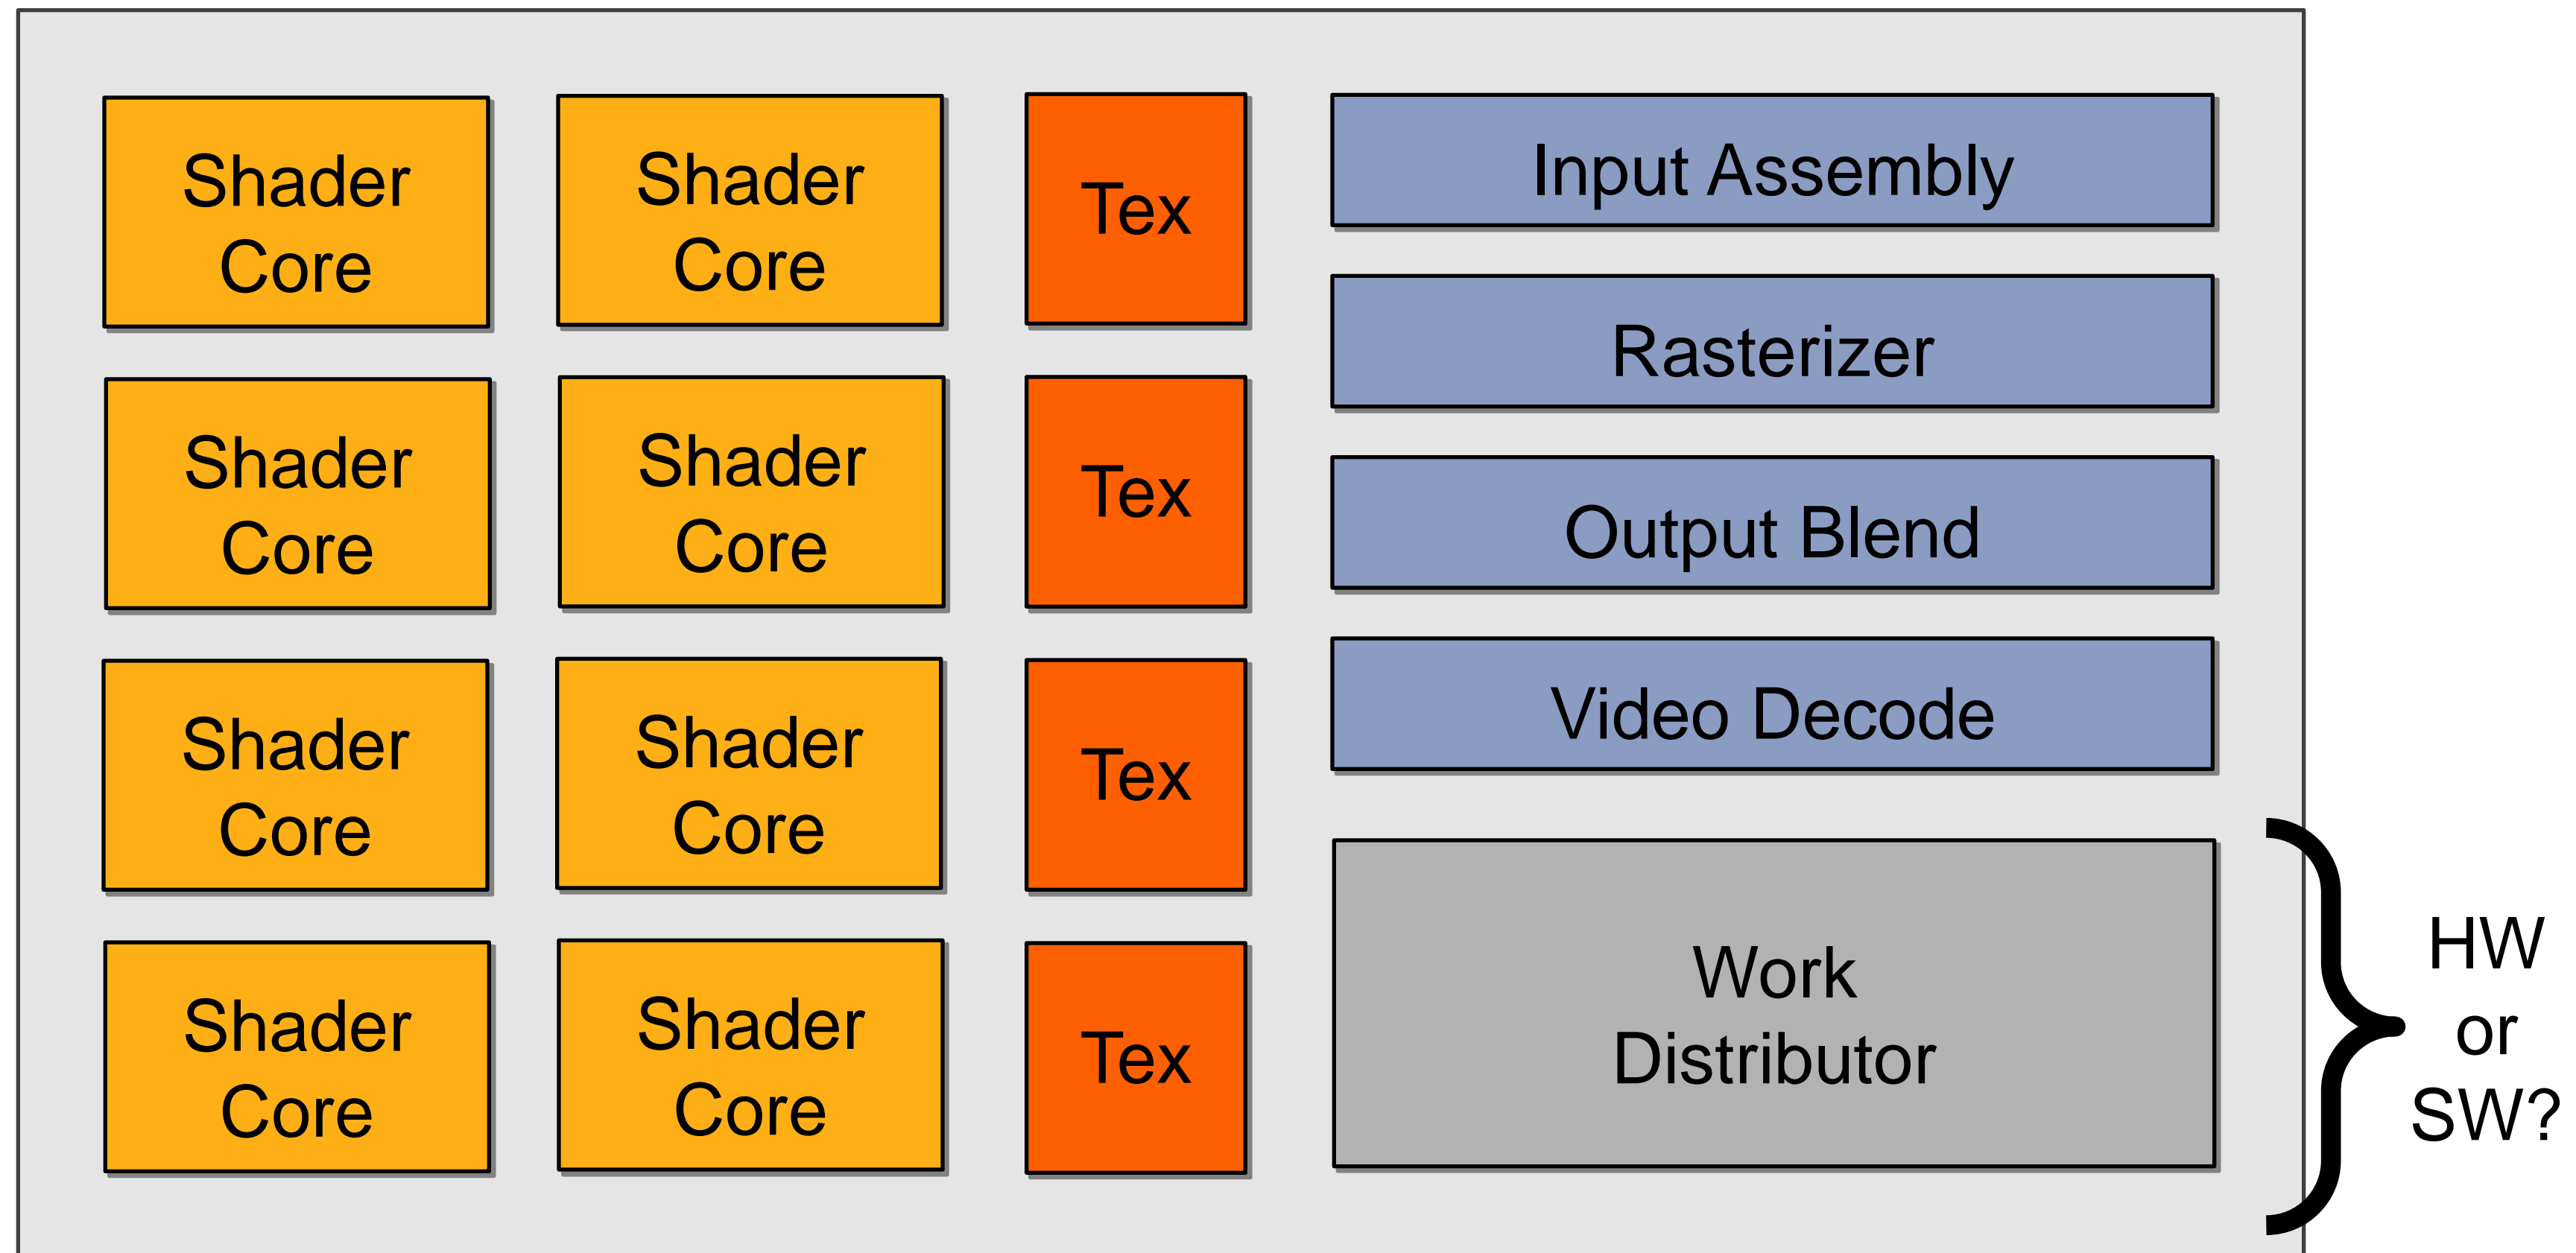
1. Three major ideas that make GPU processing cores run fast
2. Closer look at real GPU designs
  - NVIDIA GTX 580
  - AMD Radeon 6970
3. The GPU memory hierarchy: moving data to processors
4. Heterogeneous Cores

# Part 1: throughput processing

- Three key concepts behind how modern GPU processing cores run code
- Knowing these concepts will help you:
  1. Understand space of GPU core (and throughput CPU core) designs
  2. Optimize shaders/compute kernels
  3. Establish intuition: what workloads might benefit from the design of these architectures?

# What's in a GPU?

A GPU is a heterogeneous chip multi-processor (highly tuned for graphics)



# A diffuse reflectance shader

```
sampler mySamp;  
Texture2D<float3> myTex;  
float3 lightDir;  
  
float4 diffuseShader(float3 norm, float2 uv)  
{  
    float3 kd;  
    kd = myTex.Sample(mySamp, uv);  
    kd *= clamp( dot(lightDir, norm), 0.0, 1.0);  
    return float4(kd, 1.0);  
}
```

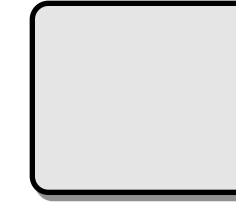
Shader programming model:

Fragments are processed  
**independently,**  
but there is no explicit parallel  
programming



# Compile shader

1 unshaded fragment input record



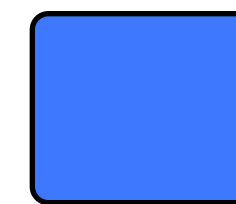
```
sampler mySamp;  
Texture2D<float3> myTex;  
float3 lightDir;  
  
float4 diffuseShader(float3 norm, float2 uv)  
{  
    float3 kd;  
    kd = myTex.Sample(mySamp, uv);  
    kd *= clamp( dot(lightDir, norm), 0.0, 1.0);  
    return float4(kd, 1.0);  
}
```



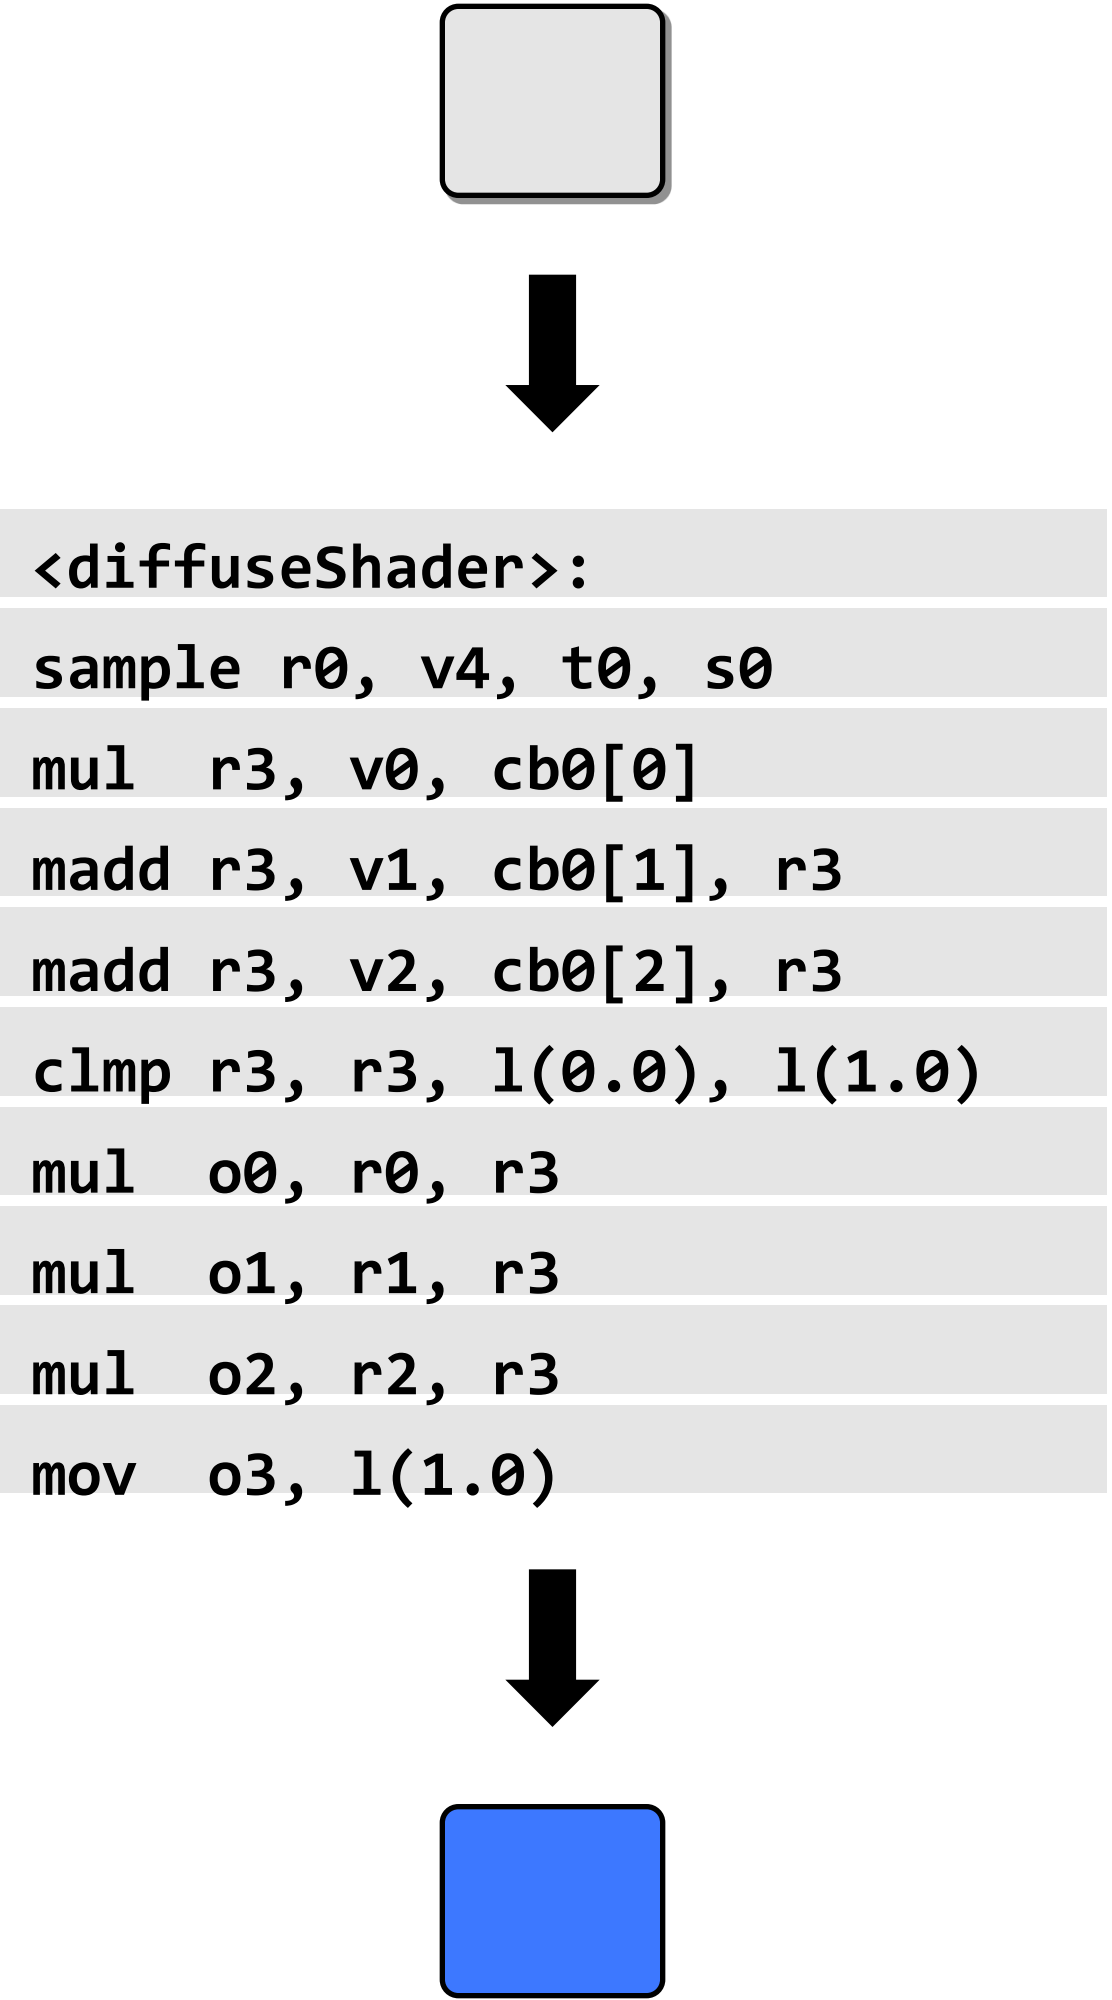
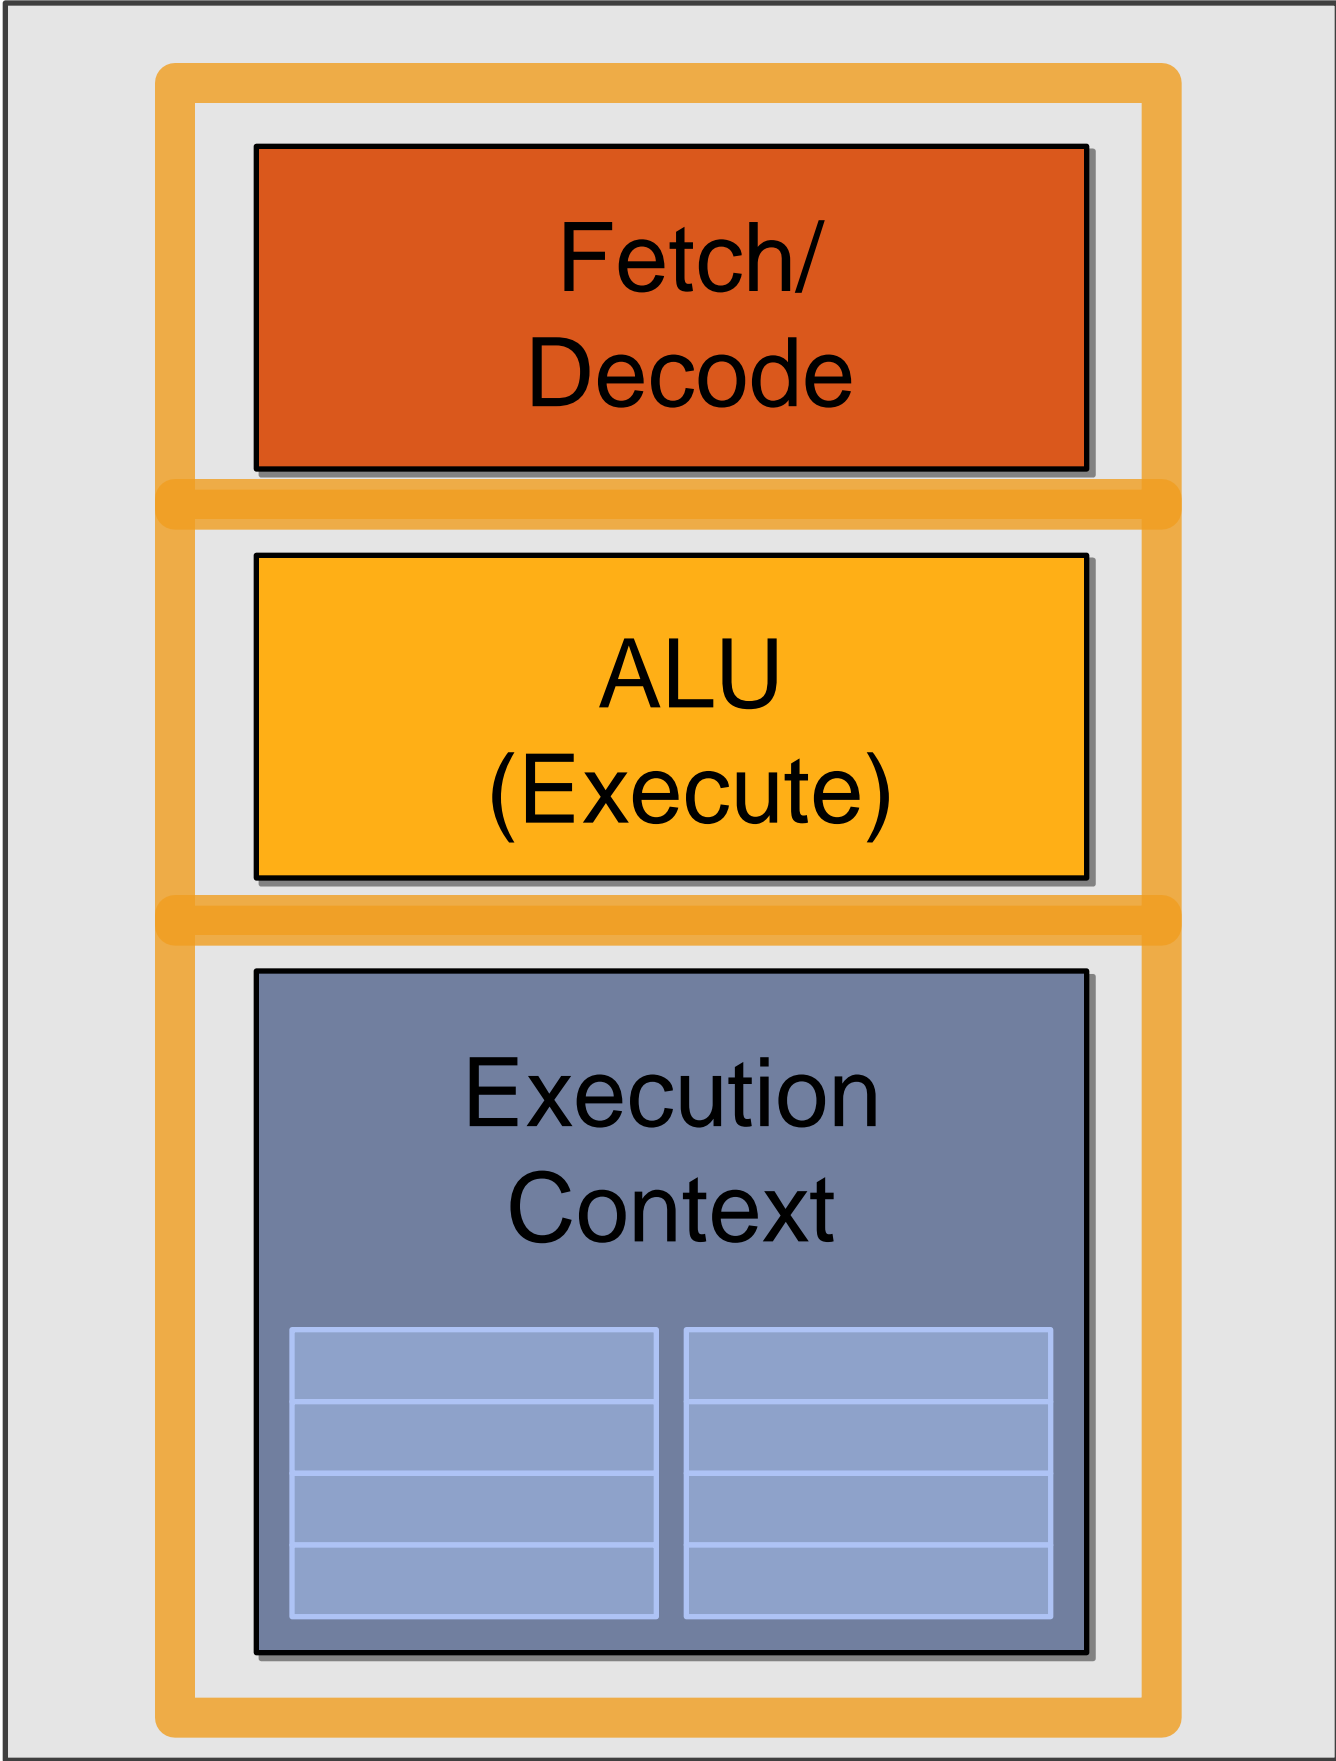
```
<diffuseShader>:  
sample r0, v4, t0, s0  
mul   r3, v0, cb0[0]  
madd  r3, v1, cb0[1], r3  
madd  r3, v2, cb0[2], r3  
clmp  r3, r3, 1(0.0), 1(1.0)  
mul   o0, r0, r3  
mul   o1, r1, r3  
mul   o2, r2, r3  
mov   o3, 1(1.0)
```



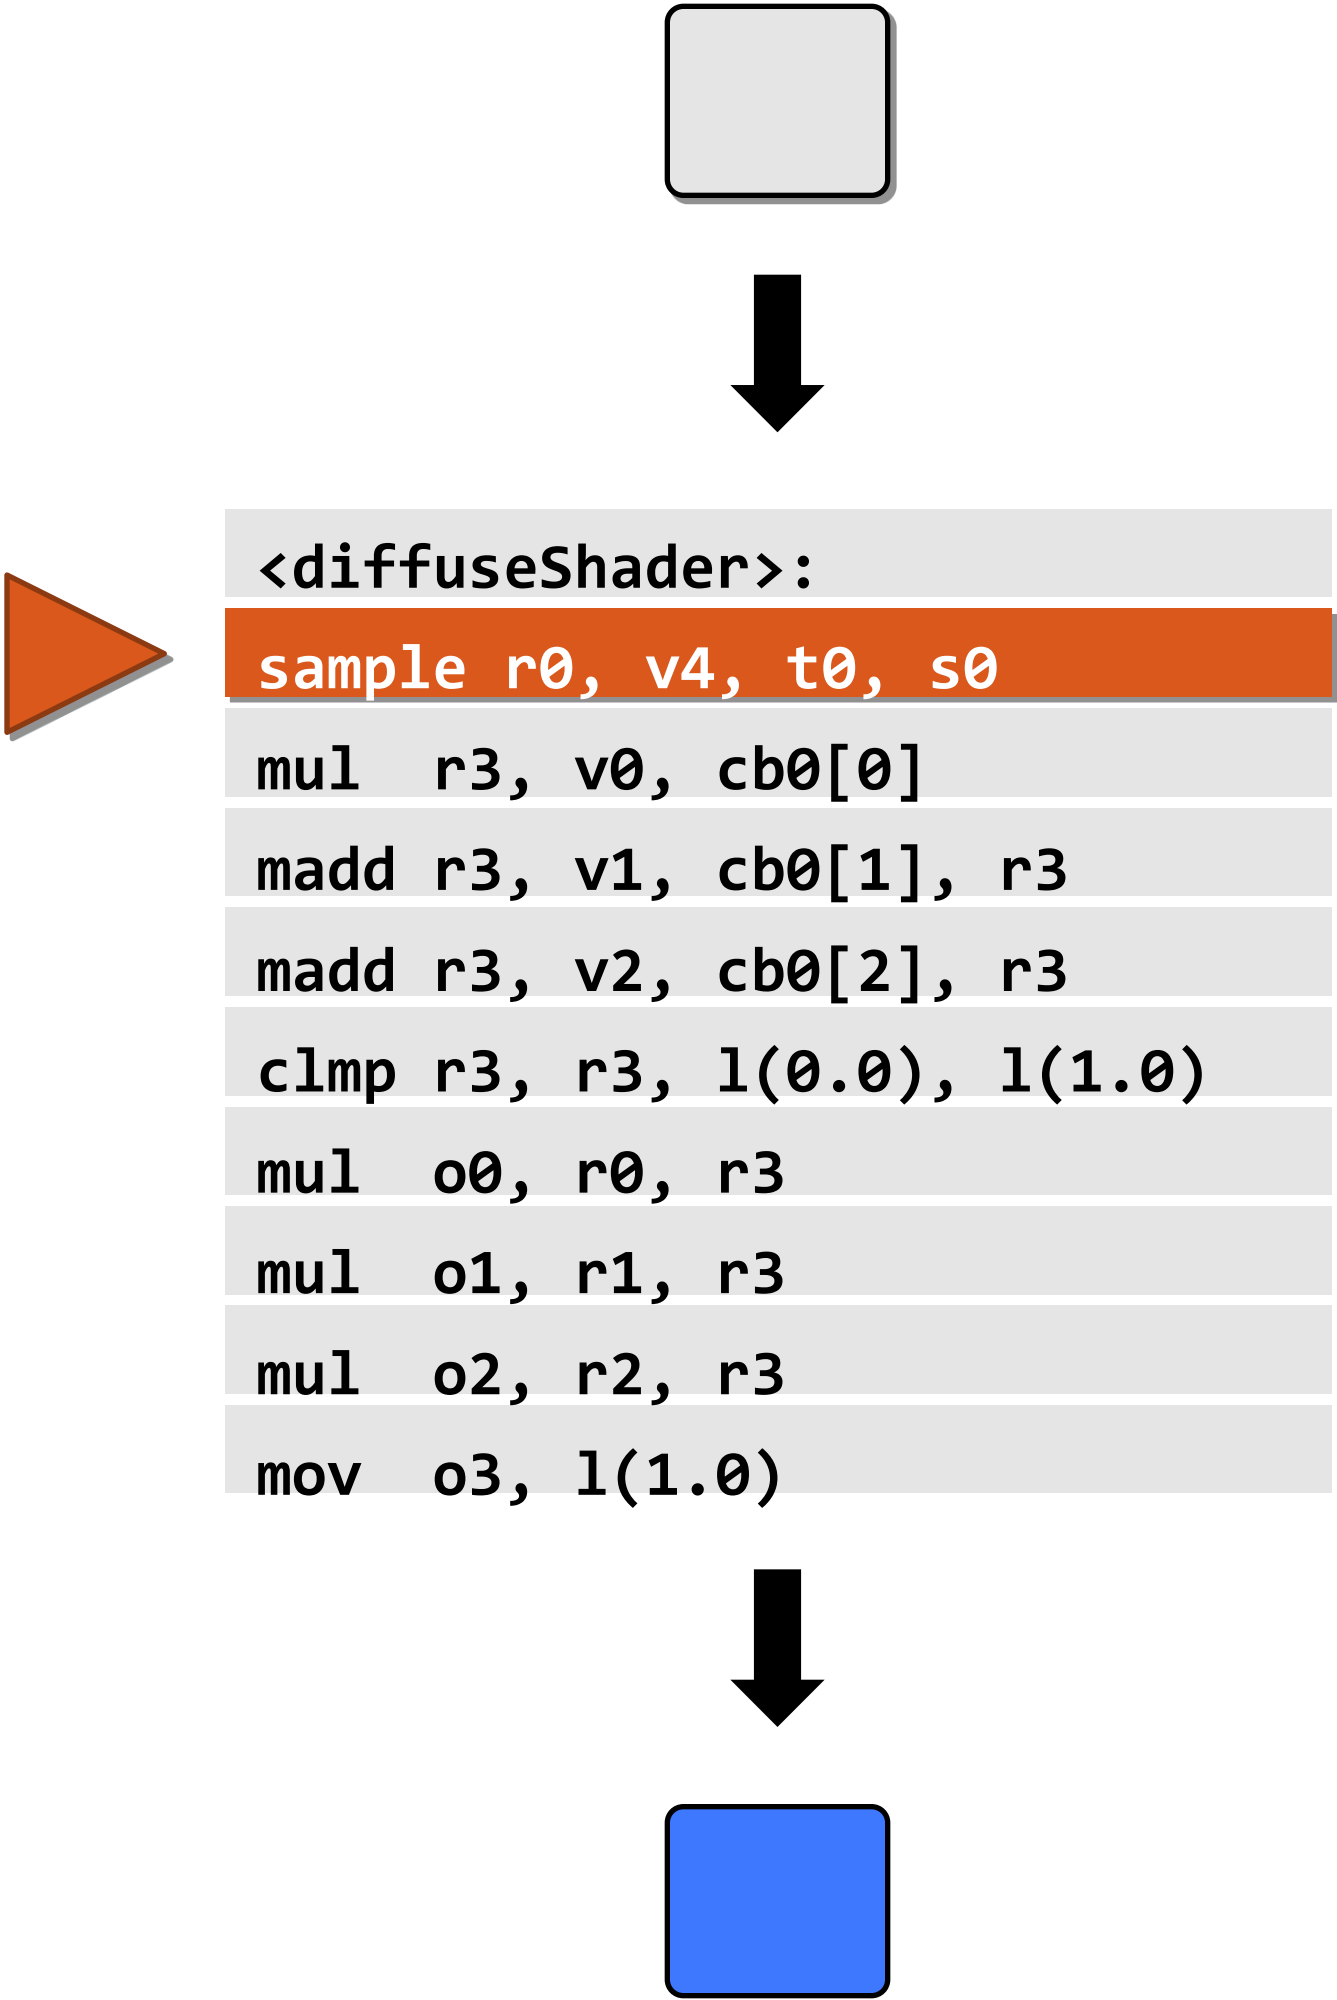
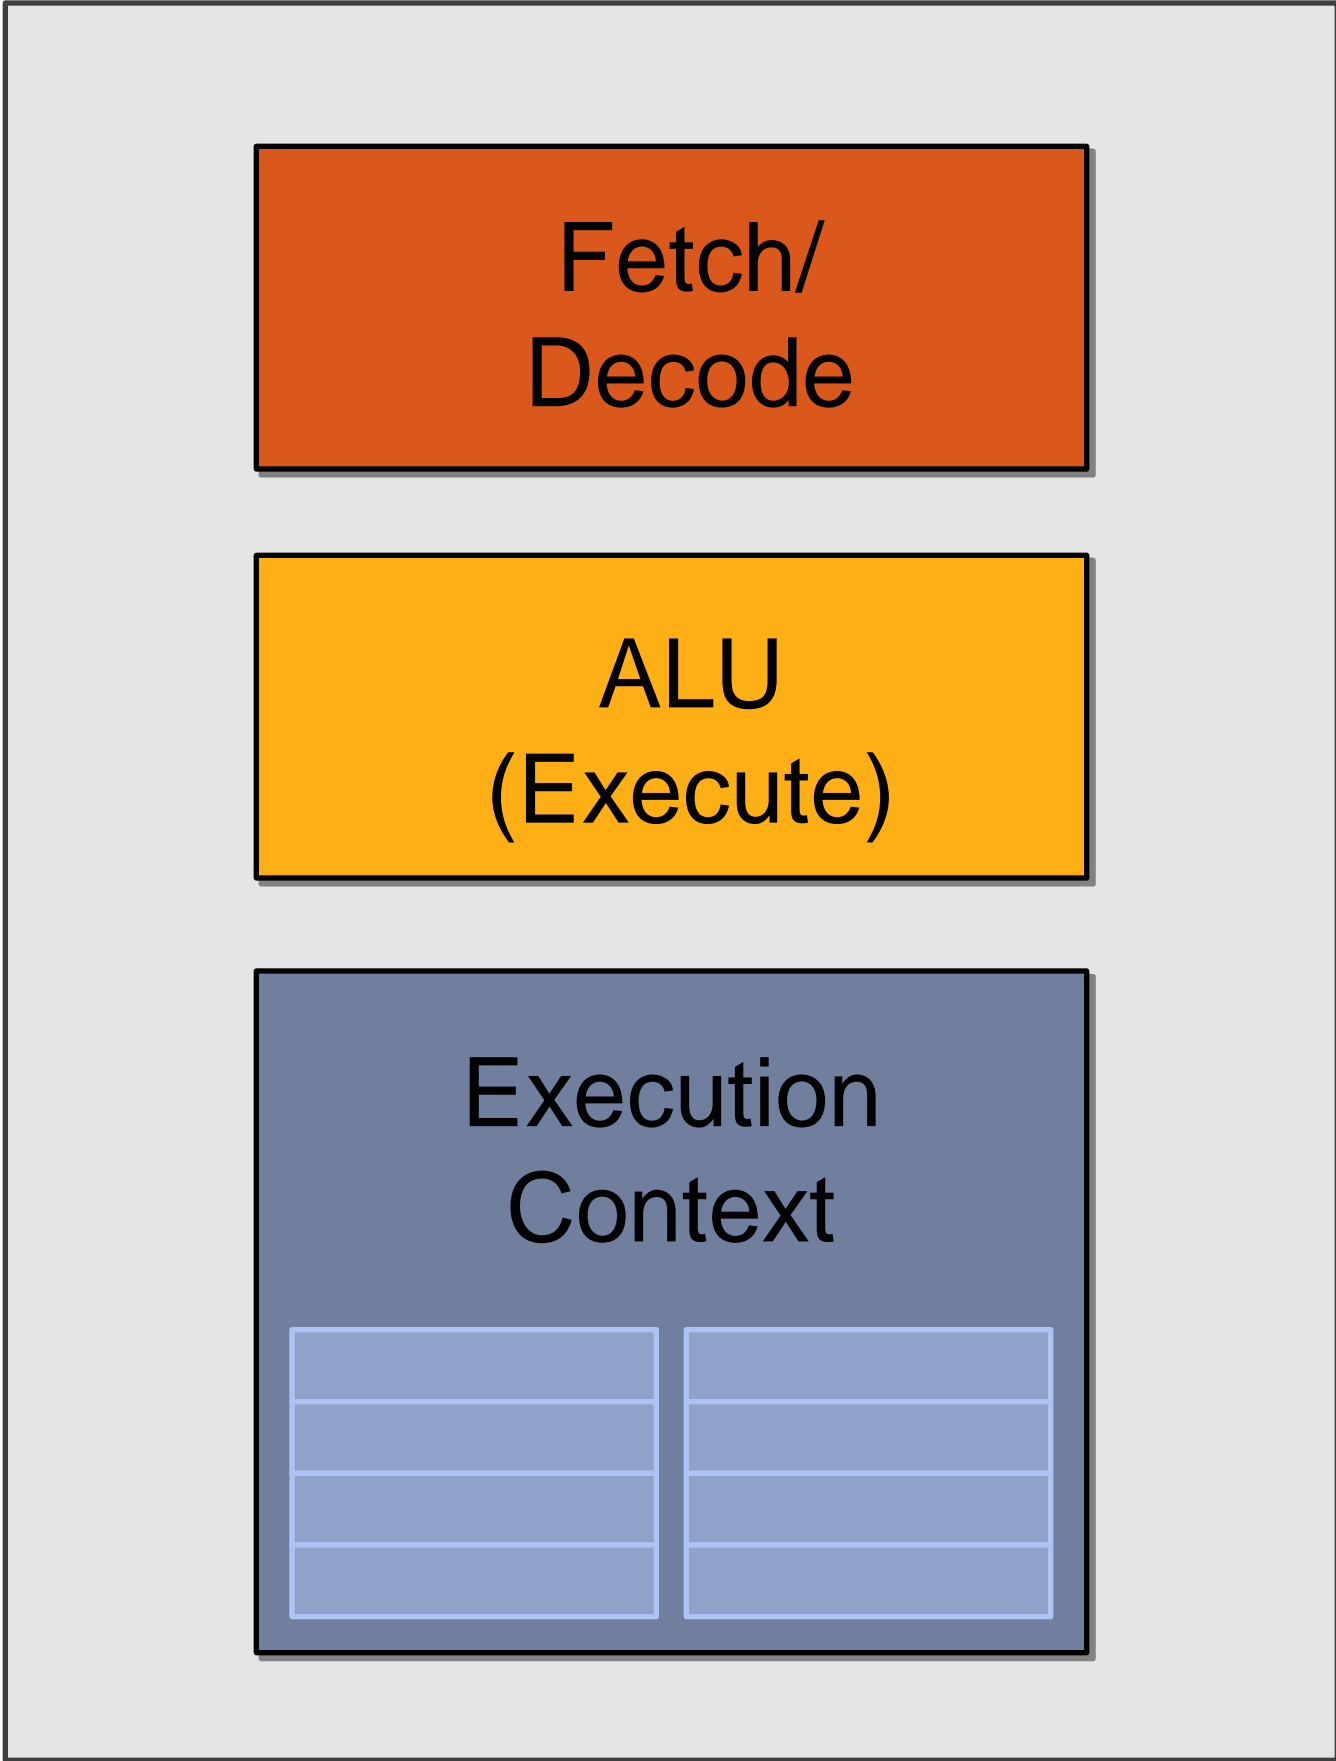
1 shaded fragment output record



# Execute shader

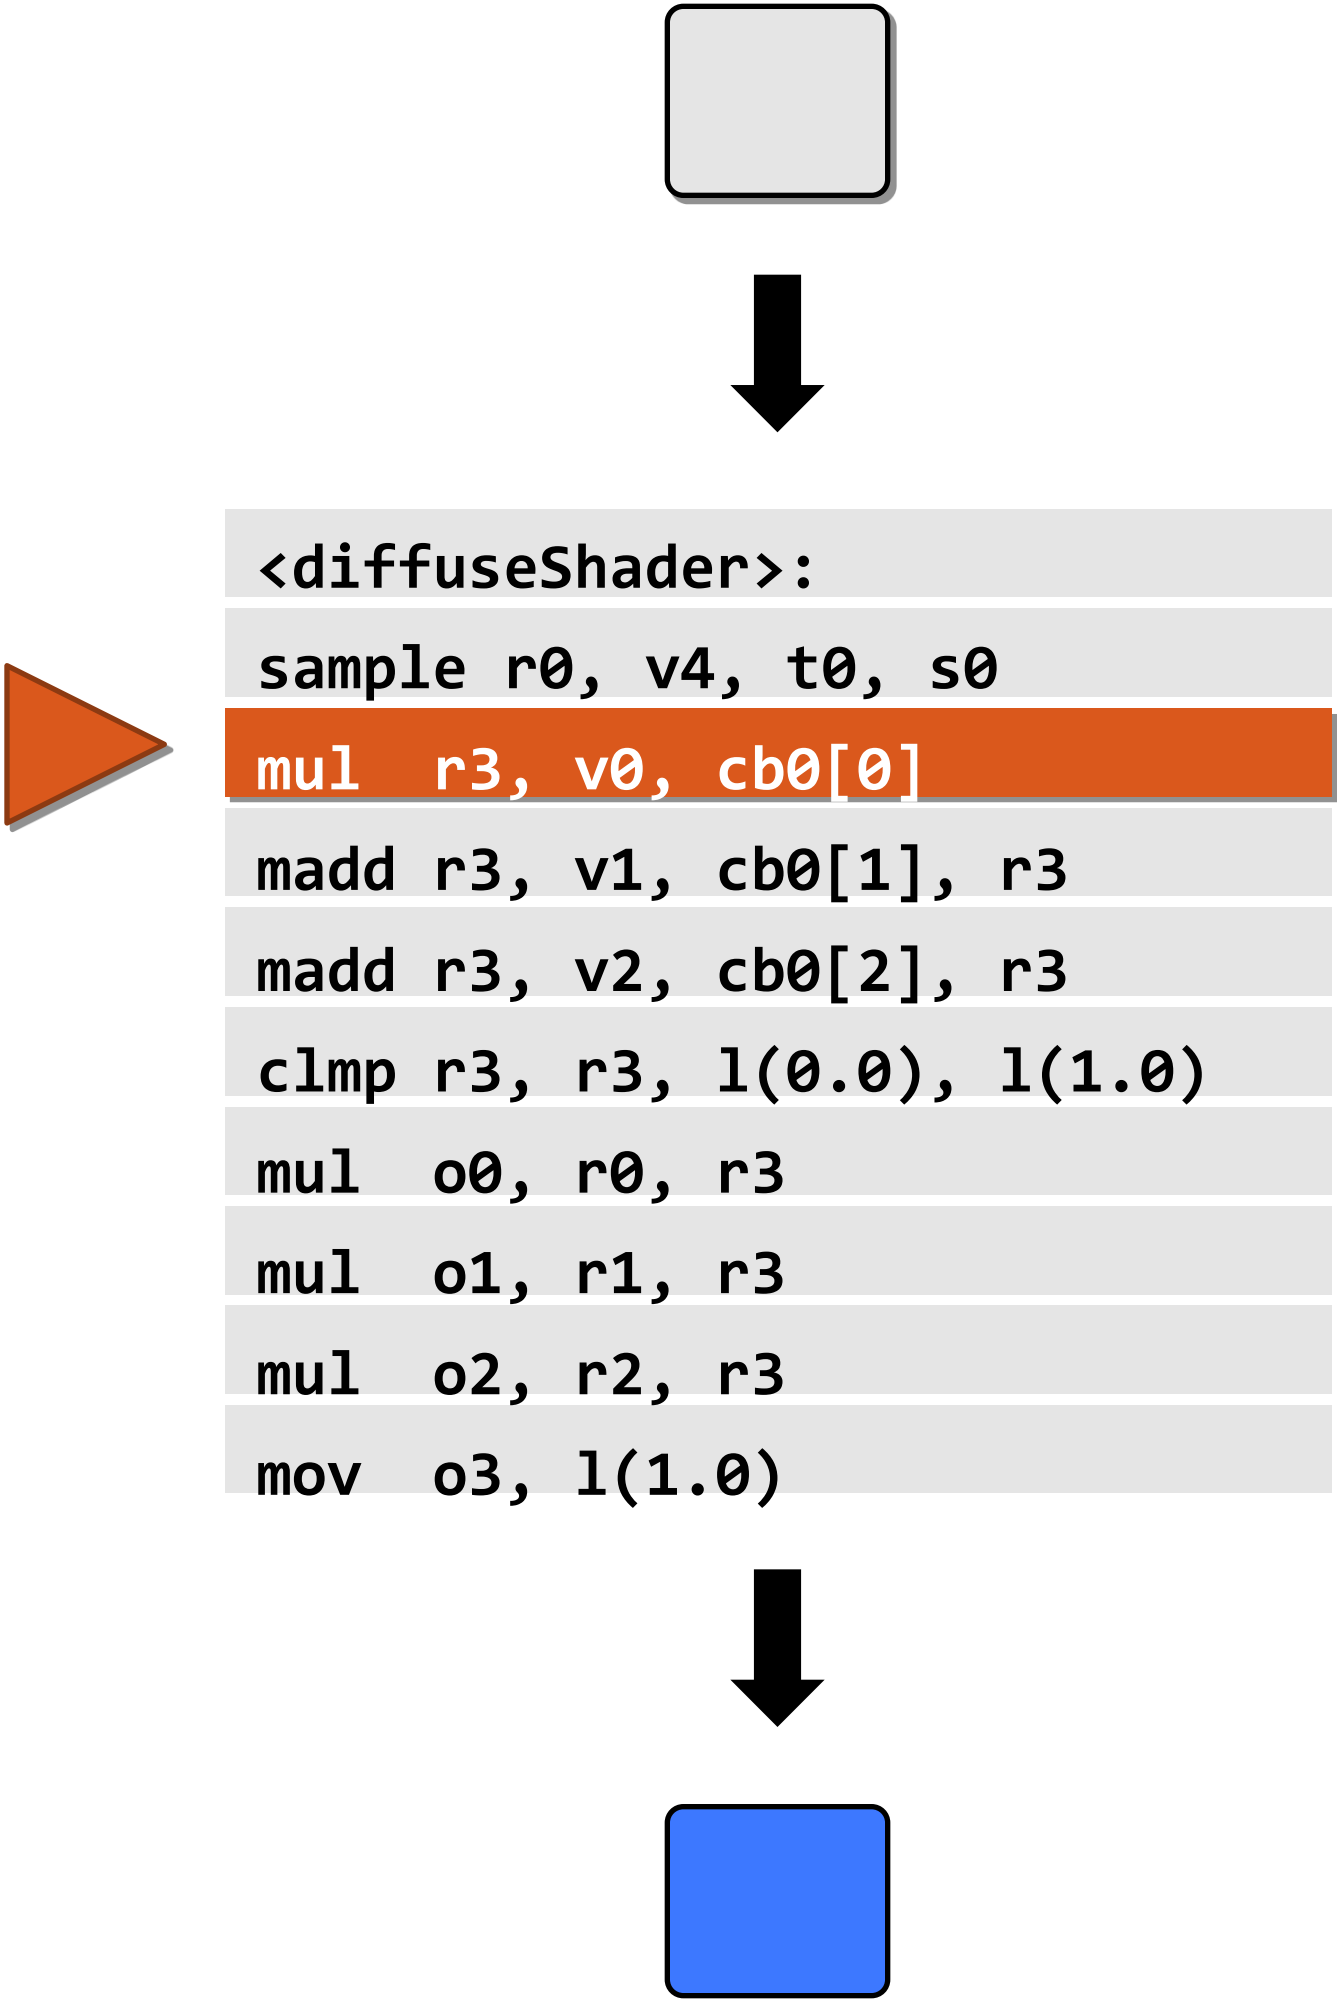
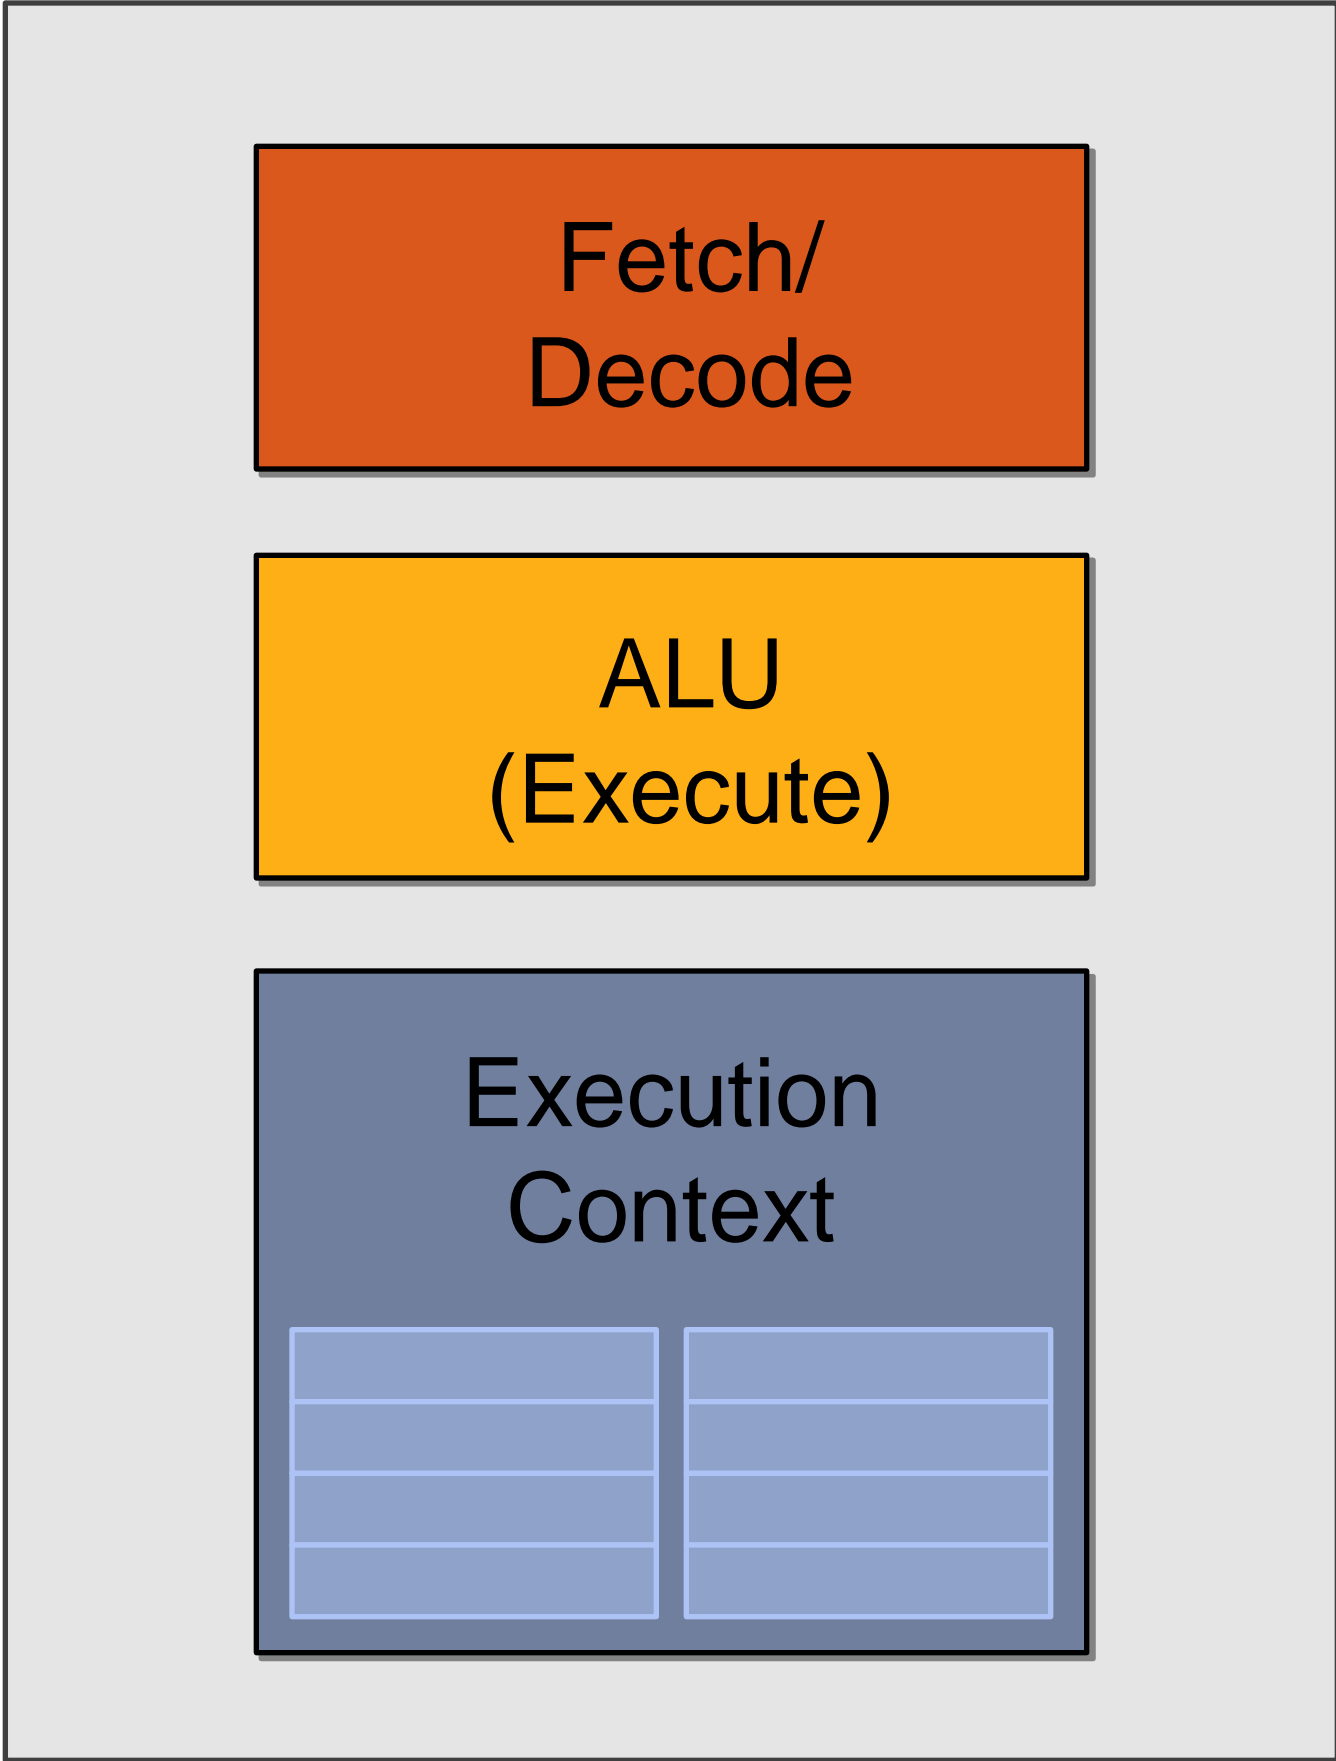


# Execute shader

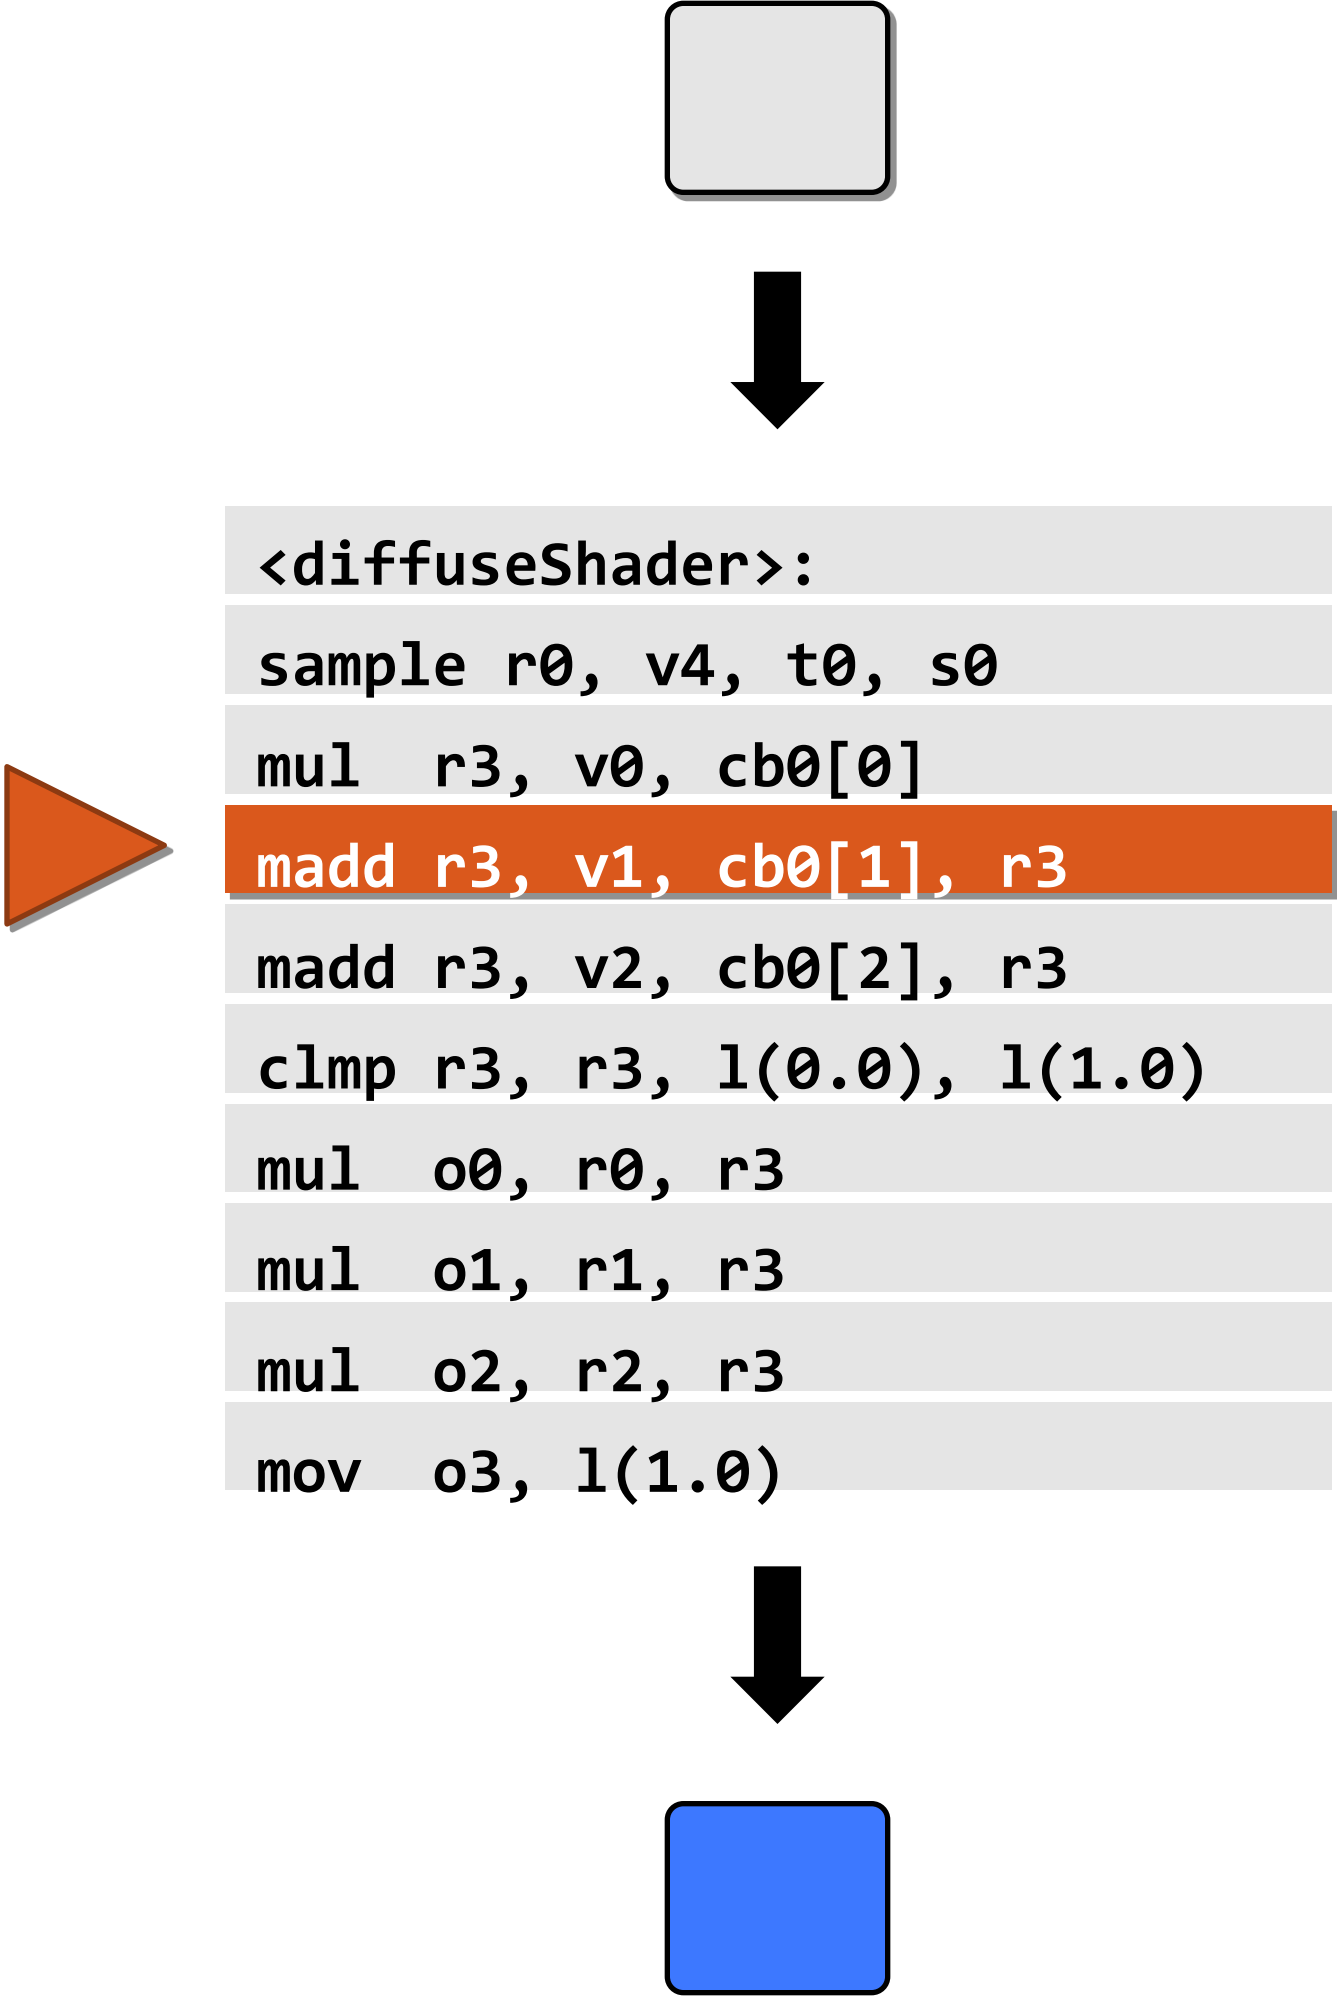
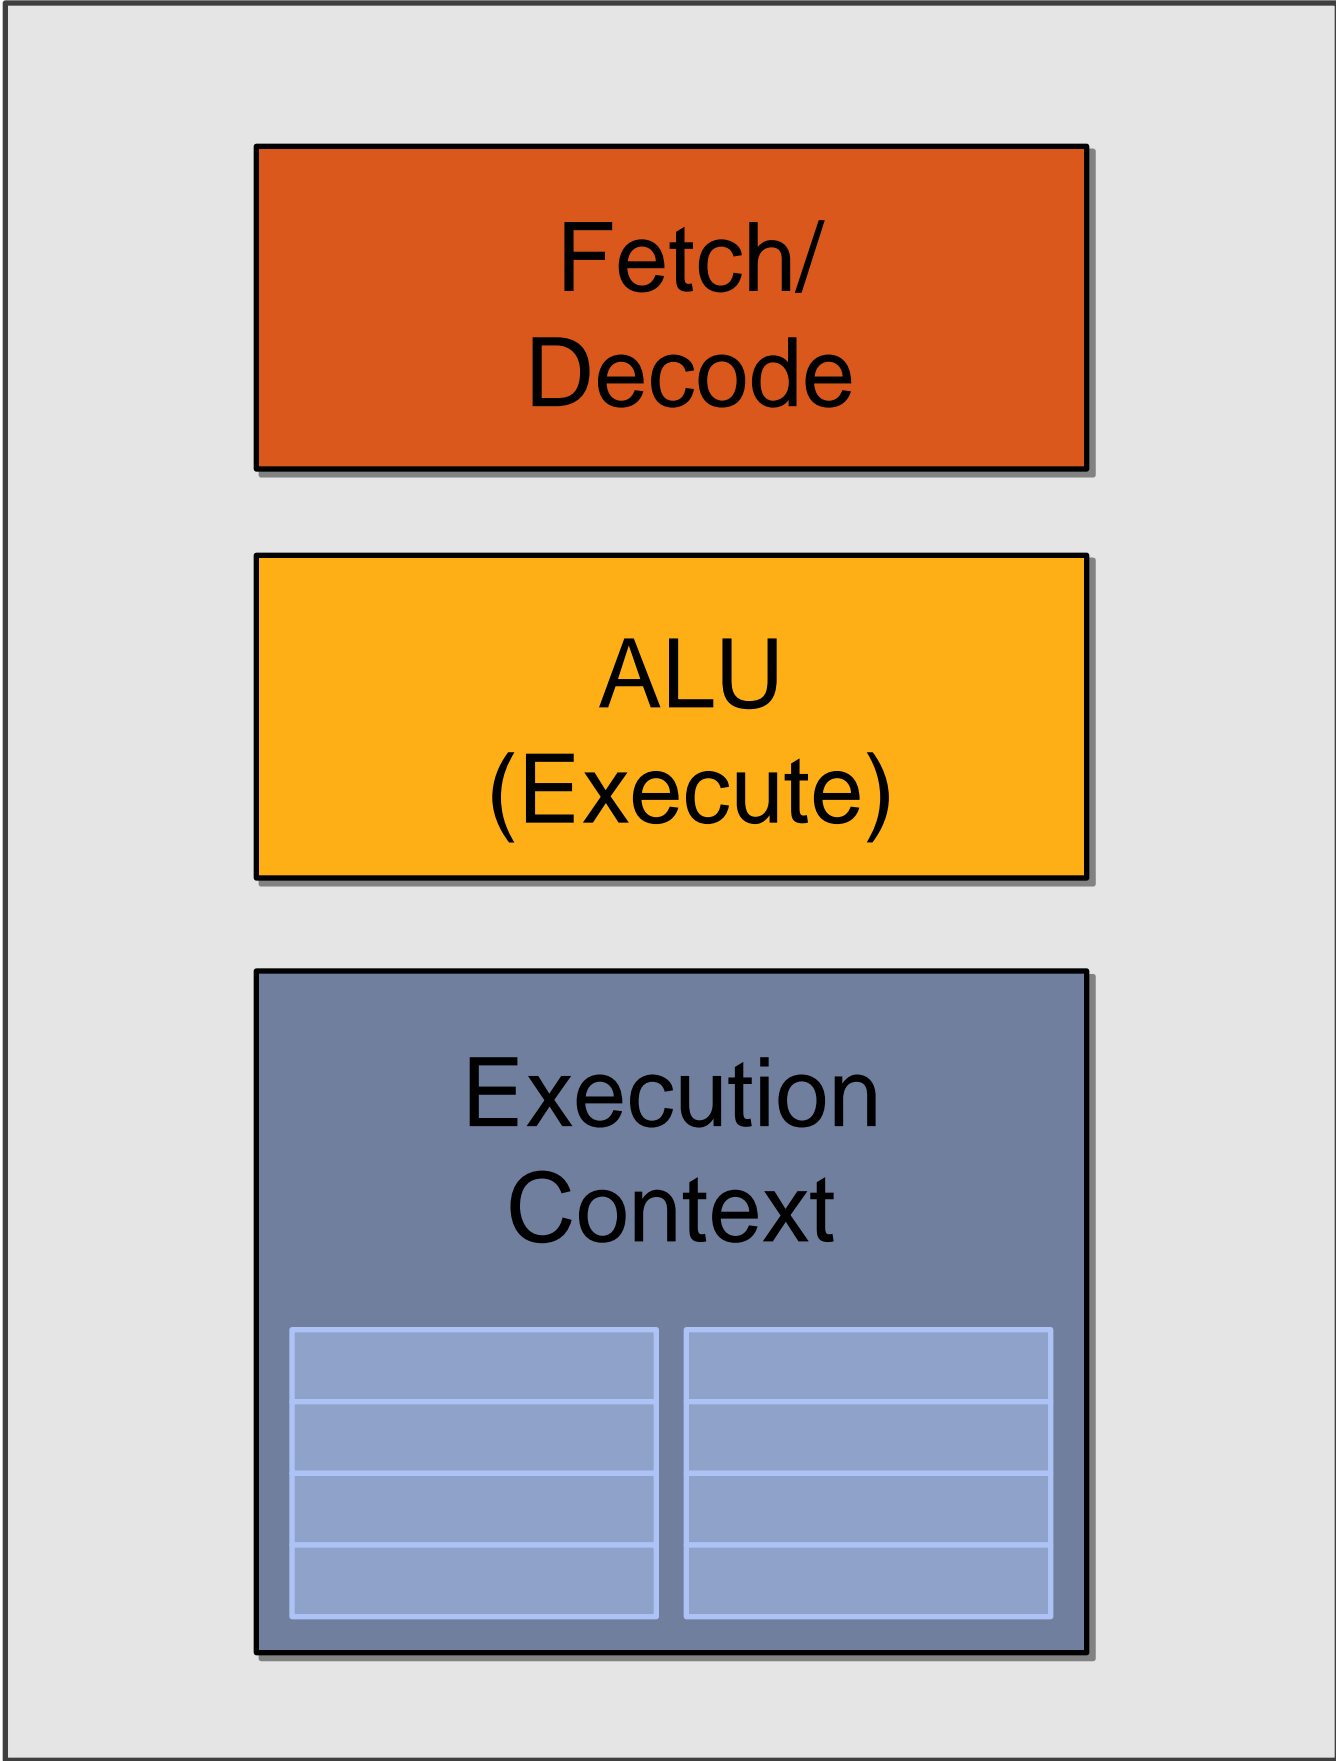




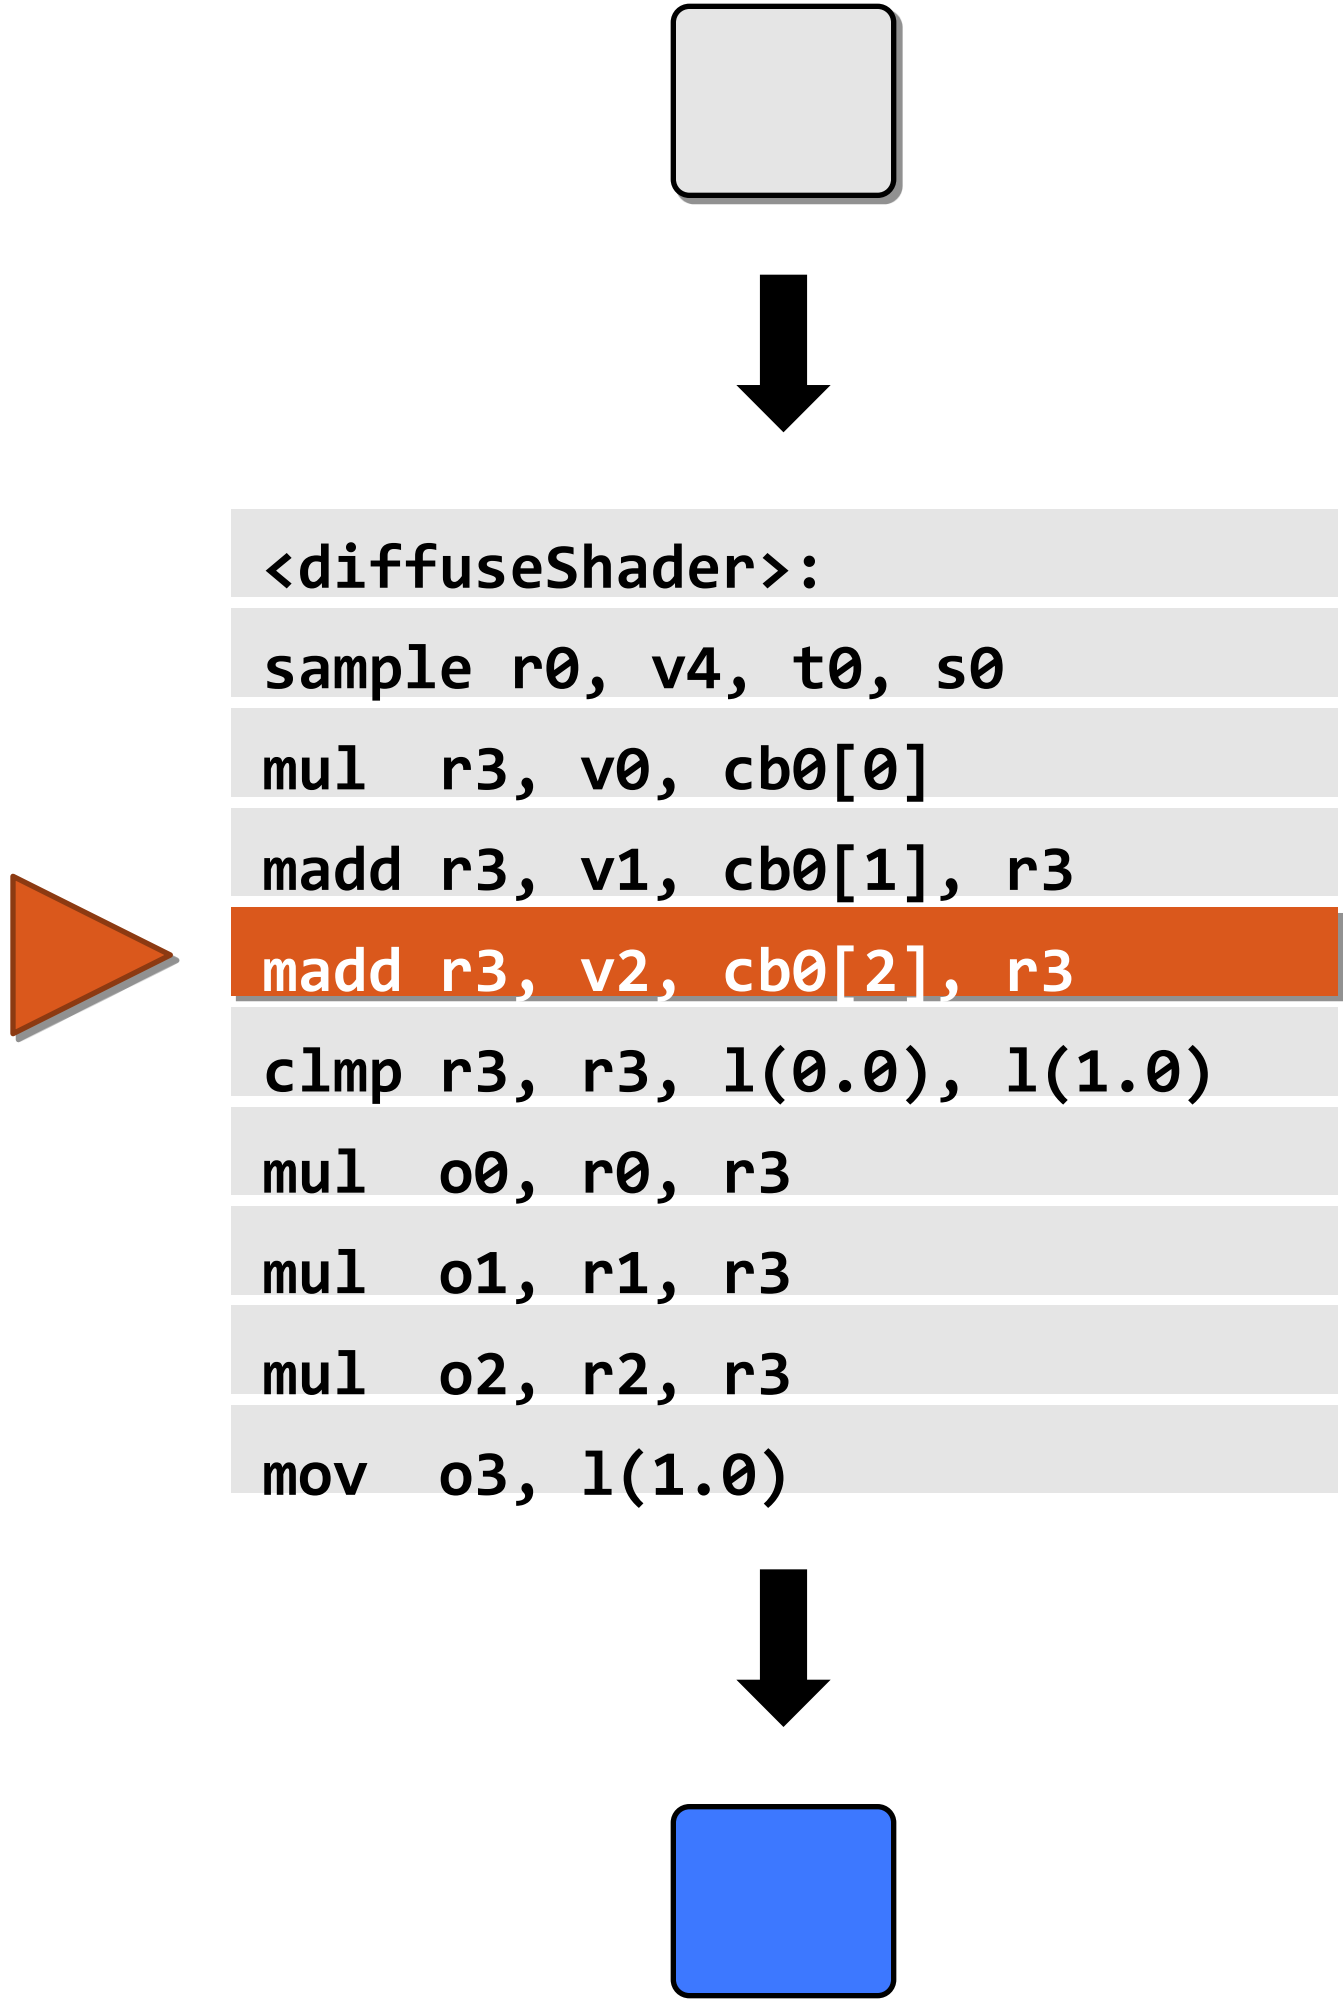
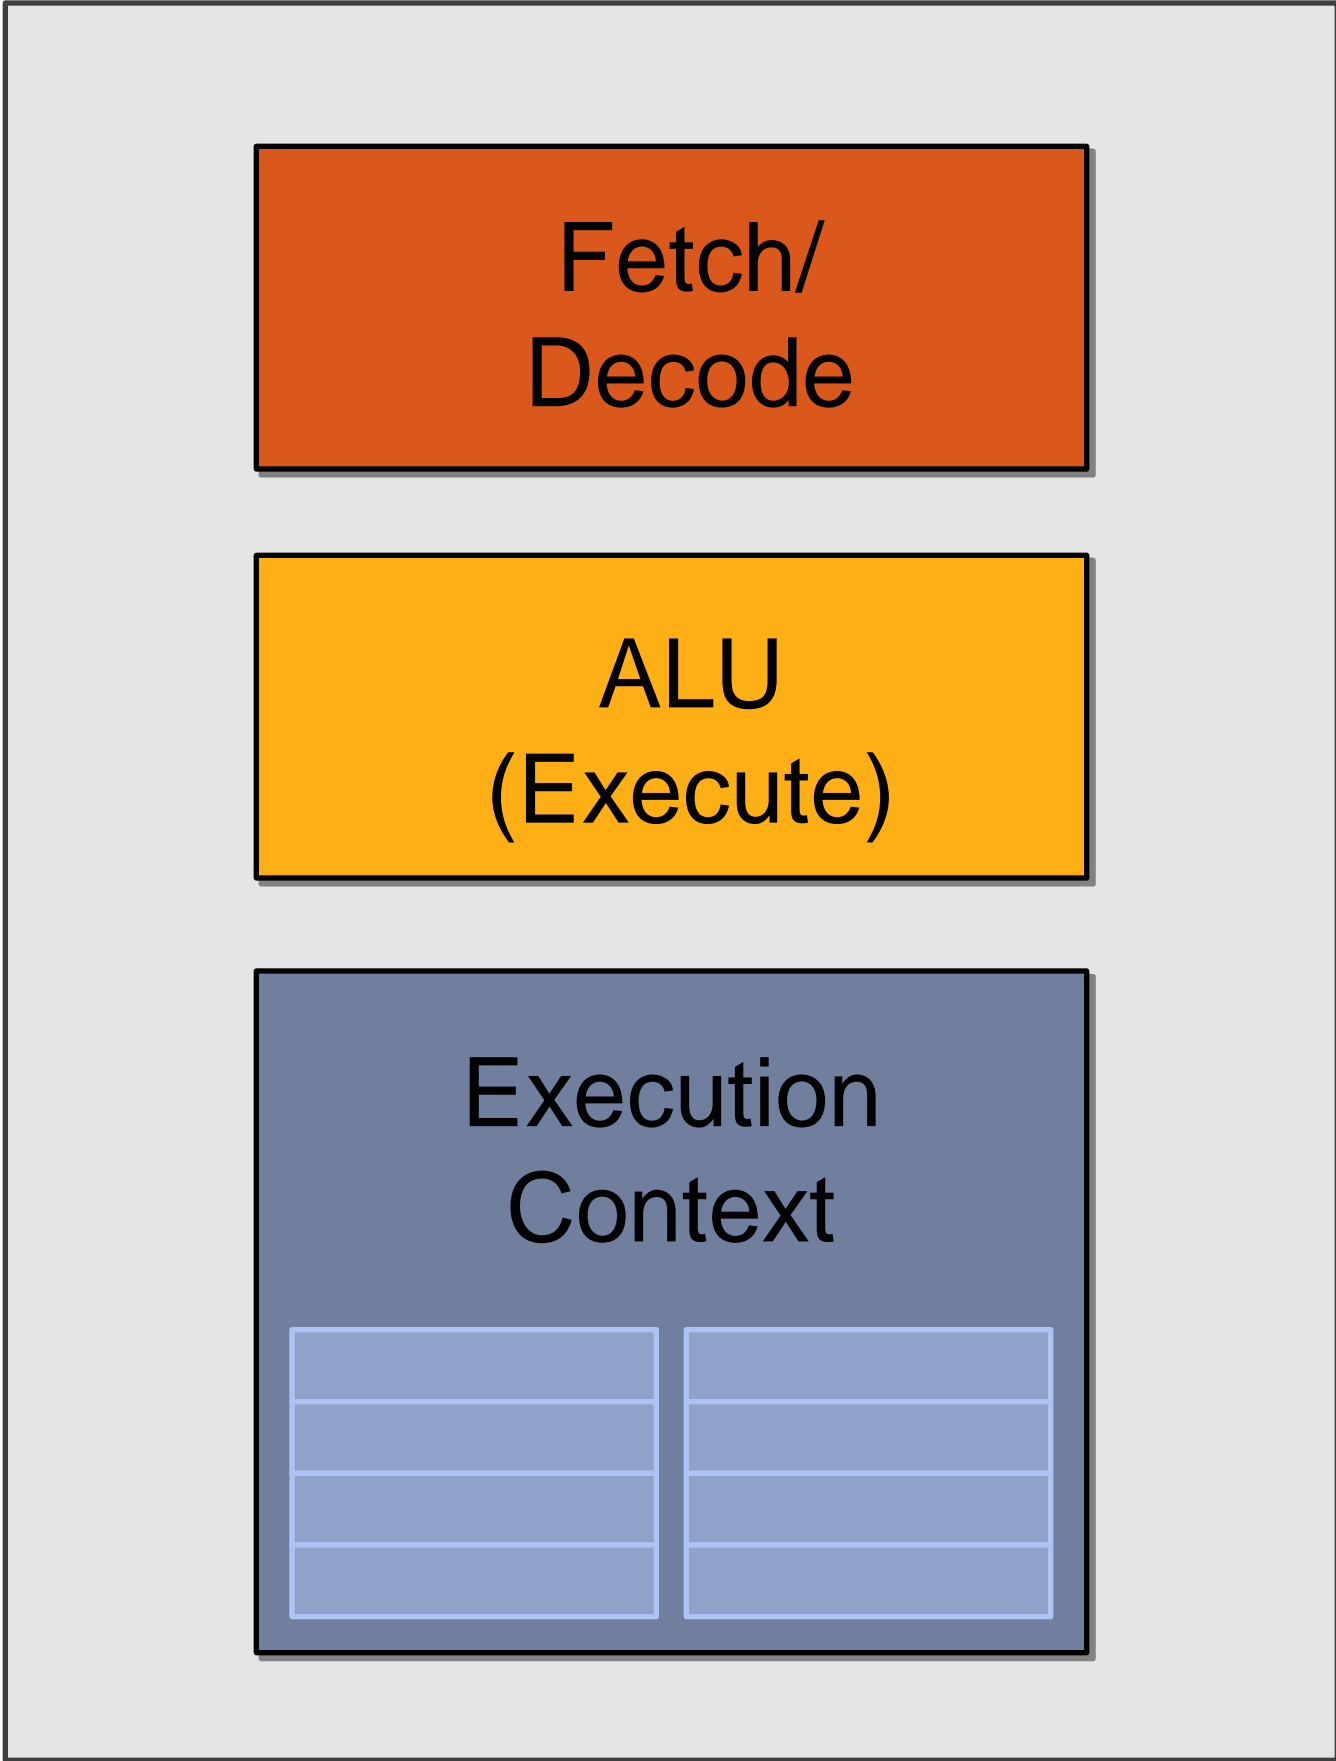
# Execute shader



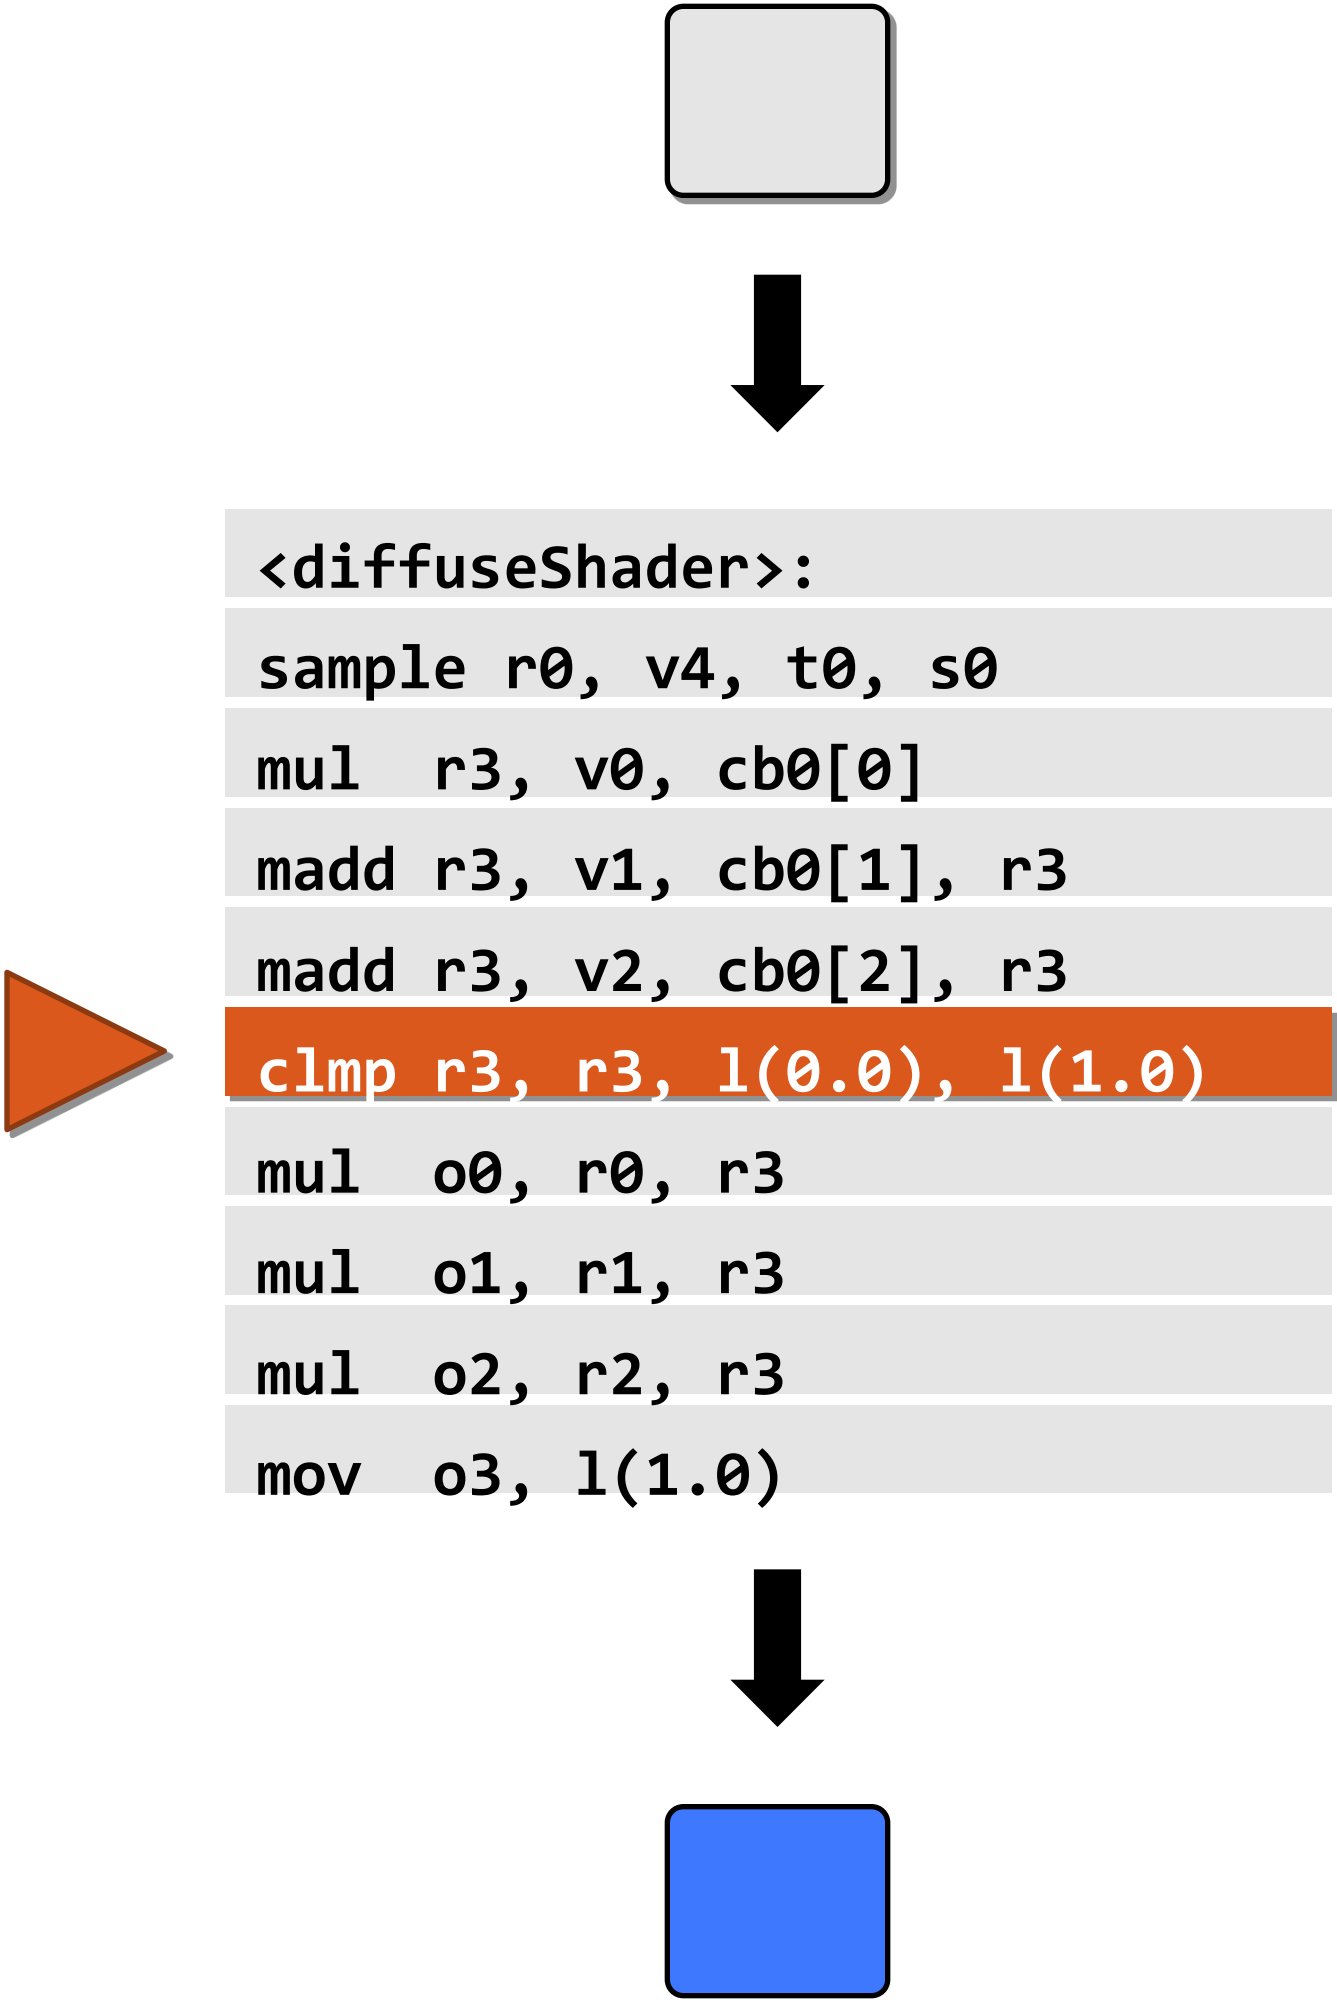
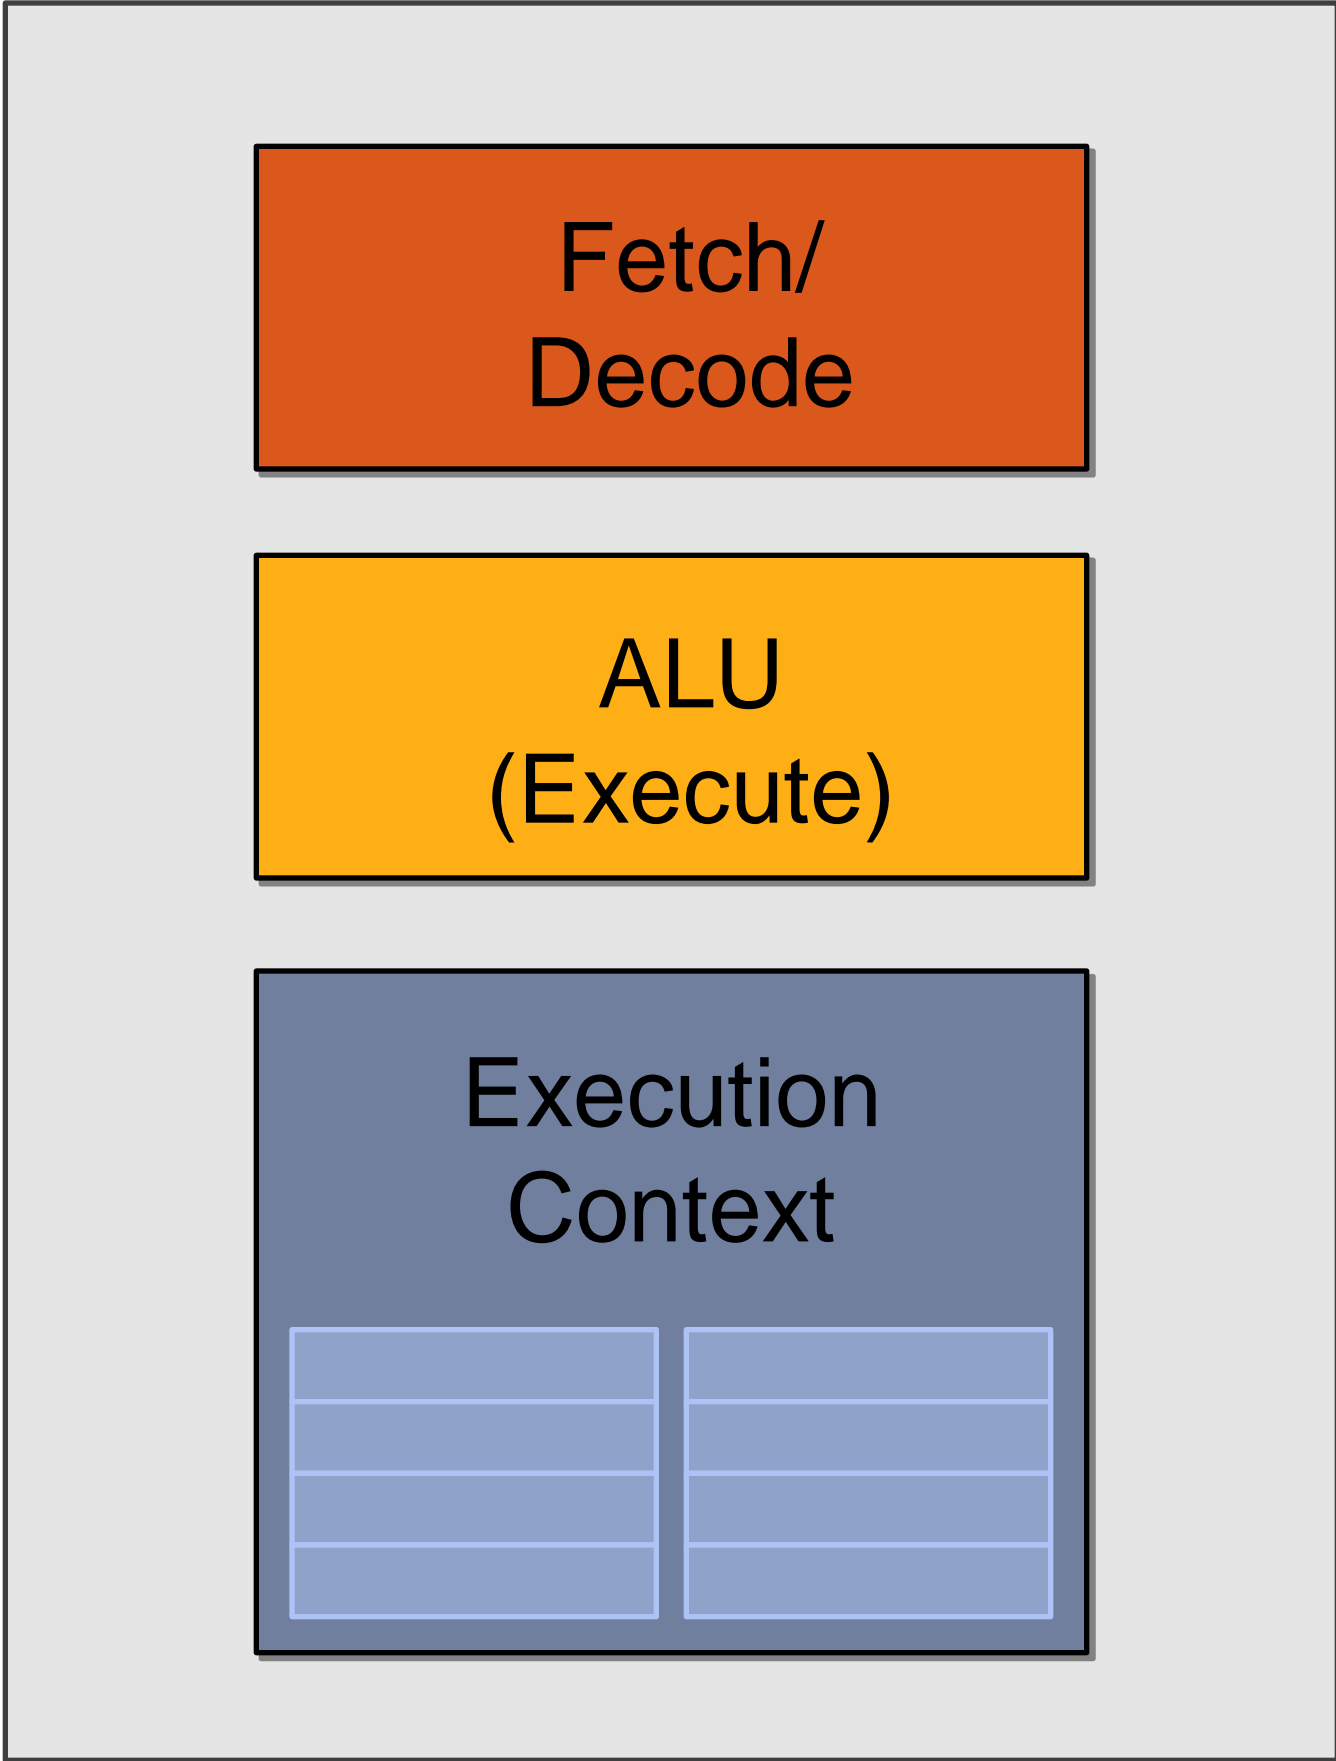
# Execute shader



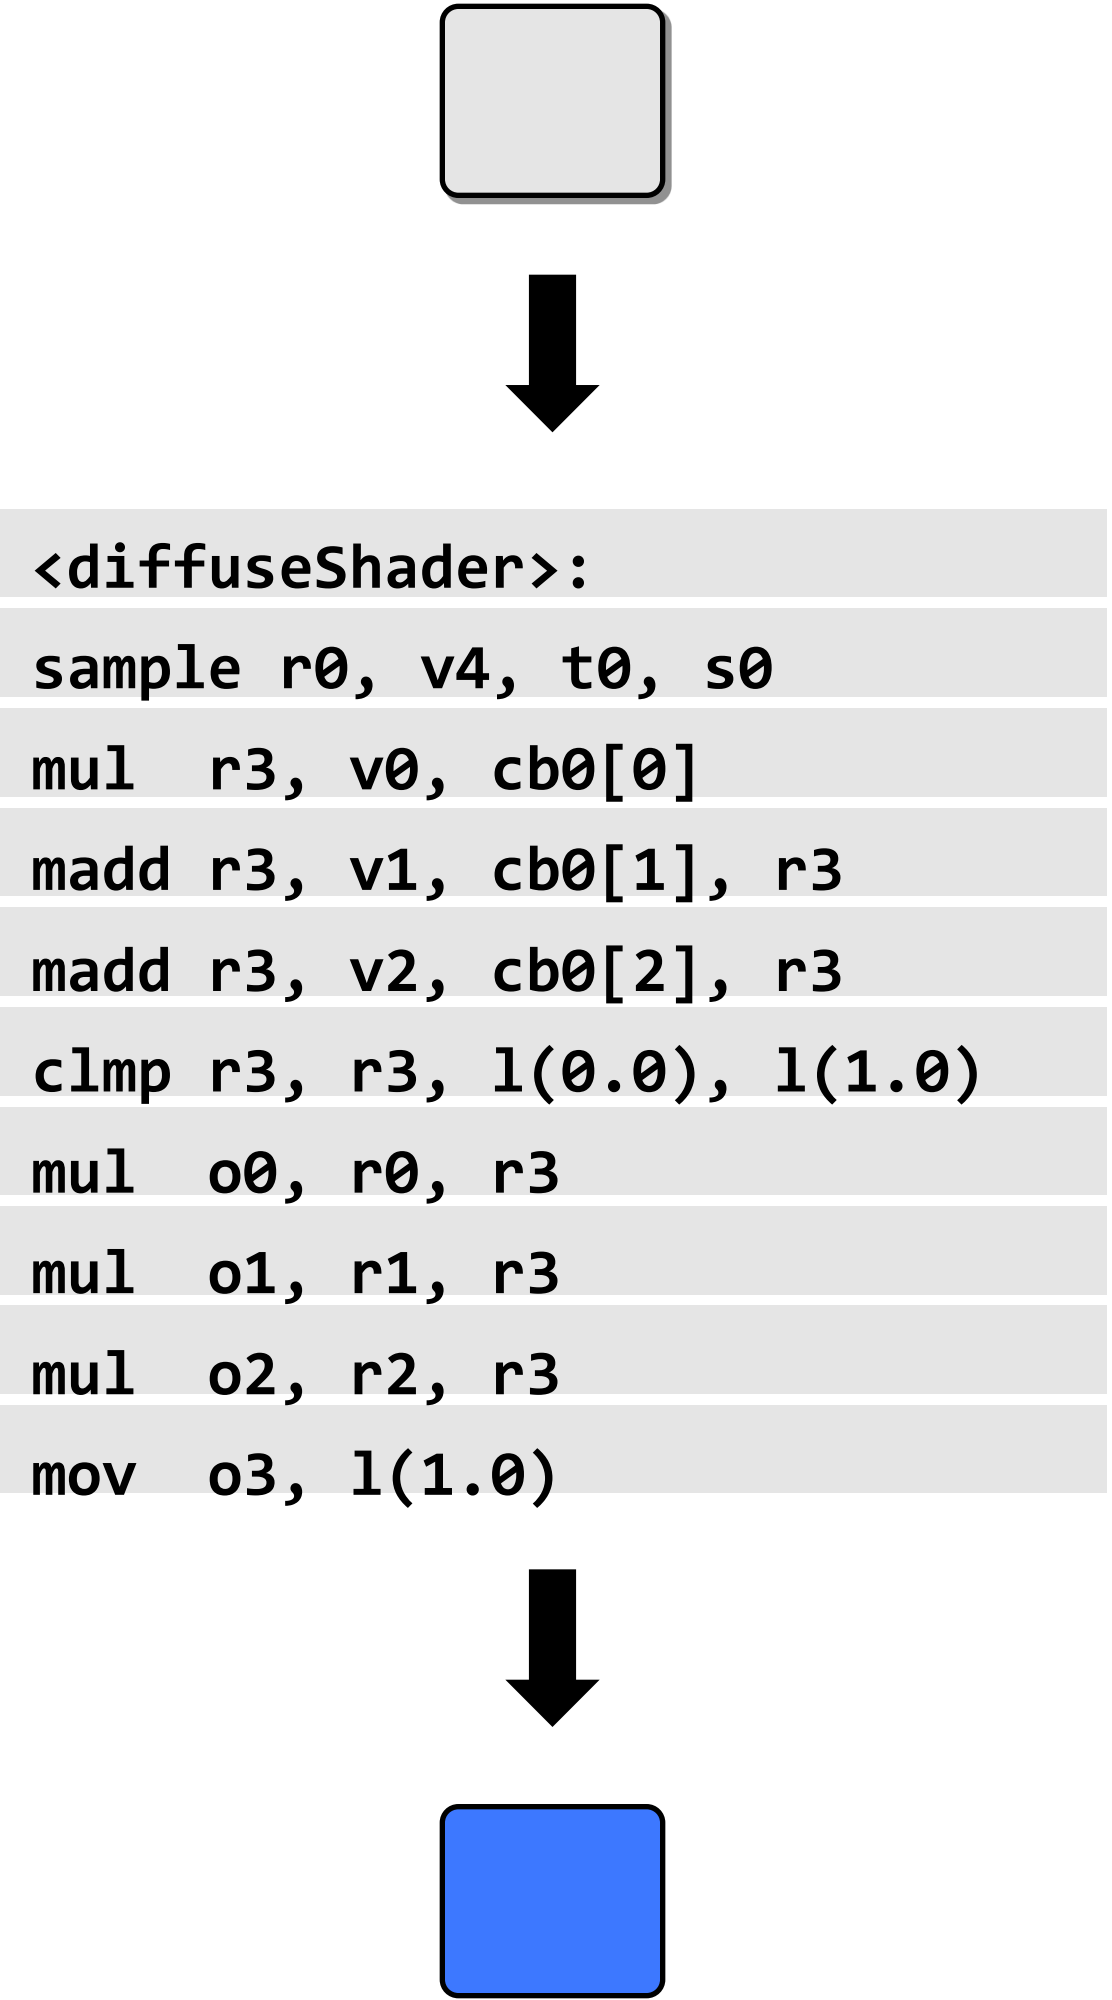
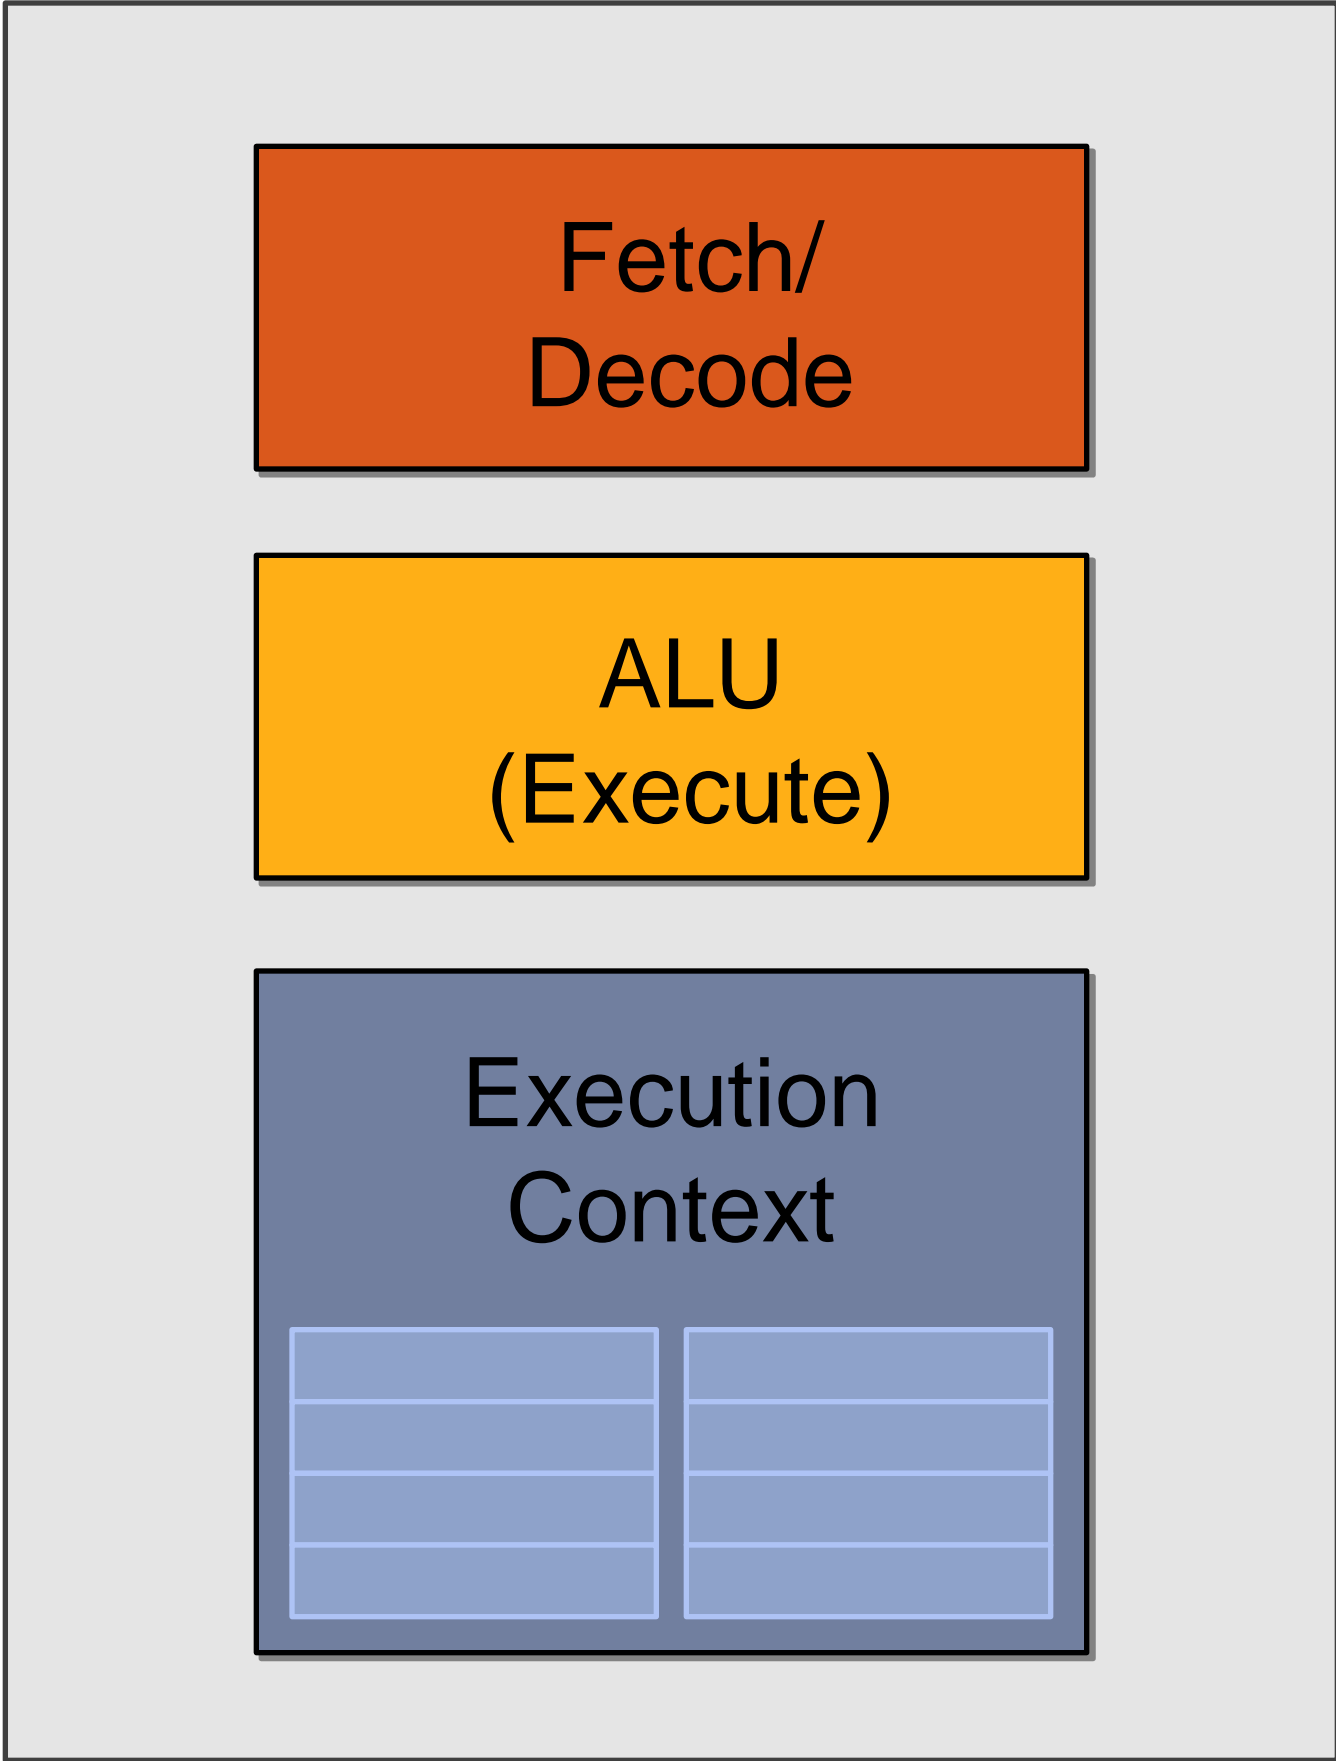
# Execute shader



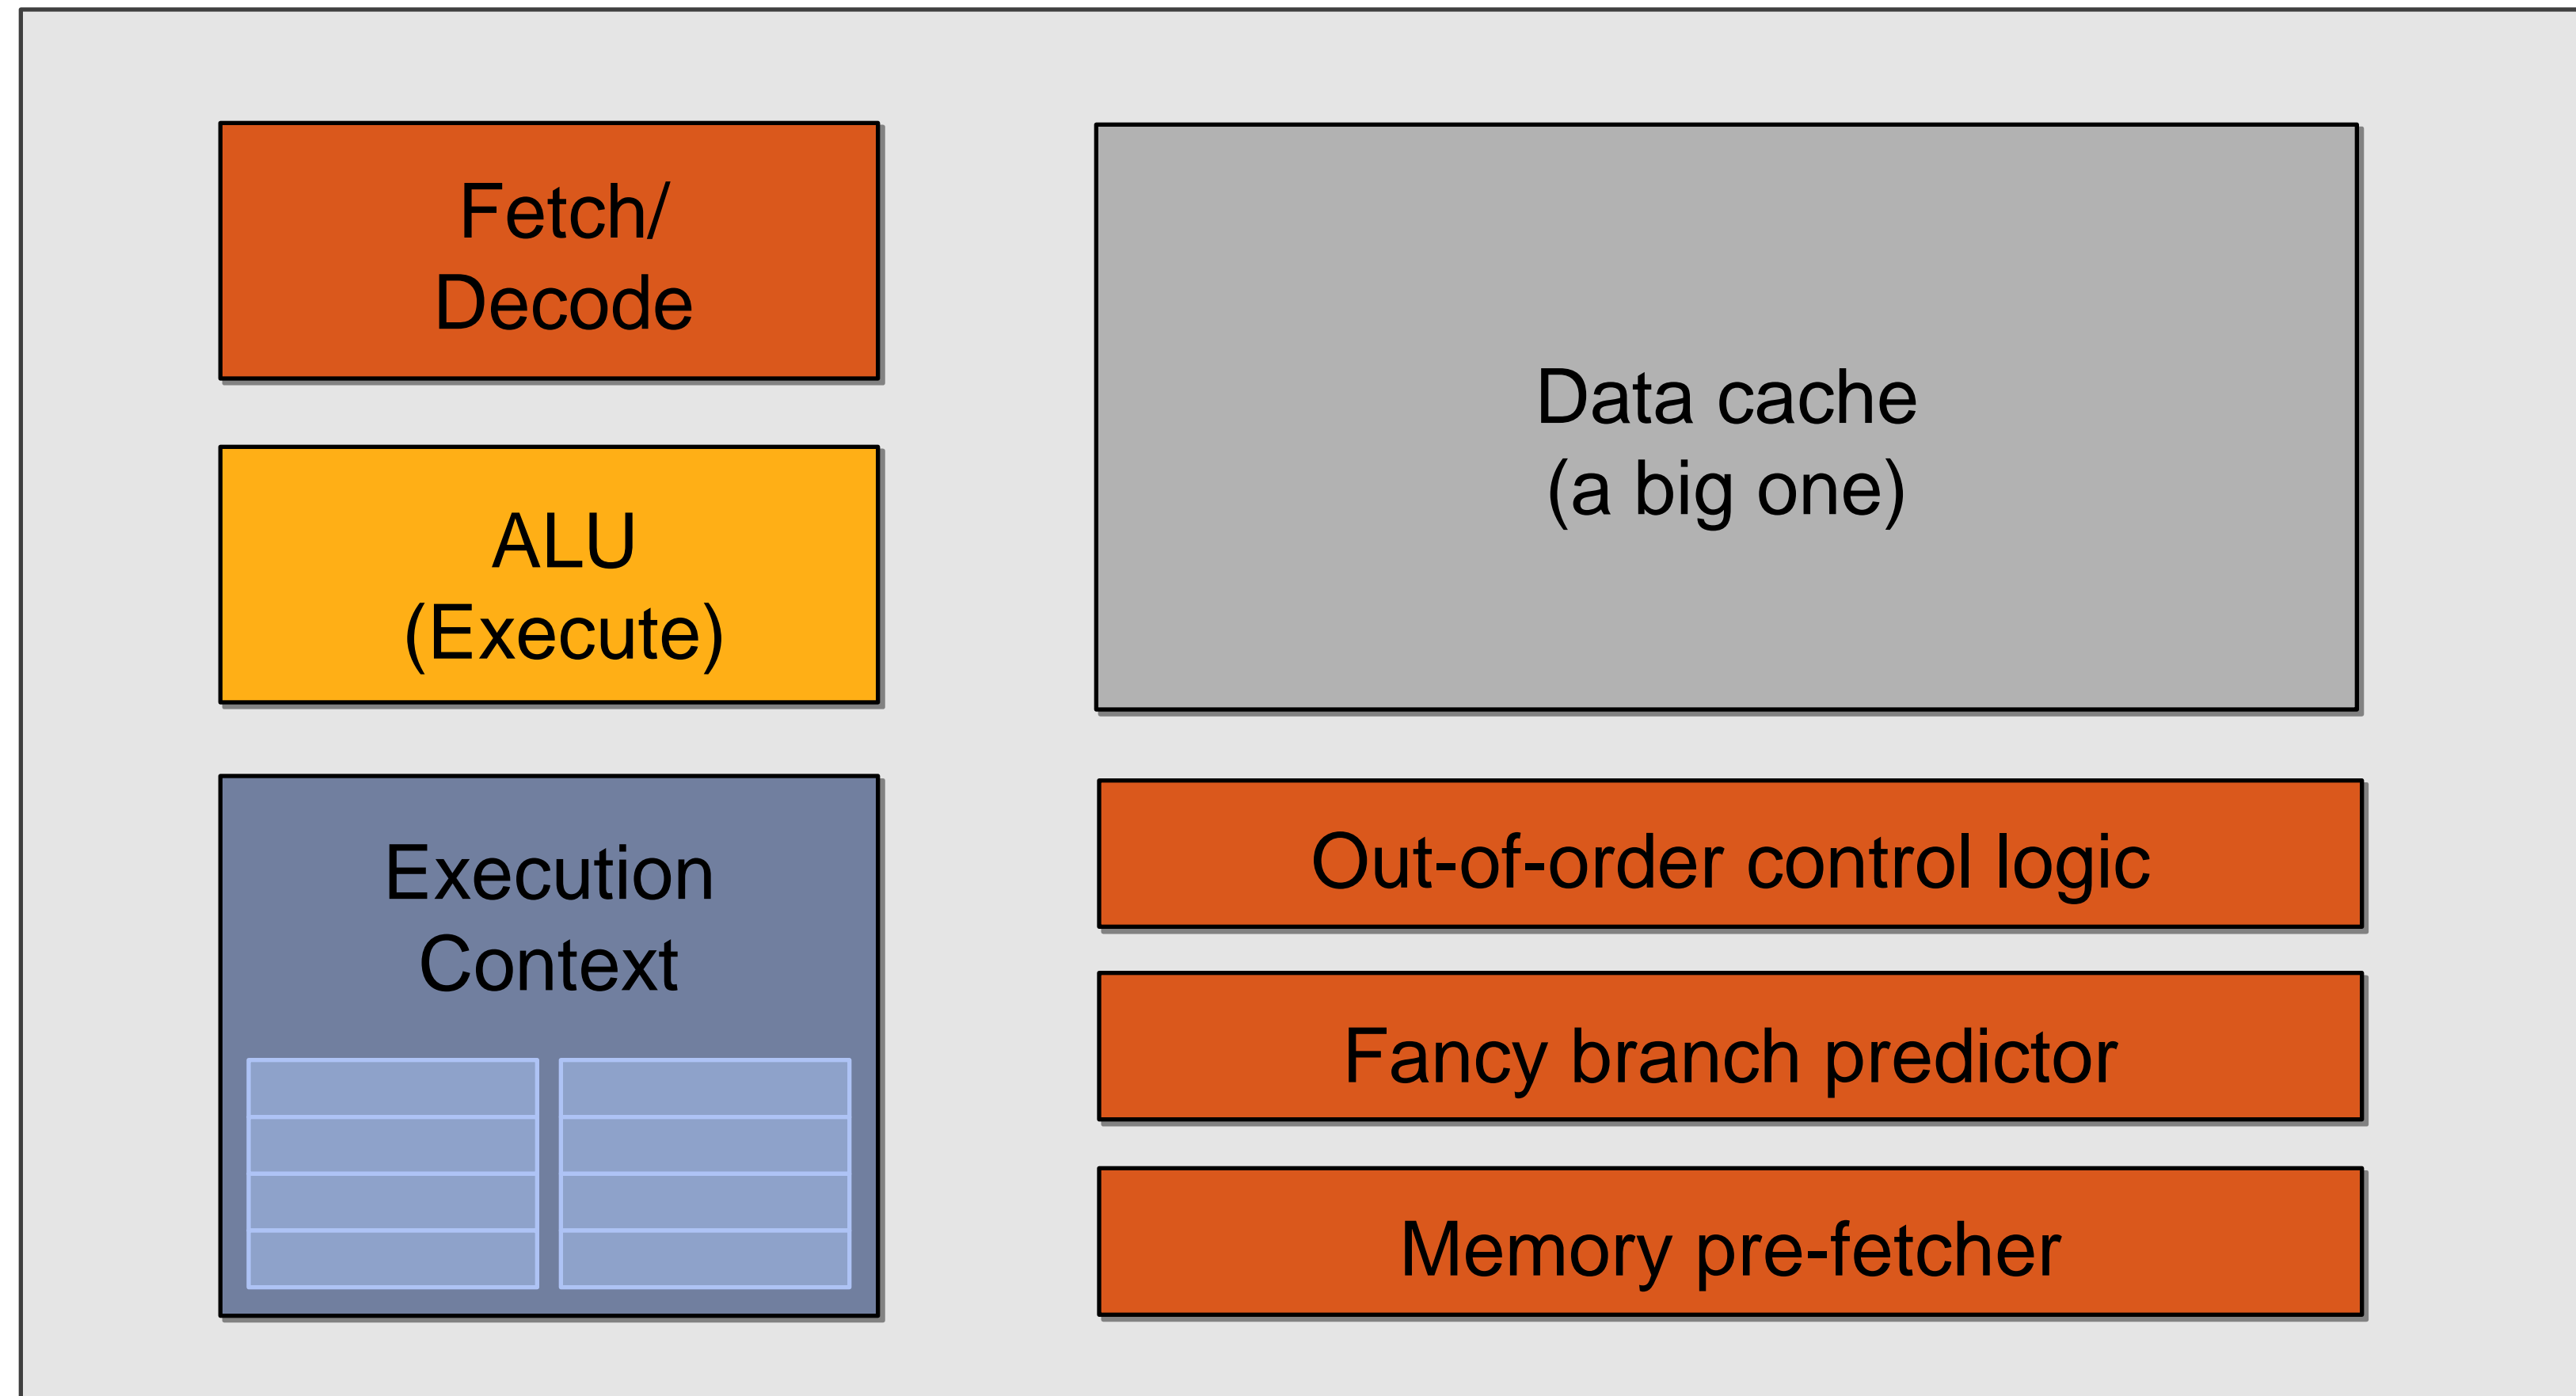
# Execute shader




# Execute shader



# “CPU-style” cores

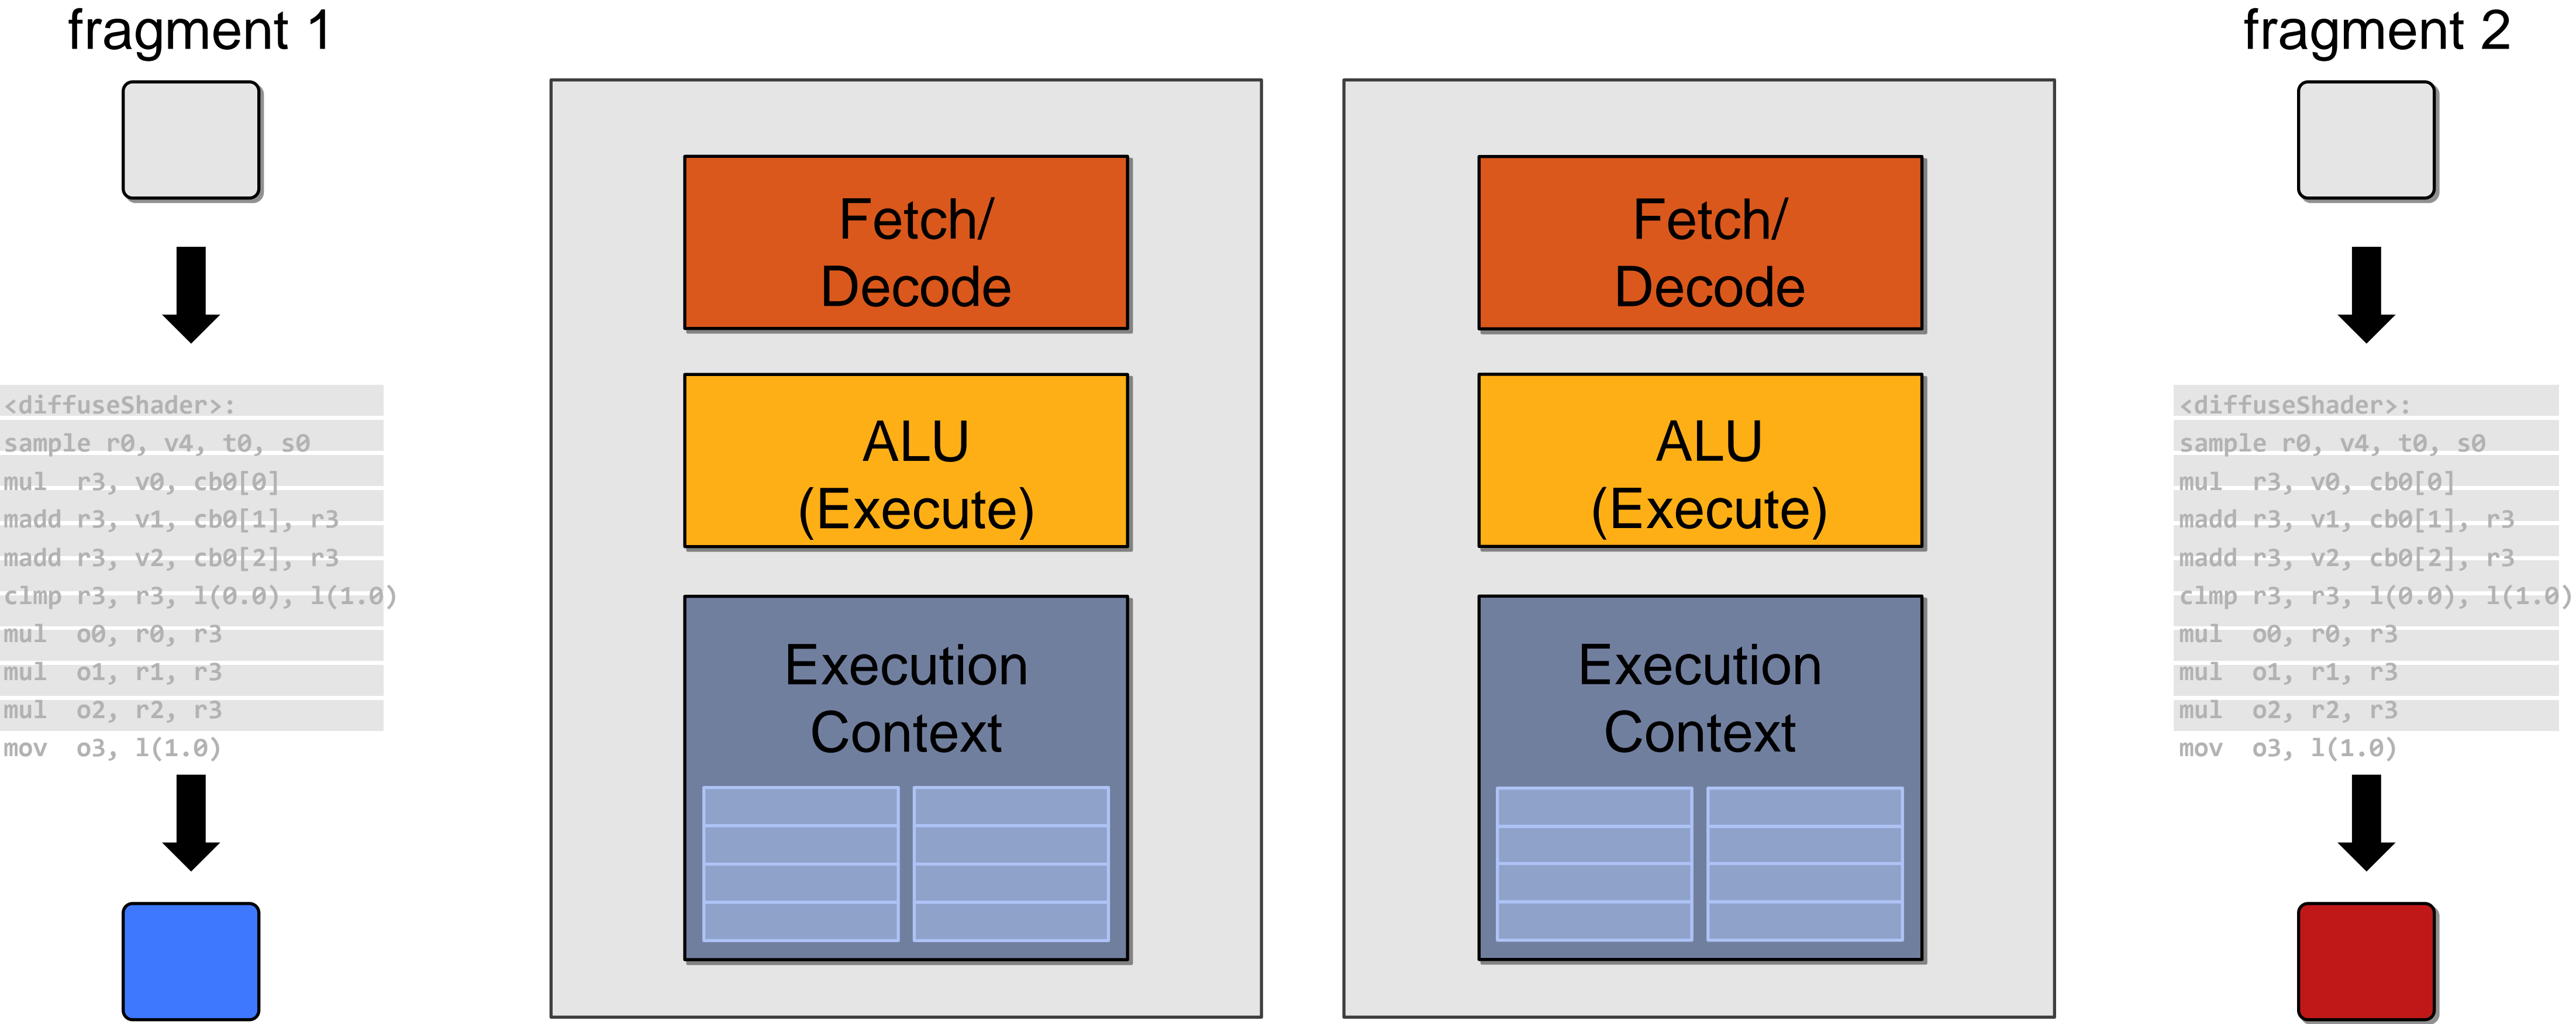


# Slimming down

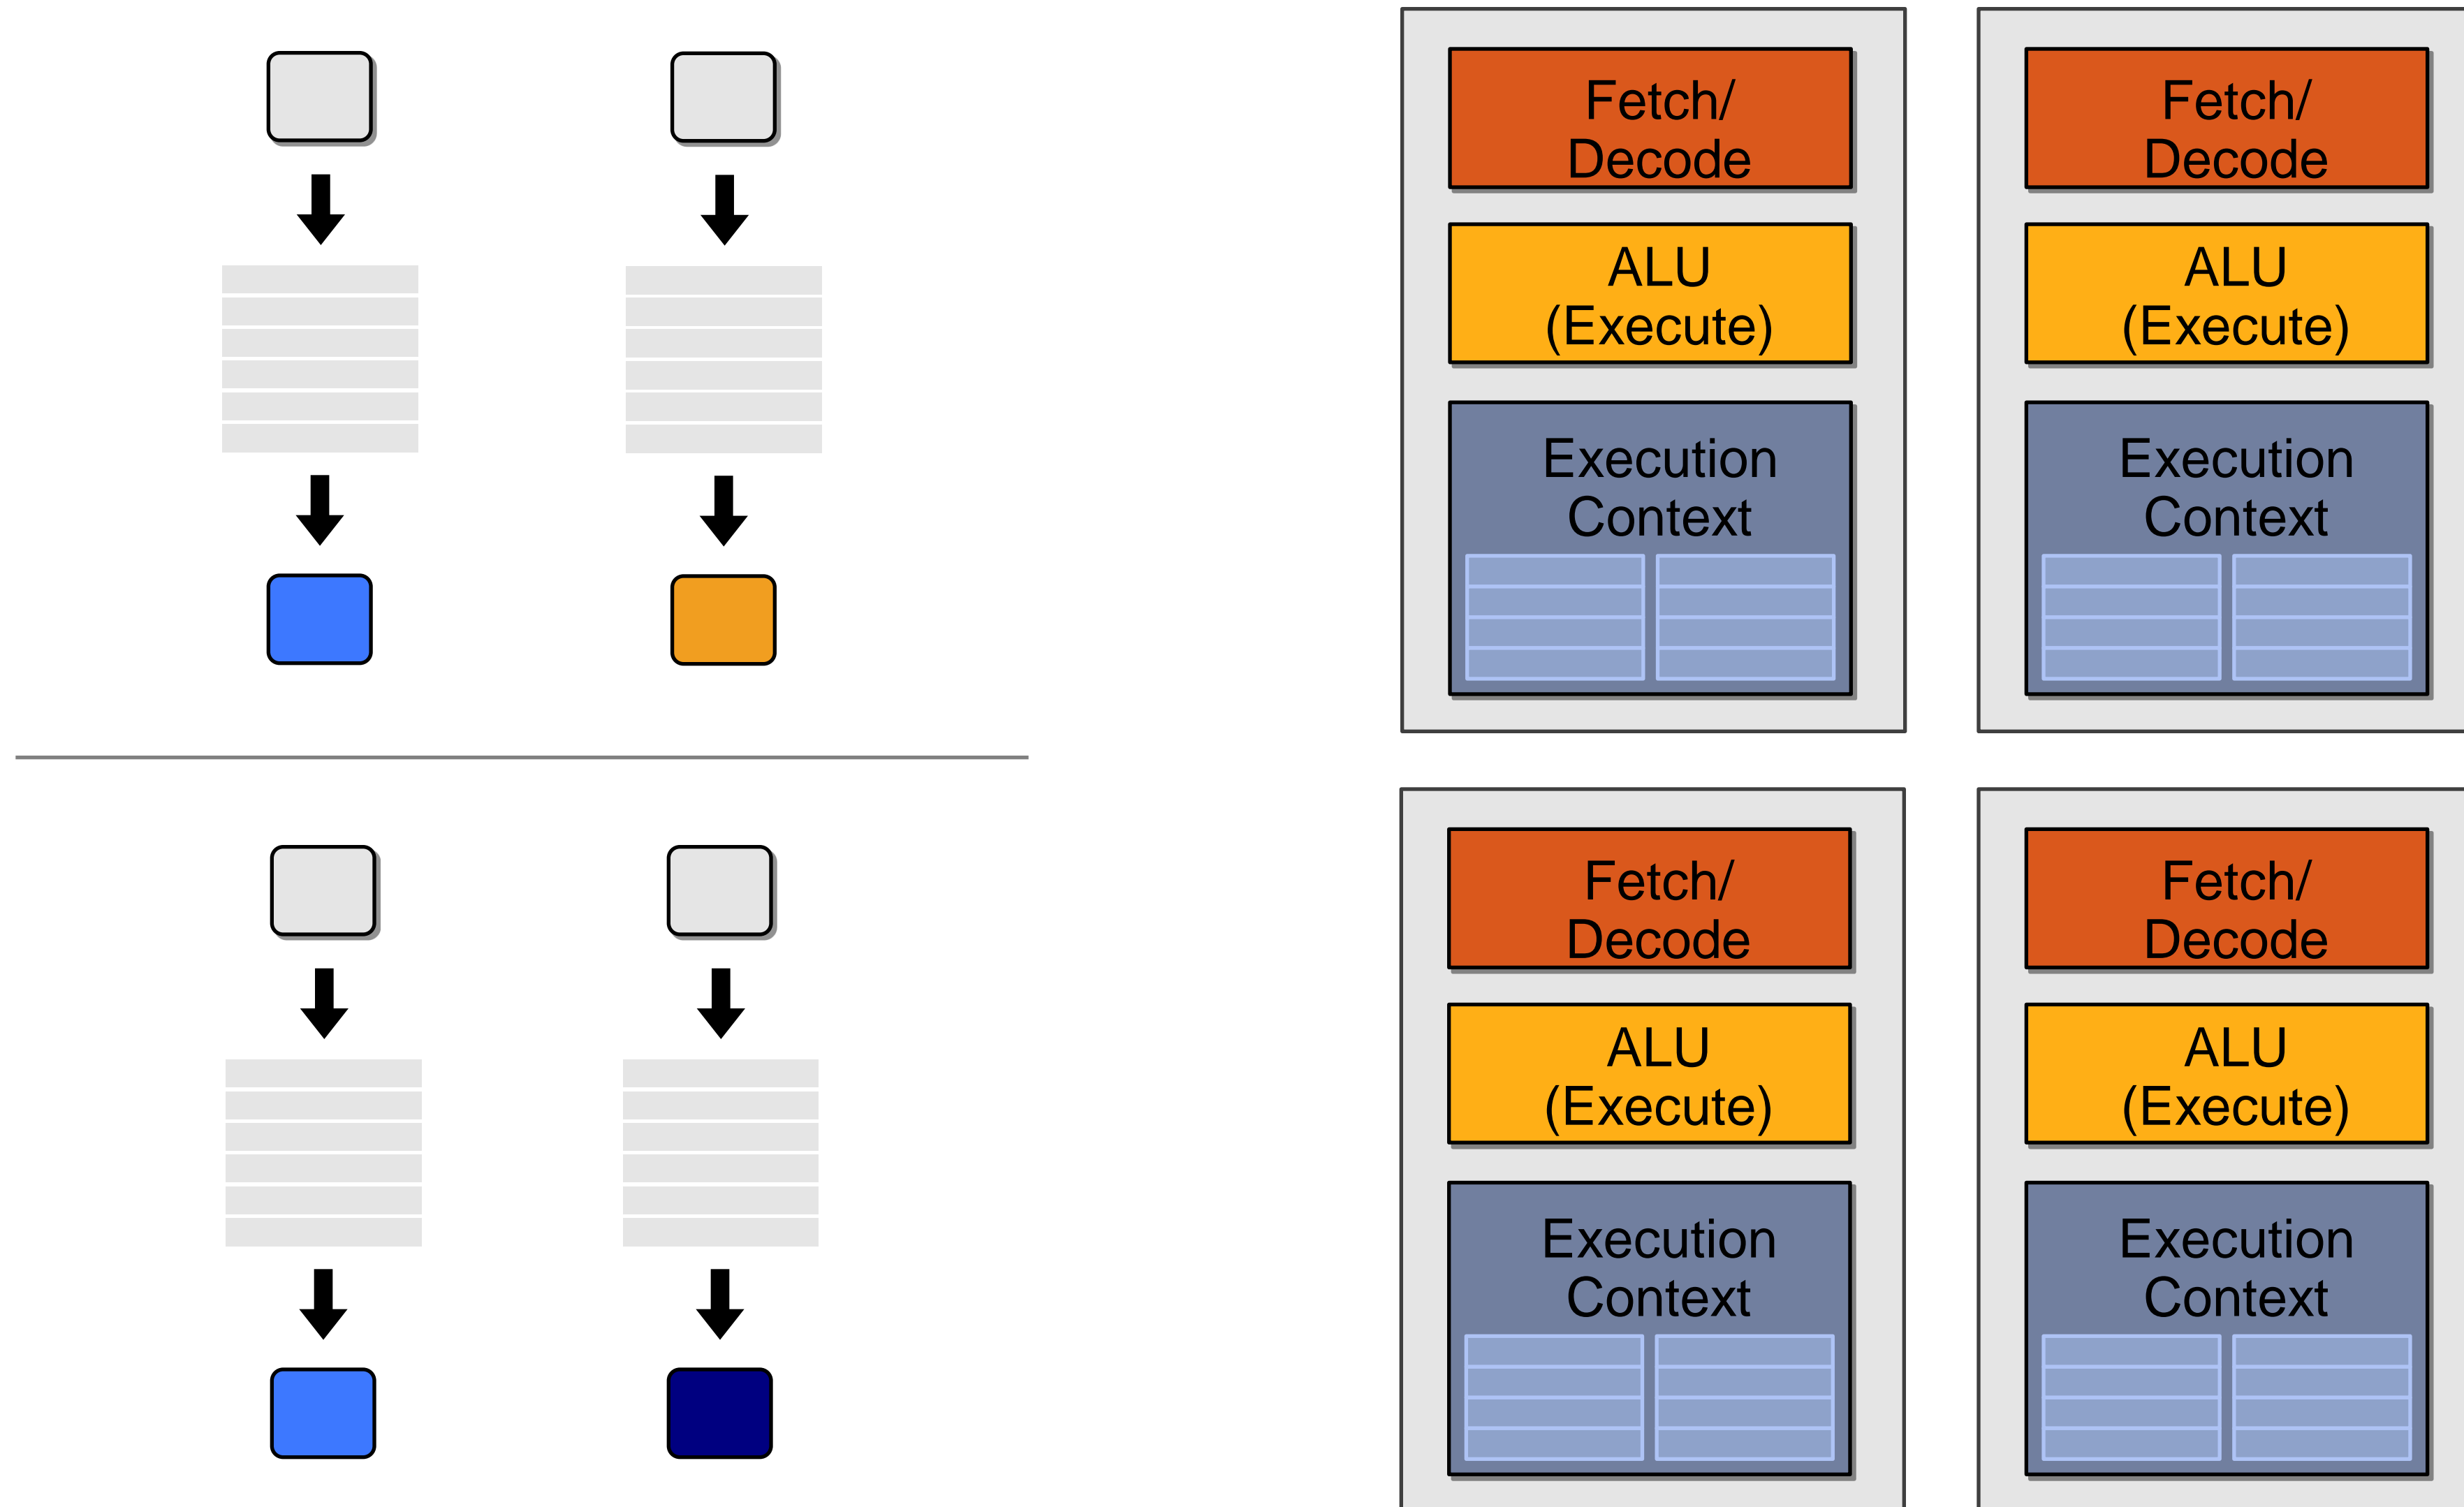






# Two cores (two fragments in parallel)



# Four cores (four fragments in parallel)

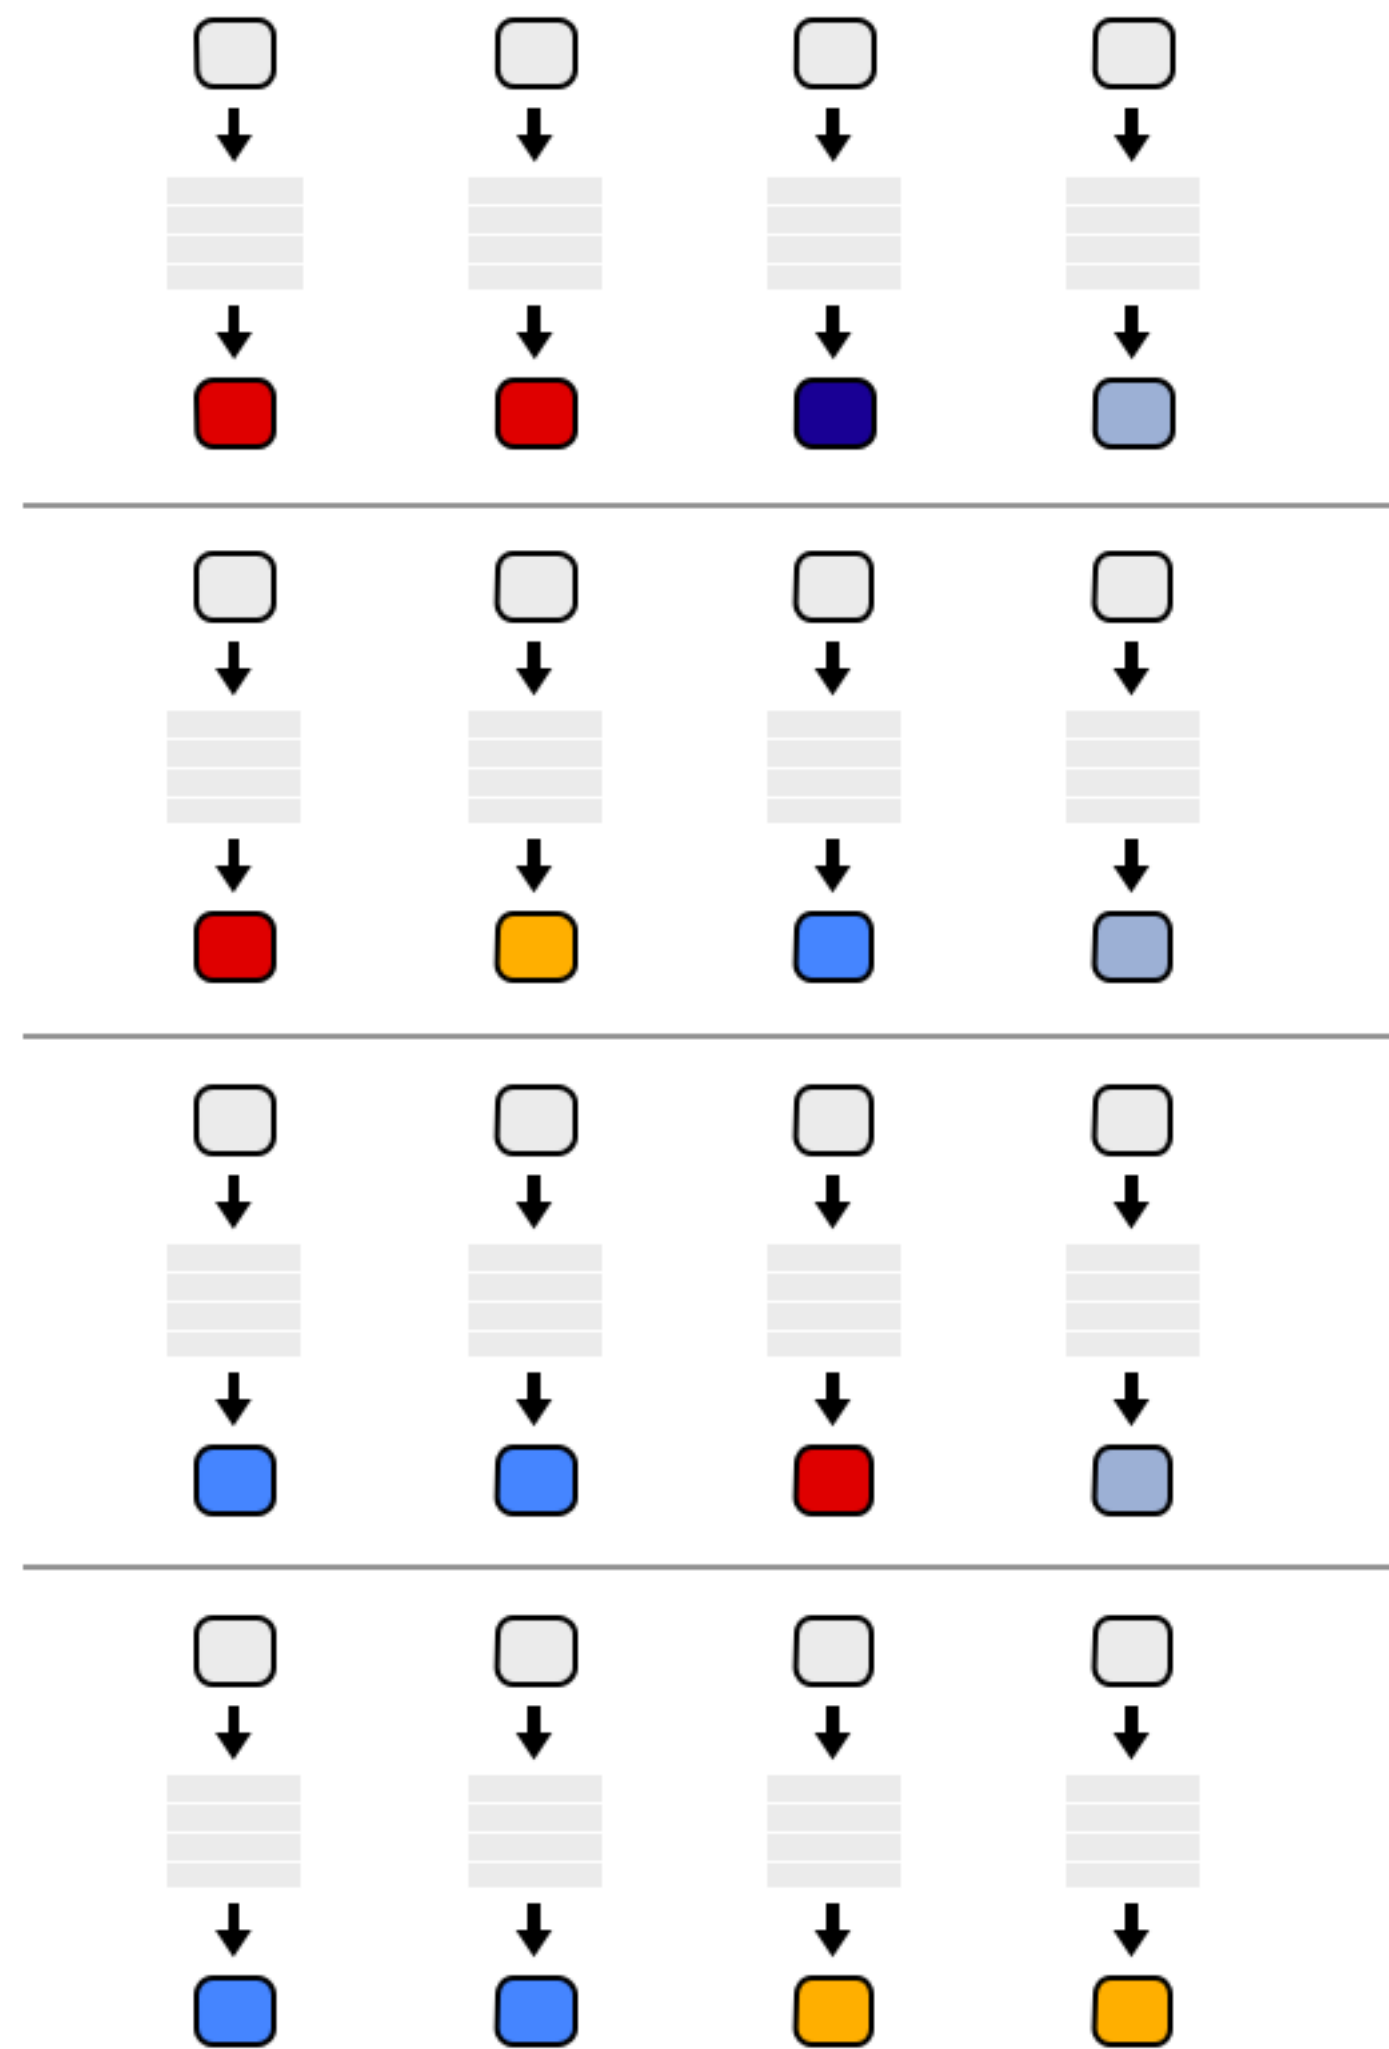


# Sixteen cores (sixteen fragments in parallel)



16 cores = 16 simultaneous instruction streams

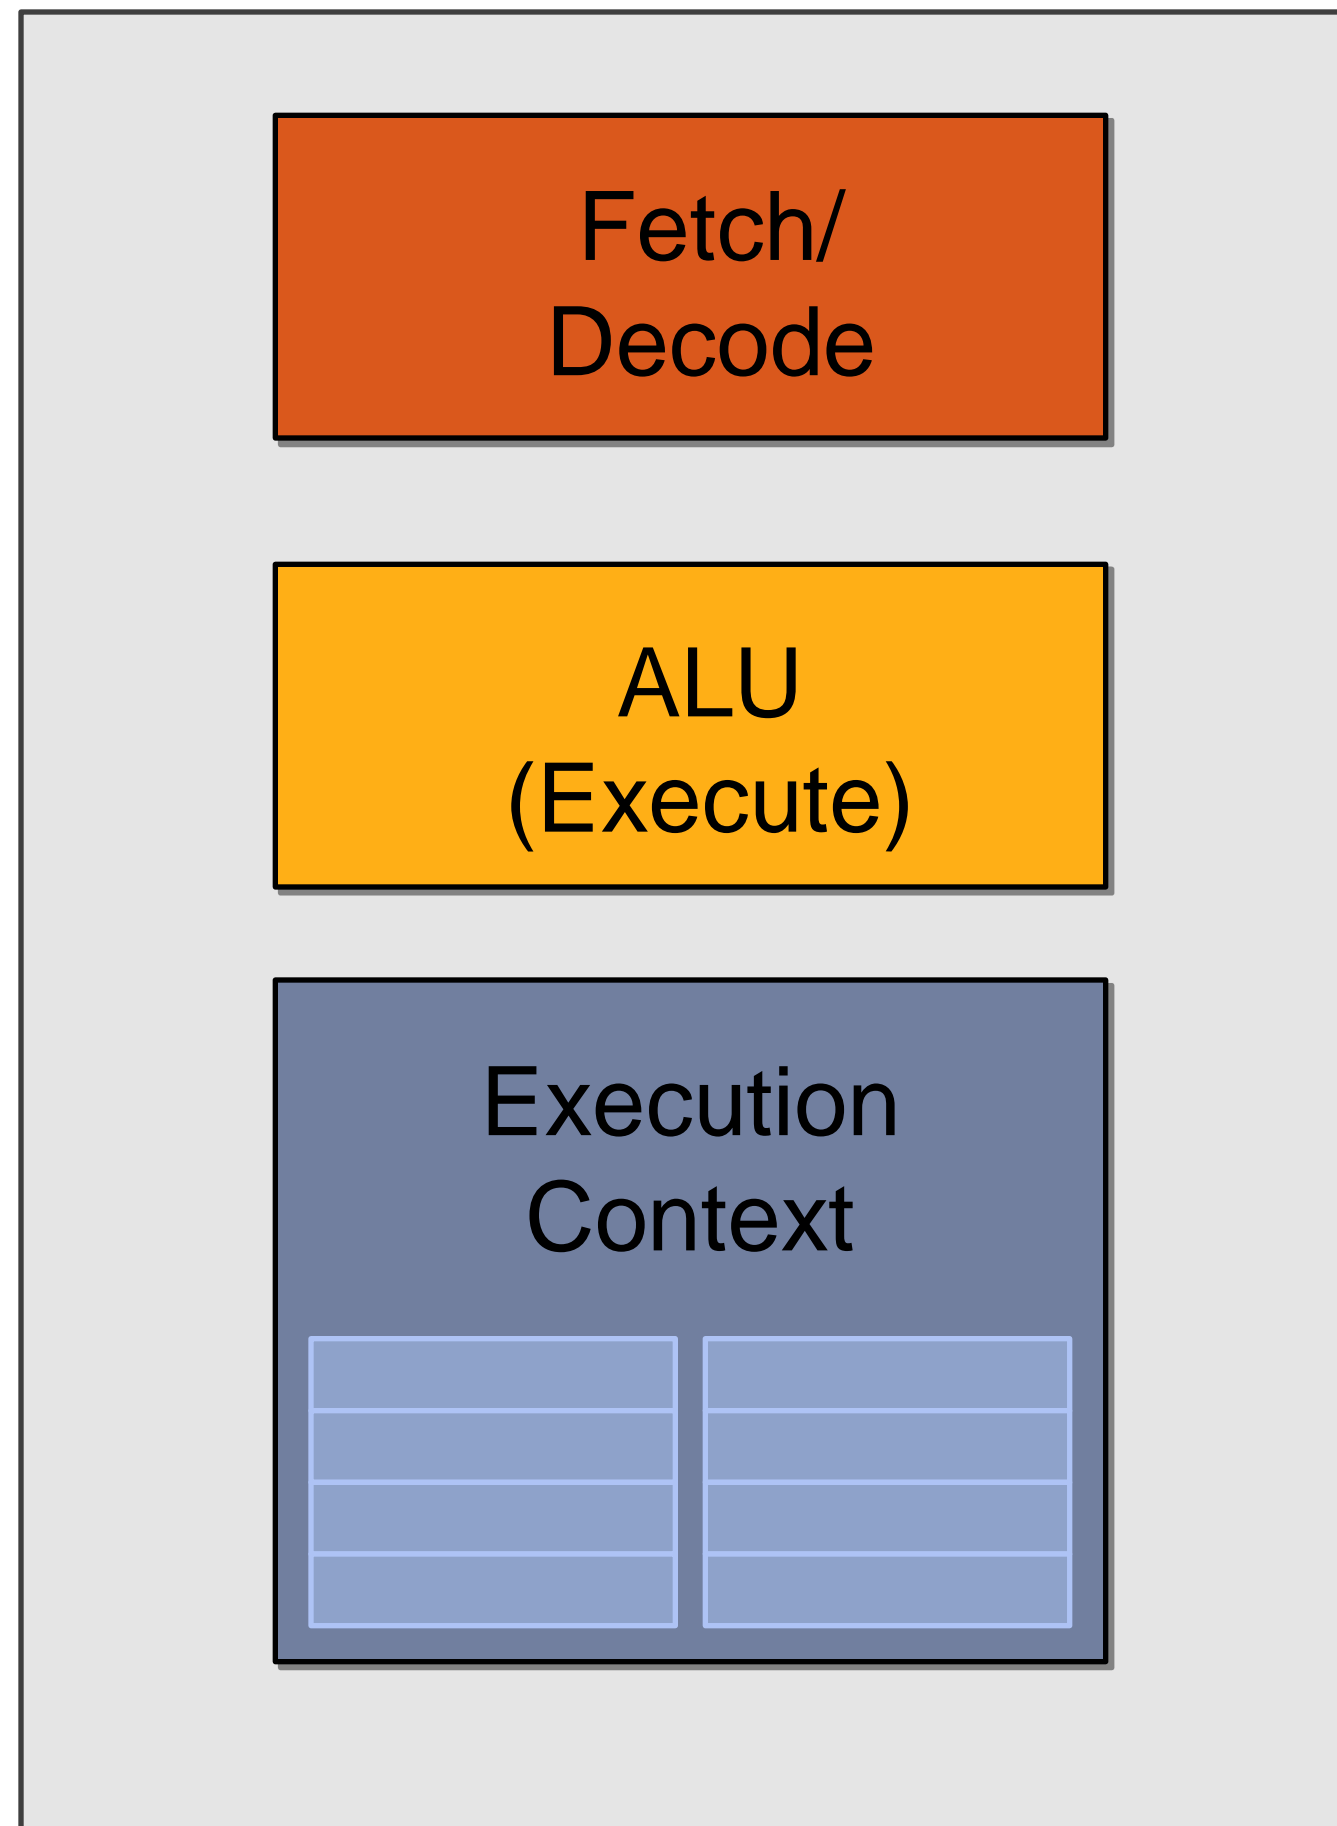
# Instruction stream sharing




But ... many fragments  
should be able to share an  
instruction stream!

```
<diffuseShader>:  
sample r0, v4, t0, s0  
mul  r3, v0, cb0[0]  
madd r3, v1, cb0[1], r3  
madd r3, v2, cb0[2], r3  
clmp r3, r3, 1(0.0), 1(1.0)  
mul  o0, r0, r3  
mul  o1, r1, r3  
mul  o2, r2, r3  
mov  o3, 1(1.0)
```

# Recall: simple processing core



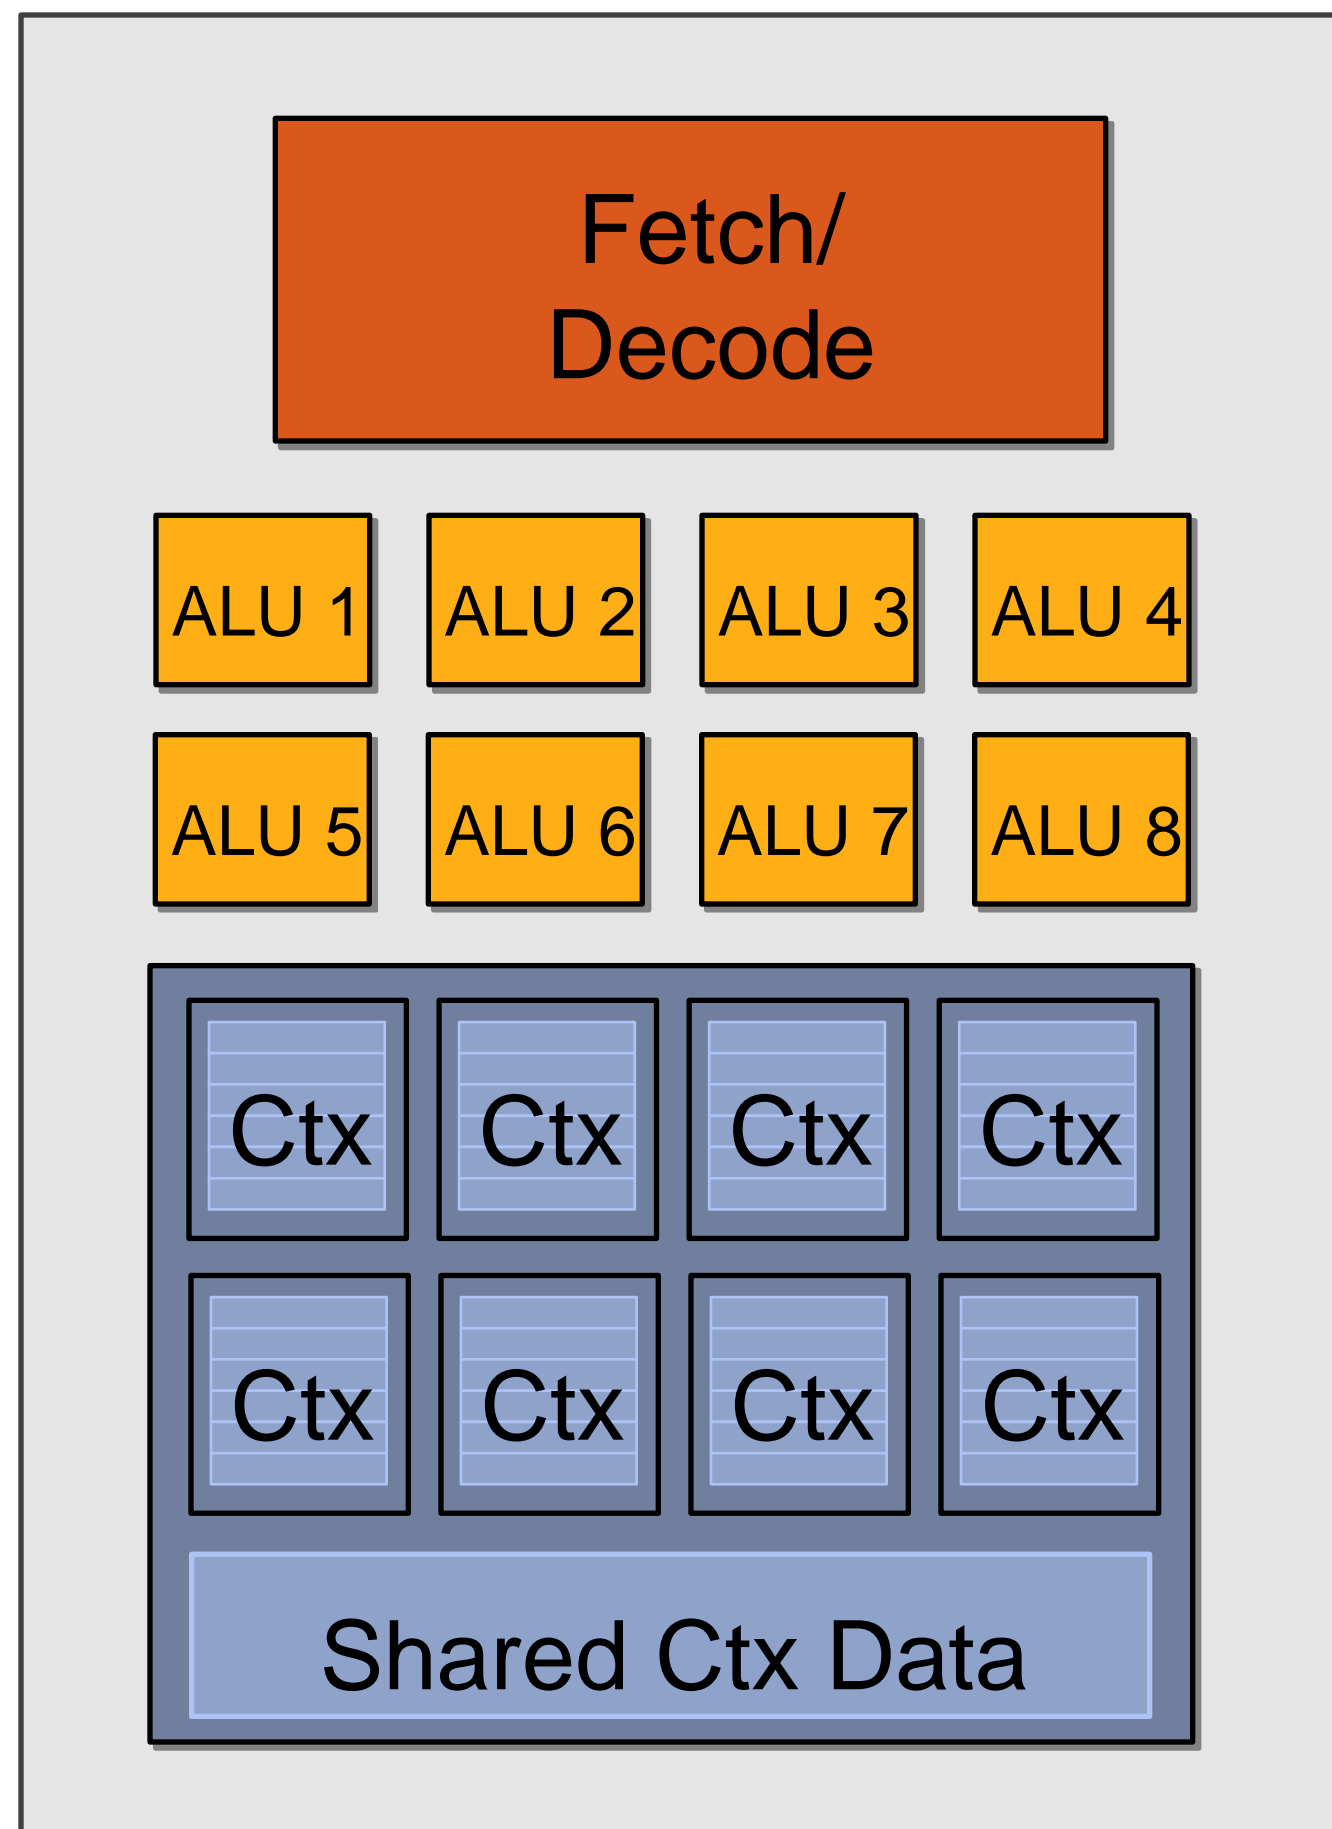
# Add ALUs



Idea #2:  
Amortize cost/complexity of  
managing an instruction  
stream across many ALUs

SIMD processing

# Modifying the shader




```
<diffuseShader>:
sample r0, v4, t0, s0
mul  r3, v0, cb0[0]
madd r3, v1, cb0[1], r3
madd r3, v2, cb0[2], r3
clmp r3, r3, 1(0.0), 1(1.0)
mul  o0, r0, r3
mul  o1, r1, r3
mul  o2, r2, r3
mov  o3, 1(1.0)
```

Original compiled shader:

Processes one fragment using  
scalar ops on scalar registers



# Modifying the shader


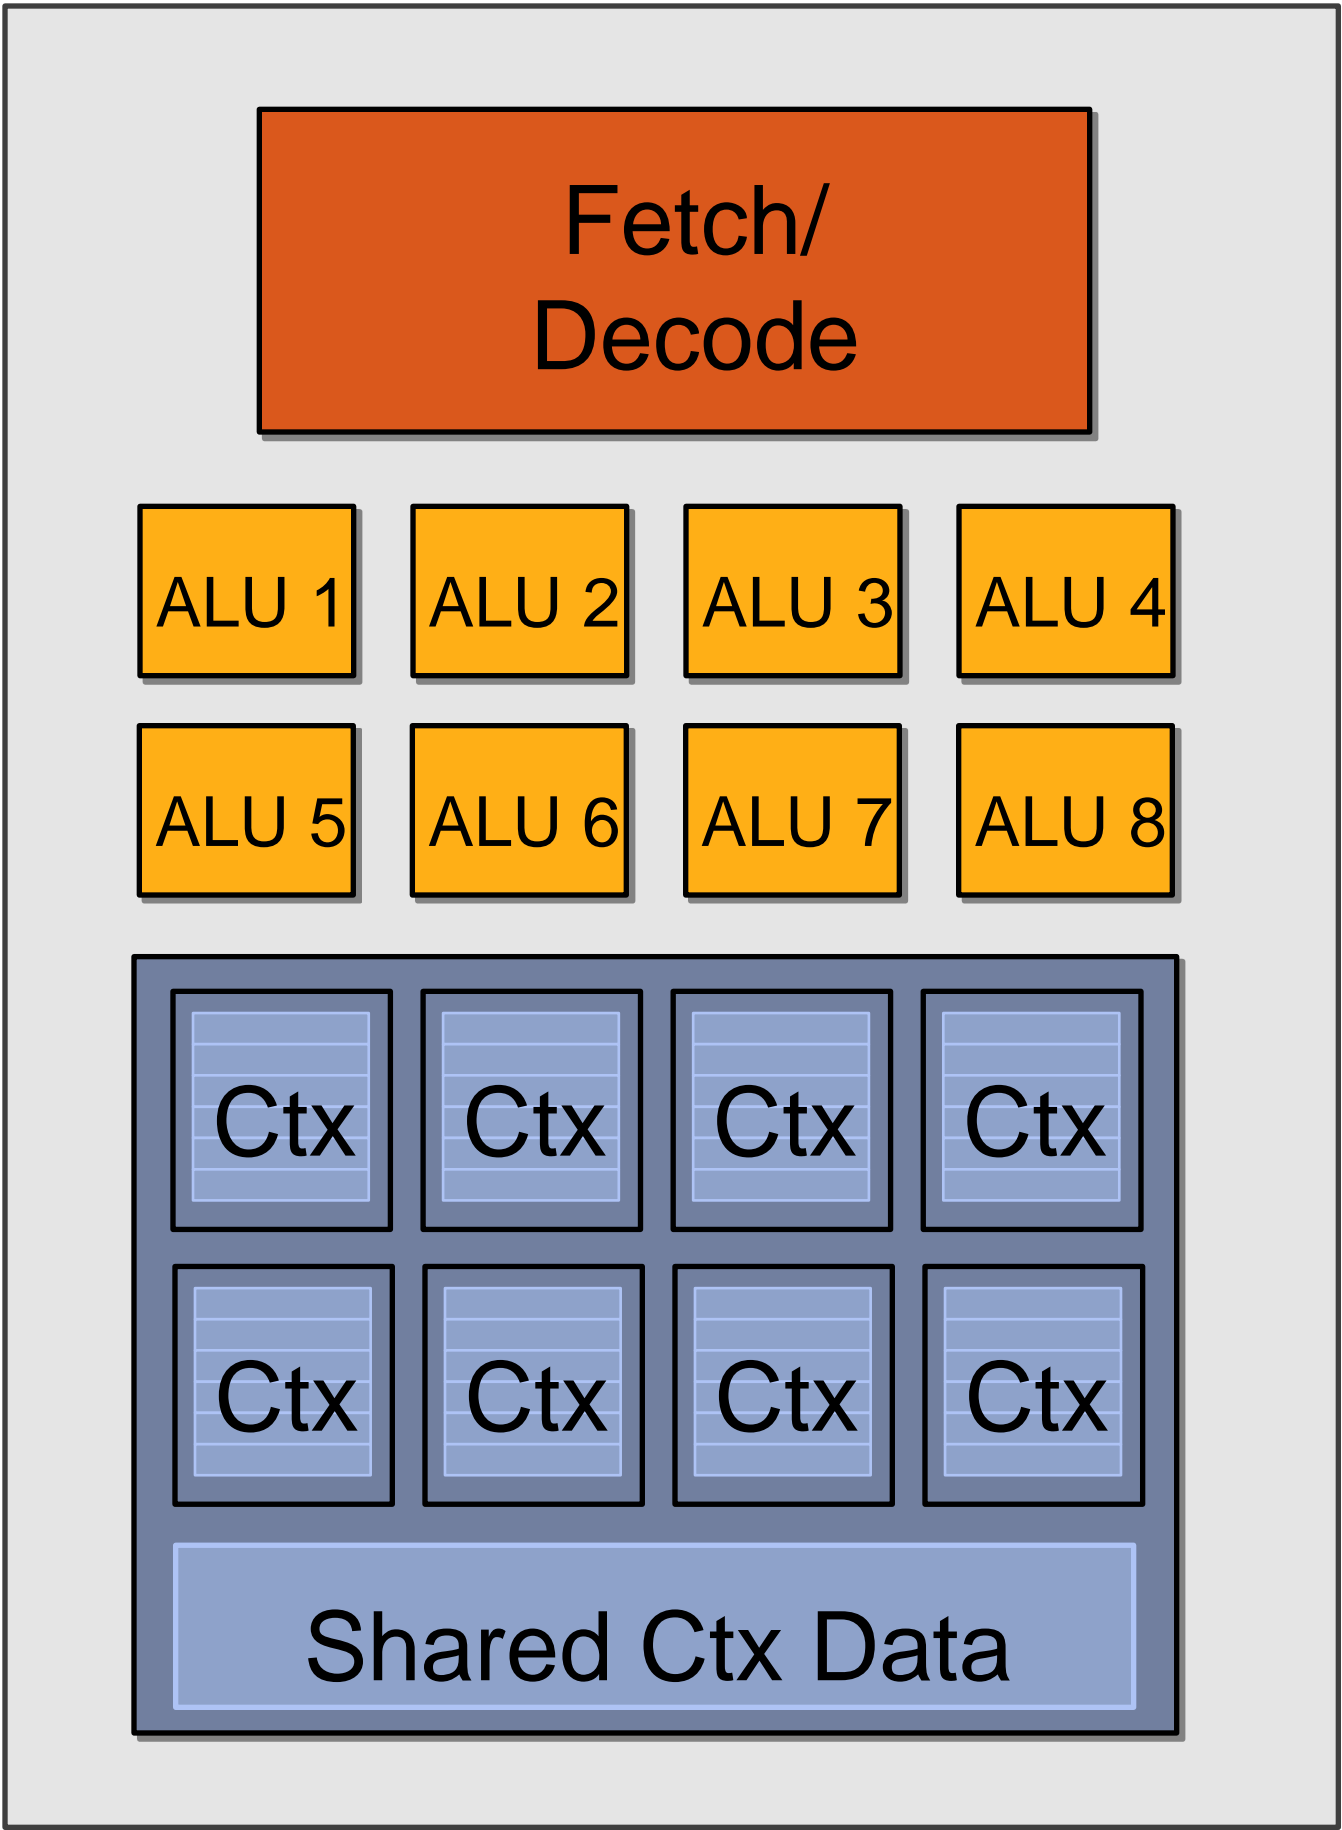


```
<VEC8_diffuseShader>:  
VEC8_sample vec_r0, vec_v4, t0, vec_s0  
VEC8_mul   vec_r3, vec_v0, cb0[0]  
VEC8_madd  vec_r3, vec_v1, cb0[1], vec_r3  
VEC8_madd  vec_r3, vec_v2, cb0[2], vec_r3  
VEC8_clmp  vec_r3, vec_r3, 1(0.0), 1(1.0)  
VEC8_mul   vec_o0, vec_r0, vec_r3  
VEC8_mul   vec_o1, vec_r1, vec_r3  
VEC8_mul   vec_o2, vec_r2, vec_r3  
VEC8_mov   o3, 1(1.0)
```


New compiled shader:

Processes eight fragments using  
vector ops on vector registers


# Modifying the shader




```
<VEC8_diffuseShader>:
VEC8_sample vec_r0, vec_v4, t0, vec_s0
VEC8_mul  vec_r3, vec_v0, cb0[0]
VEC8_madd vec_r3, vec_v1, cb0[1], vec_r3
VEC8_madd vec_r3, vec_v2, cb0[2], vec_r3
VEC8_clmp vec_r3, vec_r3, 1(0.0), 1(1.0)
VEC8_mul  vec_o0, vec_r0, vec_r3
VEC8_mul  vec_o1, vec_r1, vec_r3
VEC8_mul  vec_o2, vec_r2, vec_r3
VEC8_mov  o3, 1(1.0)
```



# 128 fragments in parallel




16 cores = 128 ALUs




, 16 simultaneous instruction streams

128 [ vertices/fragments  
primitives  
OpenCL work items ] in parallel



# But what about branches?




<unconditional  
shader code>

```
if (x > 0) {  
    y = pow(x, exp);  
    y *= Ks;  
    refl = y + Ka;  
} else {  
    x = 0;  
    refl = Ka;  
}
```

<resume unconditional  
shader code>

# But what about branches?




<unconditional  
shader code>

```
if (x > 0) {  
    y = pow(x, exp);  
    y *= Ks;  
    refl = y + Ka;  
} else {  
    x = 0;  
    refl = Ka;  
}
```

<resume unconditional  
shader code>


# But what about branches?



```
<unconditional  
shader code>  
  
if (x > 0) {  
    y = pow(x, exp);  
    y *= Ks;  
    refl = y + Ka;  
} else {  
    x = 0;  
    refl = Ka;  
}  
  
<resume unconditional  
shader code>
```



# But what about branches?



```
<unconditional  
shader code>  
  
if (x > 0) {  
    y = pow(x, exp);  
    y *= Ks;  
    refl = y + Ka;  
} else {  
    x = 0;  
    refl = Ka;  
}  
  
<resume unconditional  
shader code>
```

# Clarification

## SIMD processing does not imply SIMD instructions

- Option 1: explicit vector instructions
  - x86 SSE, AVX, Intel Larrabee
- Option 2: scalar instructions, implicit HW vectorization
  - HW determines instruction stream sharing across ALUs (amount of sharing hidden from software)
  - NVIDIA GeForce (“SIMT” warps), ATI Radeon architectures (“wavefronts”)



In practice: 16 to 64 fragments share an instruction stream.

# Stalls!

Stalls occur when a core cannot run the next instruction because of a dependency on a previous operation.

Texture access latency = 100's to 1000's of cycles

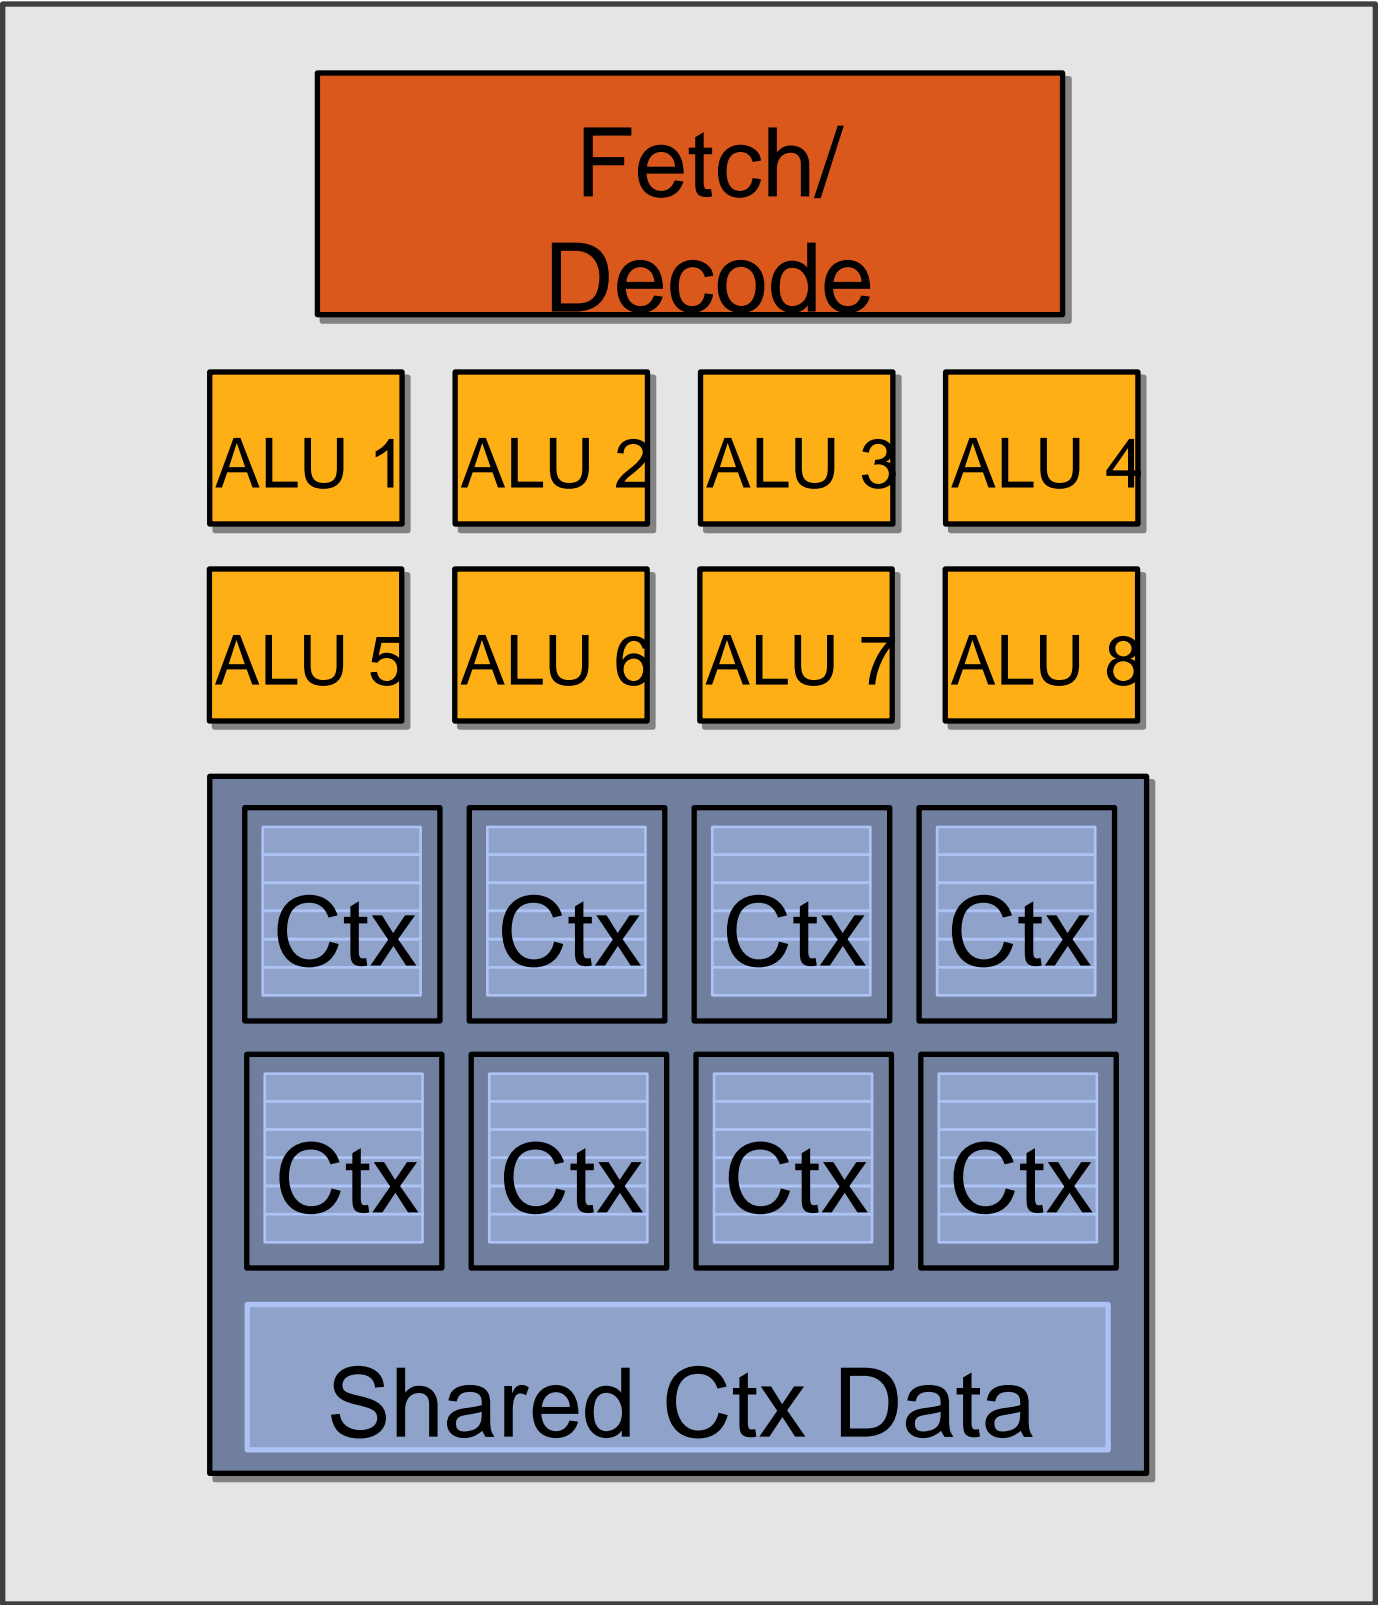

We've removed the fancy caches and logic that helps avoid stalls.

But we have **LOTS** of independent fragments.

**Idea #3:**

Interleave processing of many fragments on a single core to avoid stalls caused by high latency operations.

# Hiding shader stalls



# Hiding shader stalls


Time (clocks)

Frag 1 ... 8

Frag 9 ... 16

Frag 17 ... 24


Frag 25 ... 32



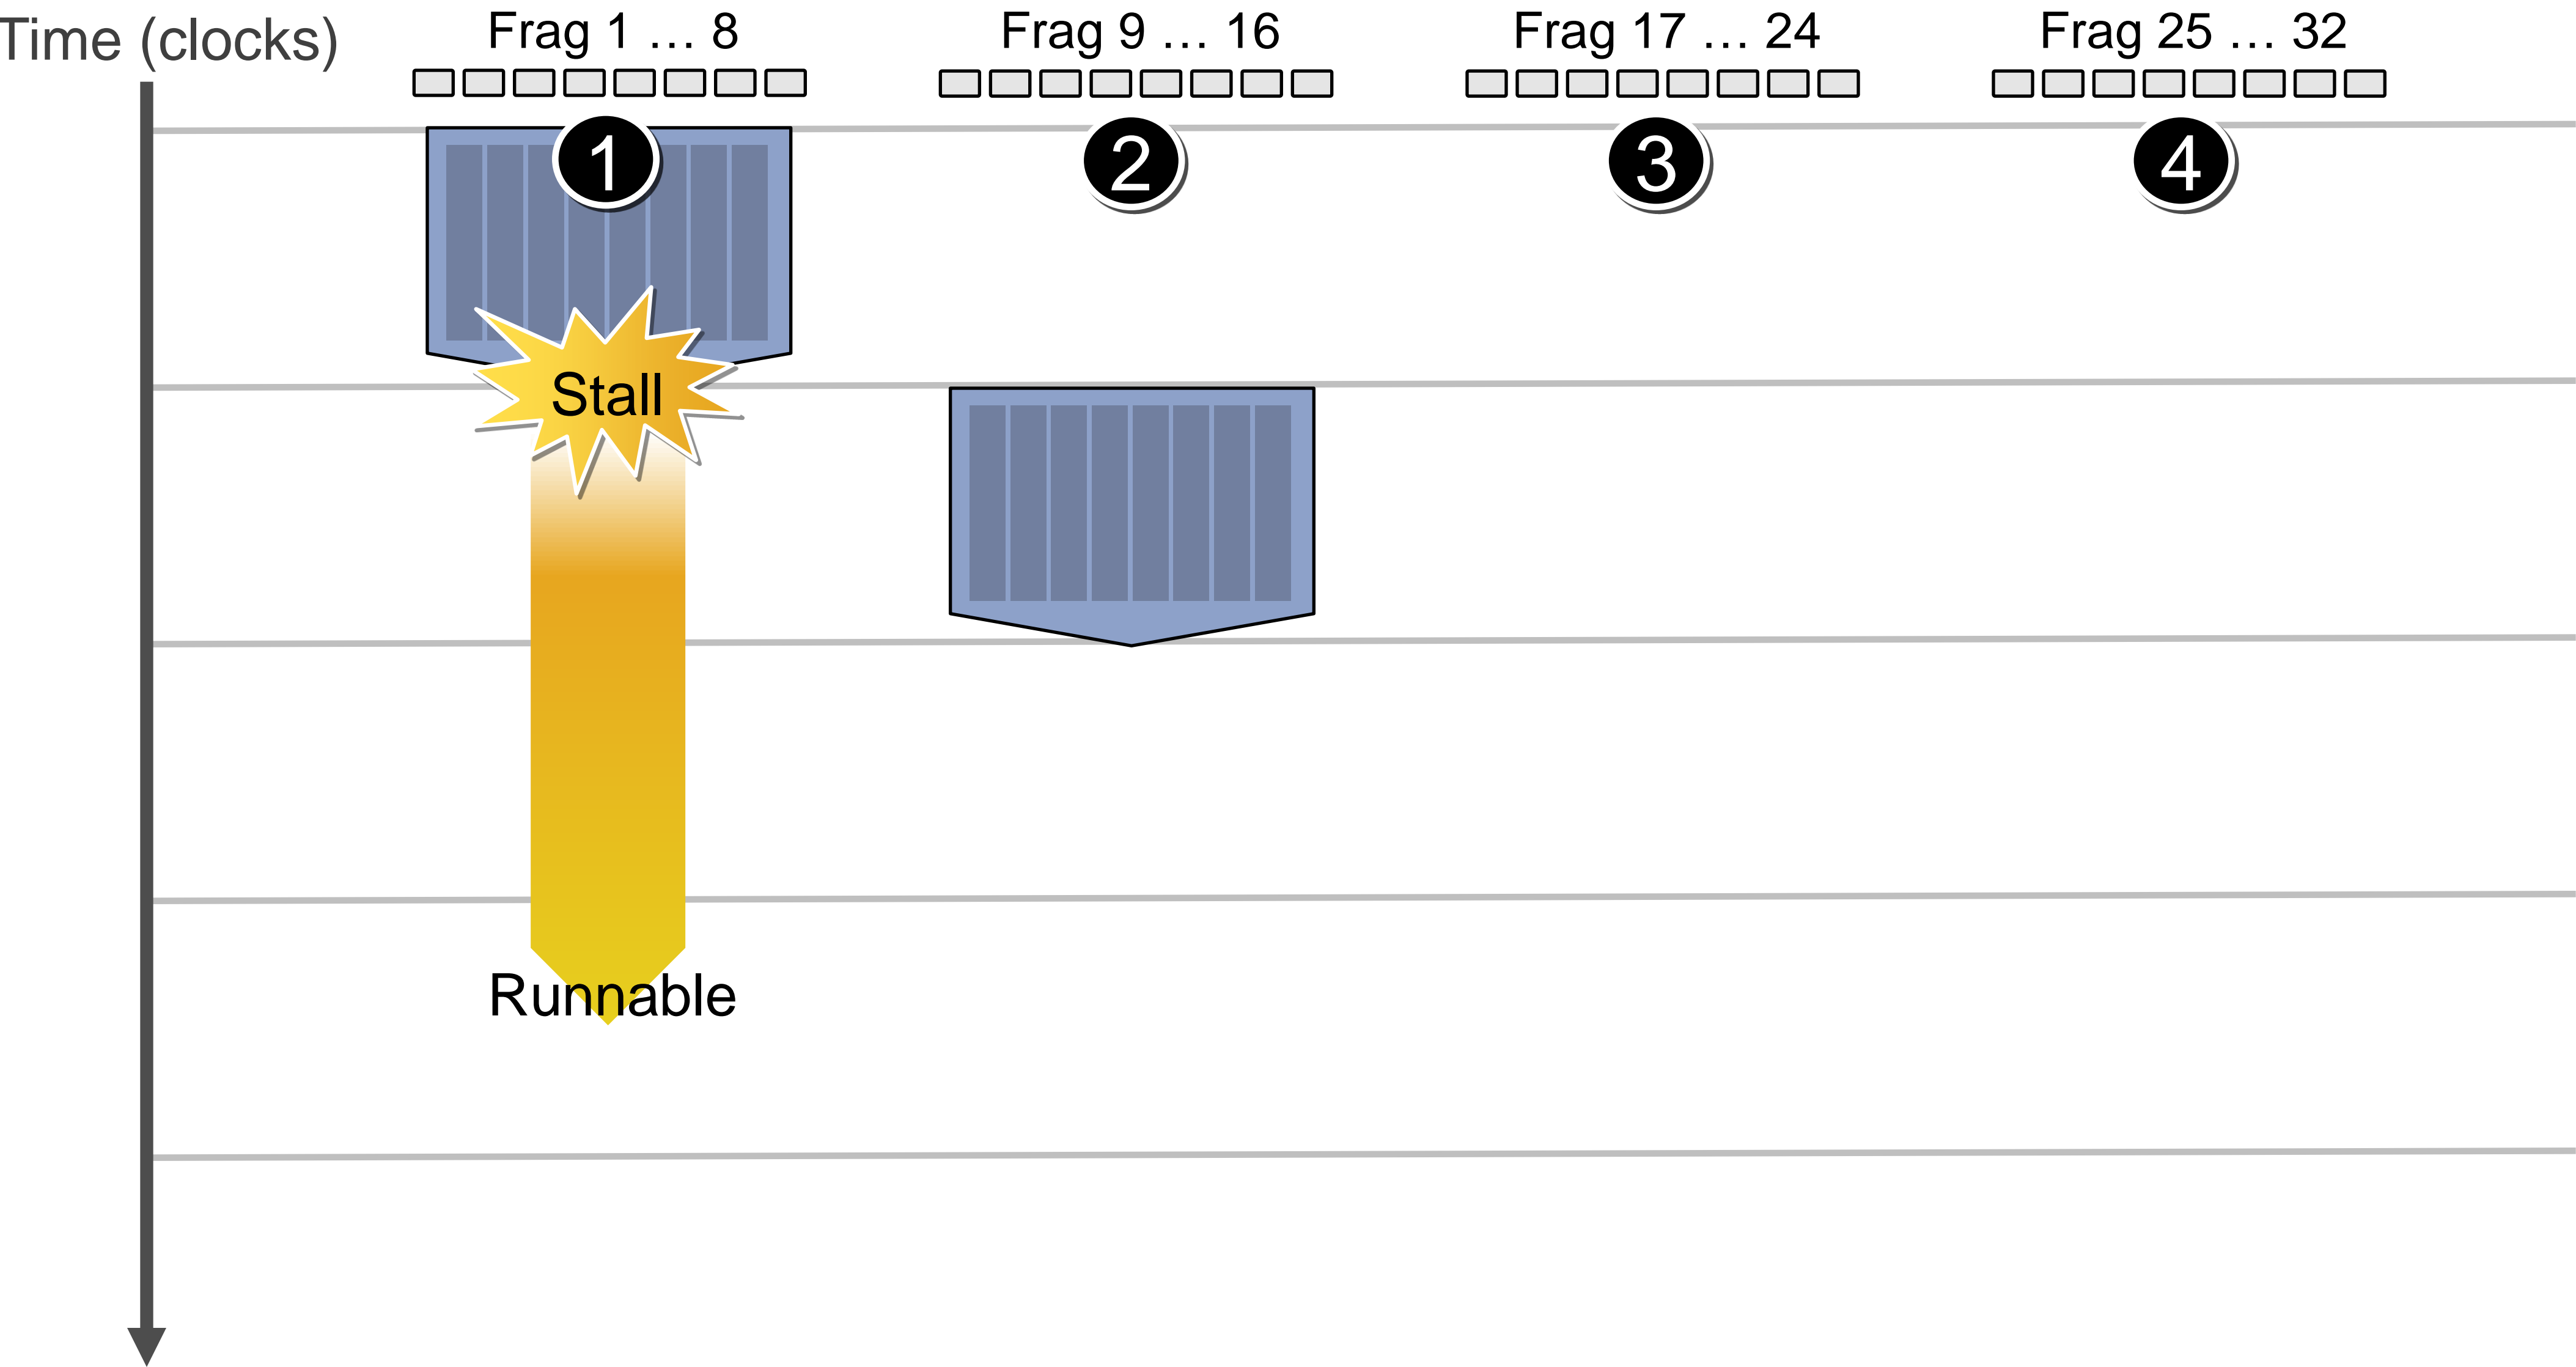
2

3

4

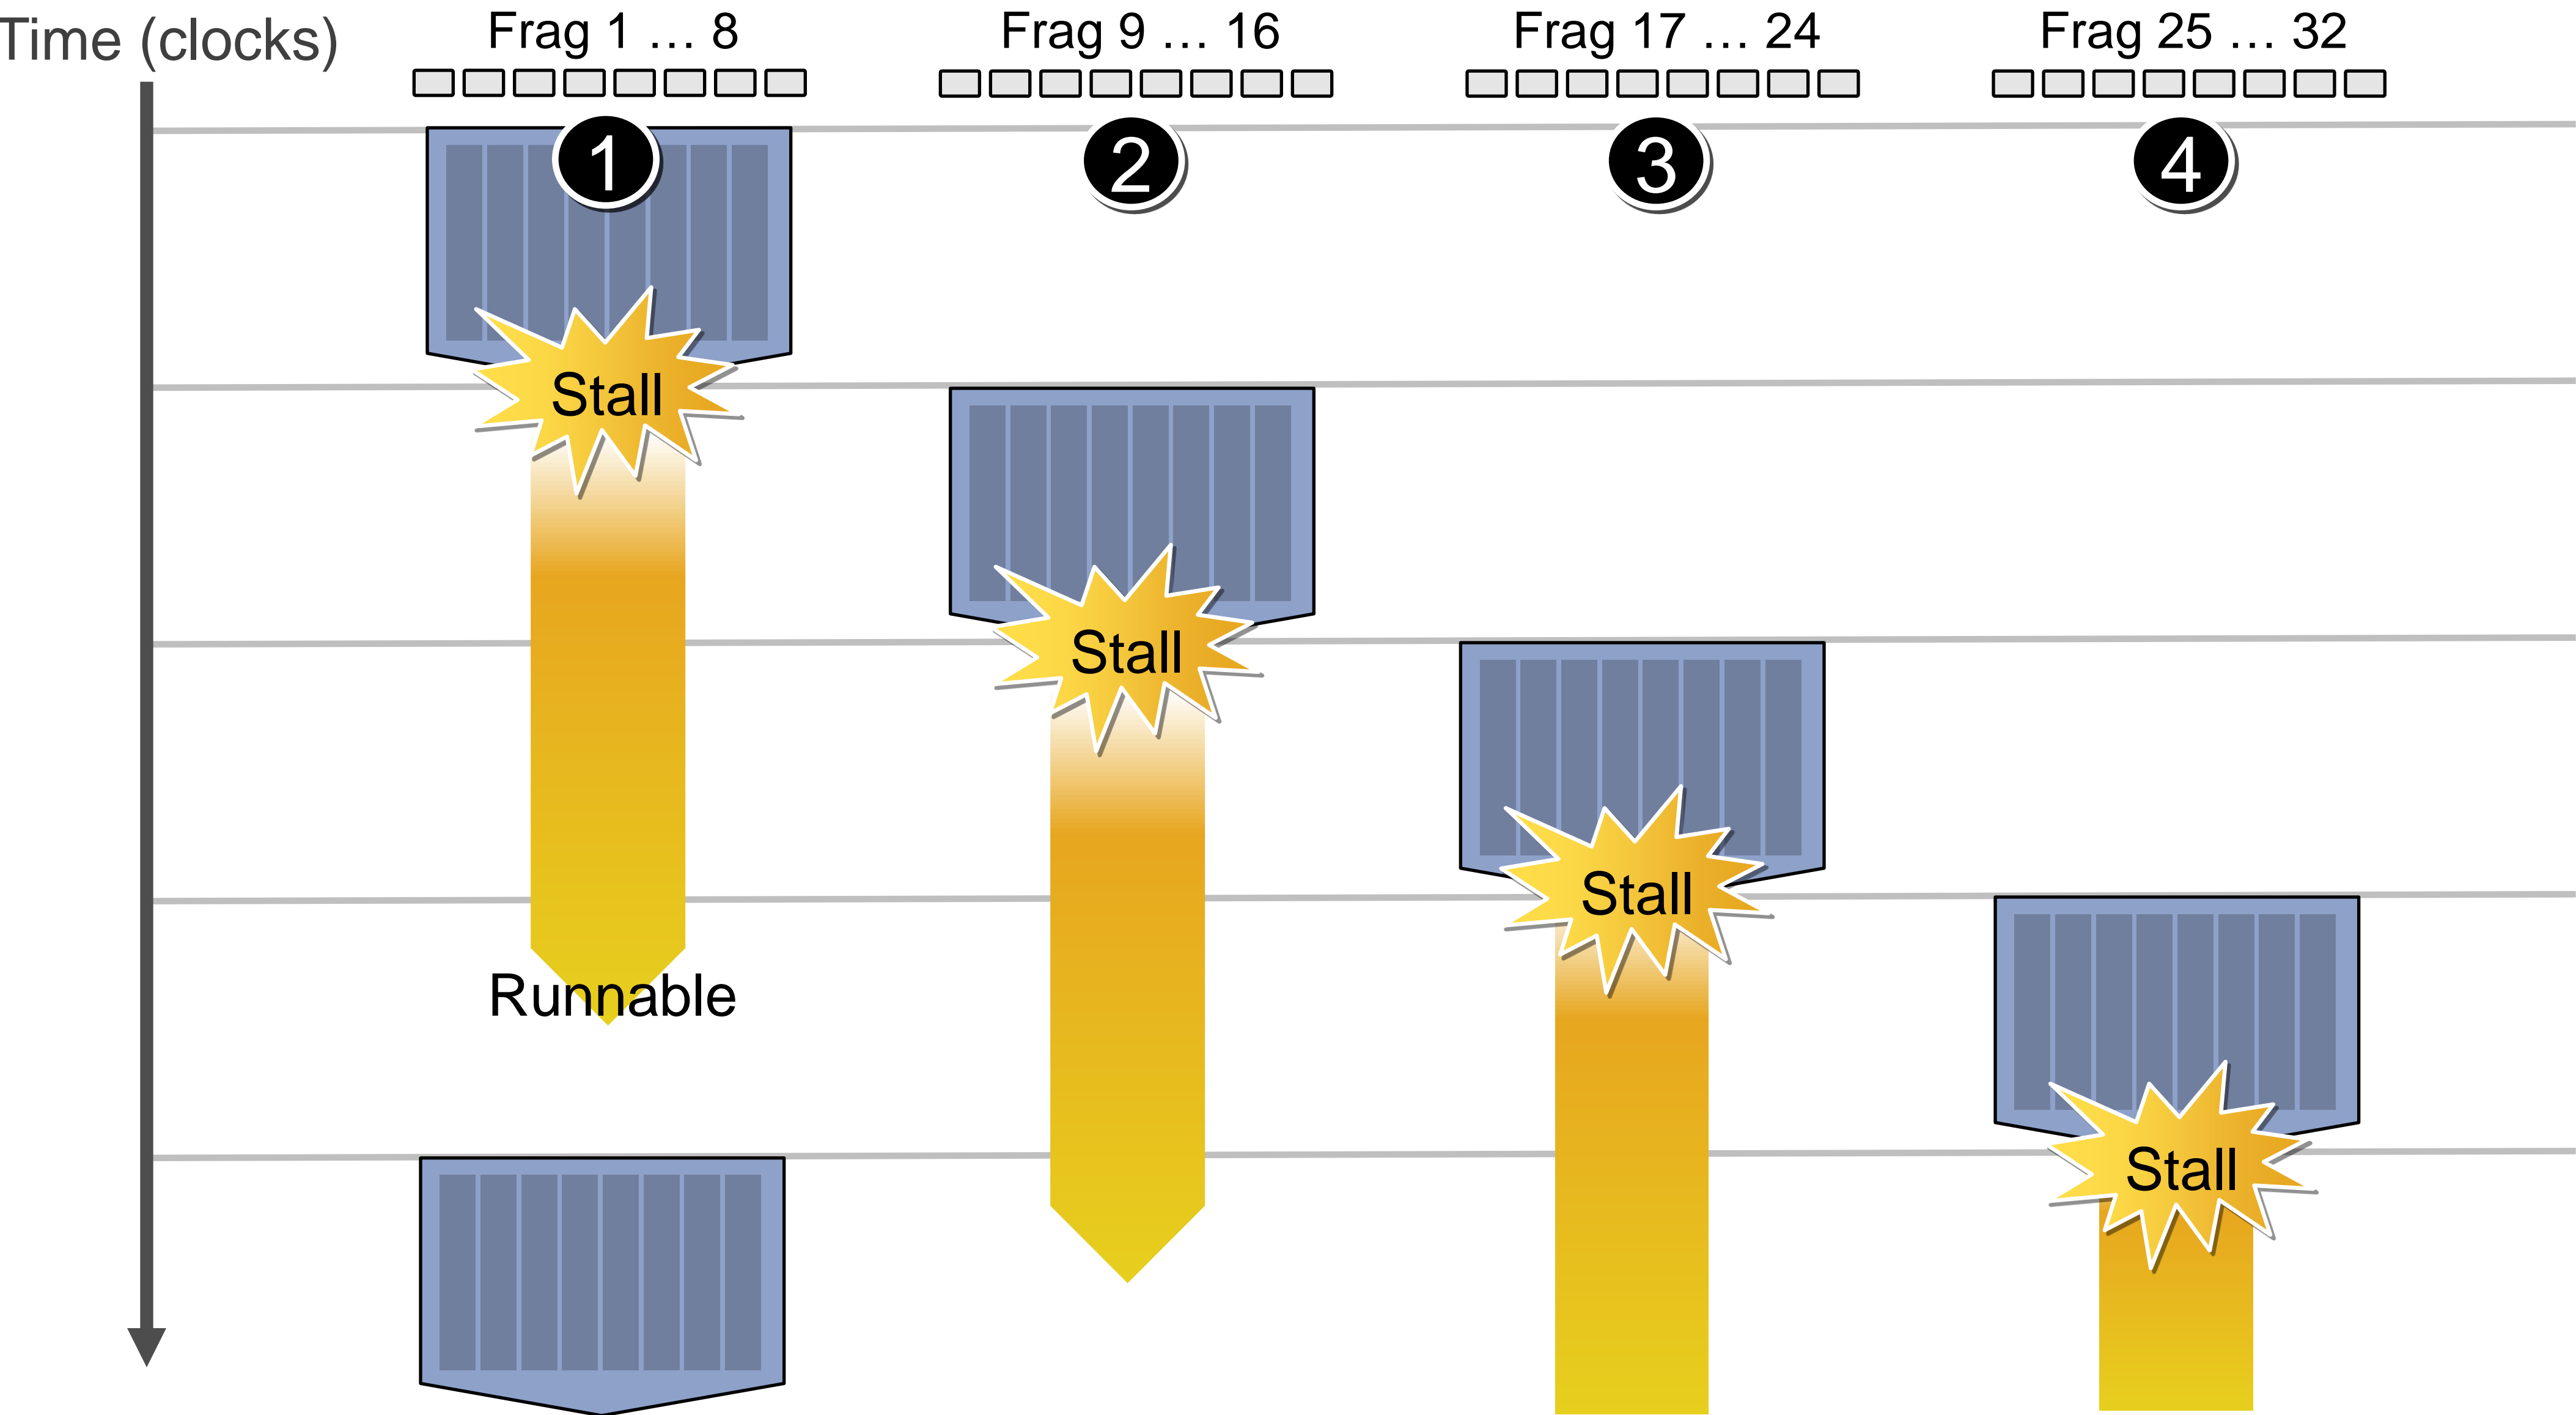


# Hiding shader stalls

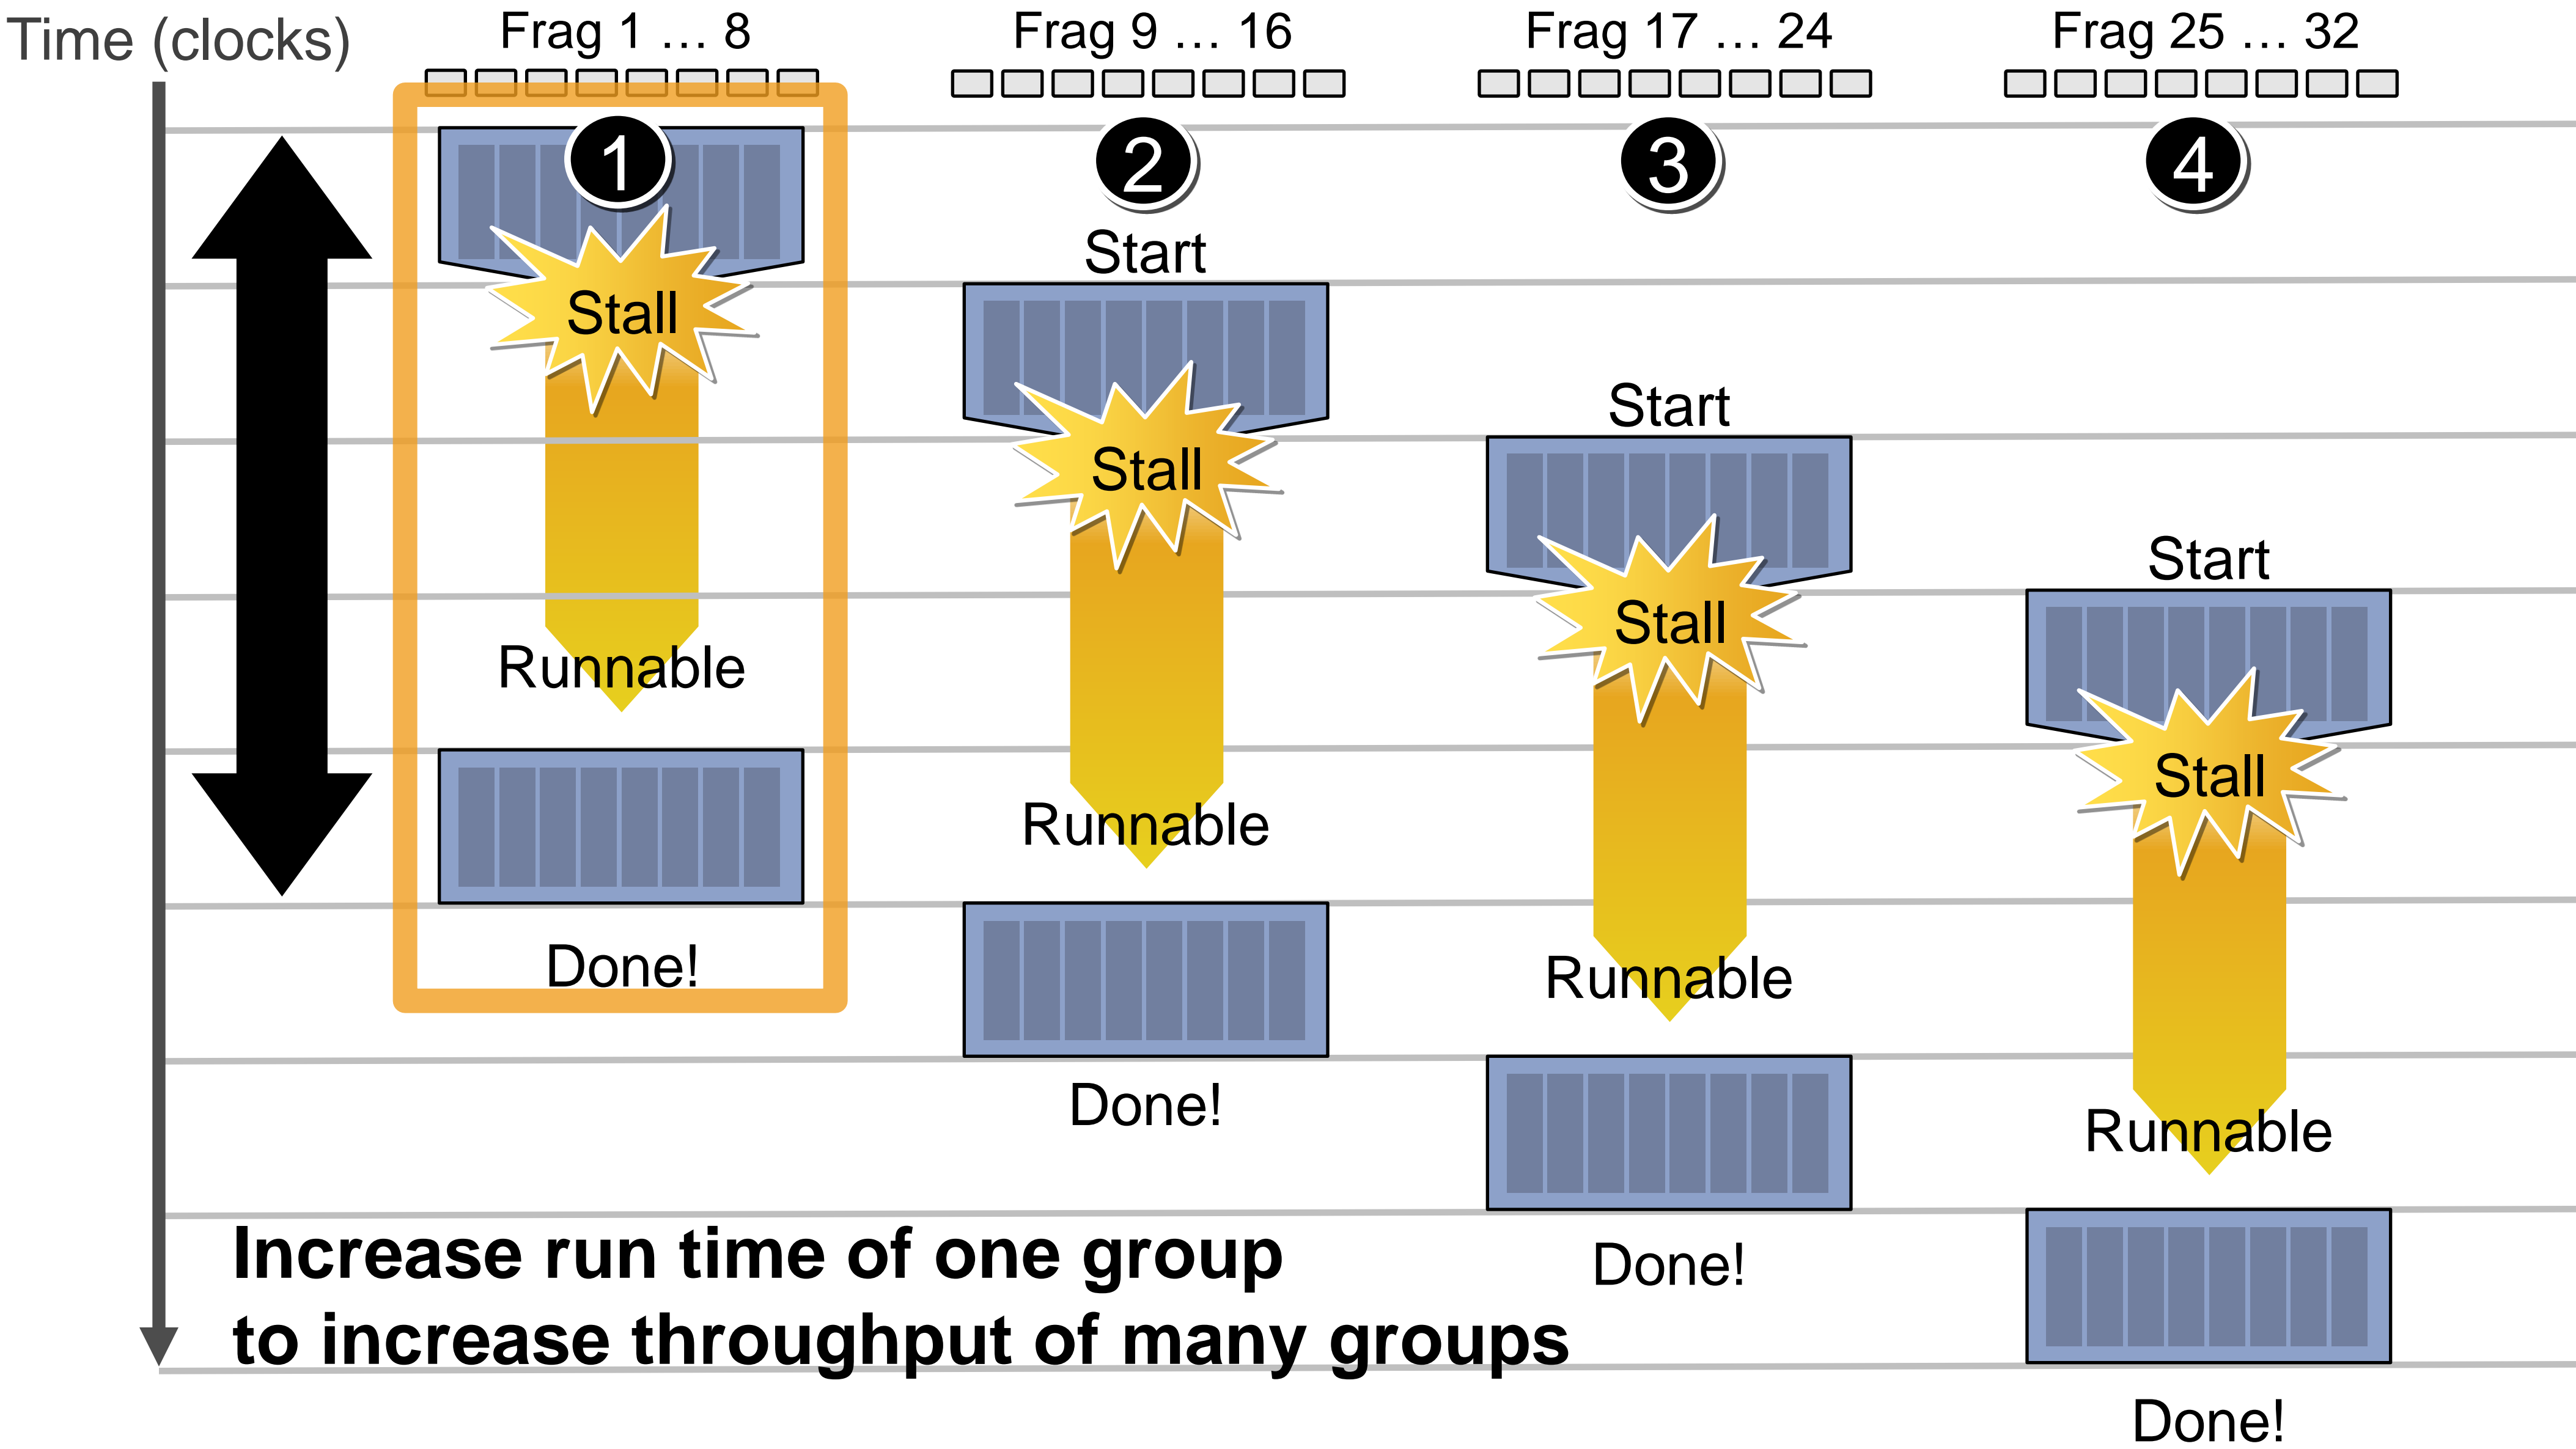





# Hiding shader stalls




# Throughput!




# Storing contexts



# Eighteen small contexts (maximal latency hiding)




# Twelve medium contexts



# Four large contexts

(low latency hiding ability)



# Clarification

Interleaving between contexts can be managed by hardware or software (or both!)

- NVIDIA / AMD Radeon GPUs
  - HW schedules / manages all contexts (lots of them)
  - Special on-chip storage holds fragment state
- Intel Larrabee
  - HW manages four x86 (big) contexts at fine granularity
  - SW scheduling interleaves many groups of fragments on each HW context
  - L1-L2 cache holds fragment state (as determined by SW)

# Example chip

16 cores

8 mul-add ALUs per core  
(128 total)

16 simultaneous  
instruction streams

64 concurrent (but interleaved)  
instruction streams

512 concurrent fragments

= 256 GFLOPs (@ 1GHz)





# Summary: three key ideas

1. Use many “slimmed down cores” to run in parallel
2. Pack cores full of ALUs (by sharing instruction stream across groups of fragments)
  - Option 1: Explicit SIMD vector instructions
  - Option 2: Implicit sharing managed by hardware
3. Avoid latency stalls by interleaving execution of many groups of fragments
  - When one group stalls, work on another group

Part 2:

Putting the three ideas into practice:  
A closer look at real GPUs

NVIDIA GeForce GTX 580

AMD Radeon HD 6970

# Disclaimer


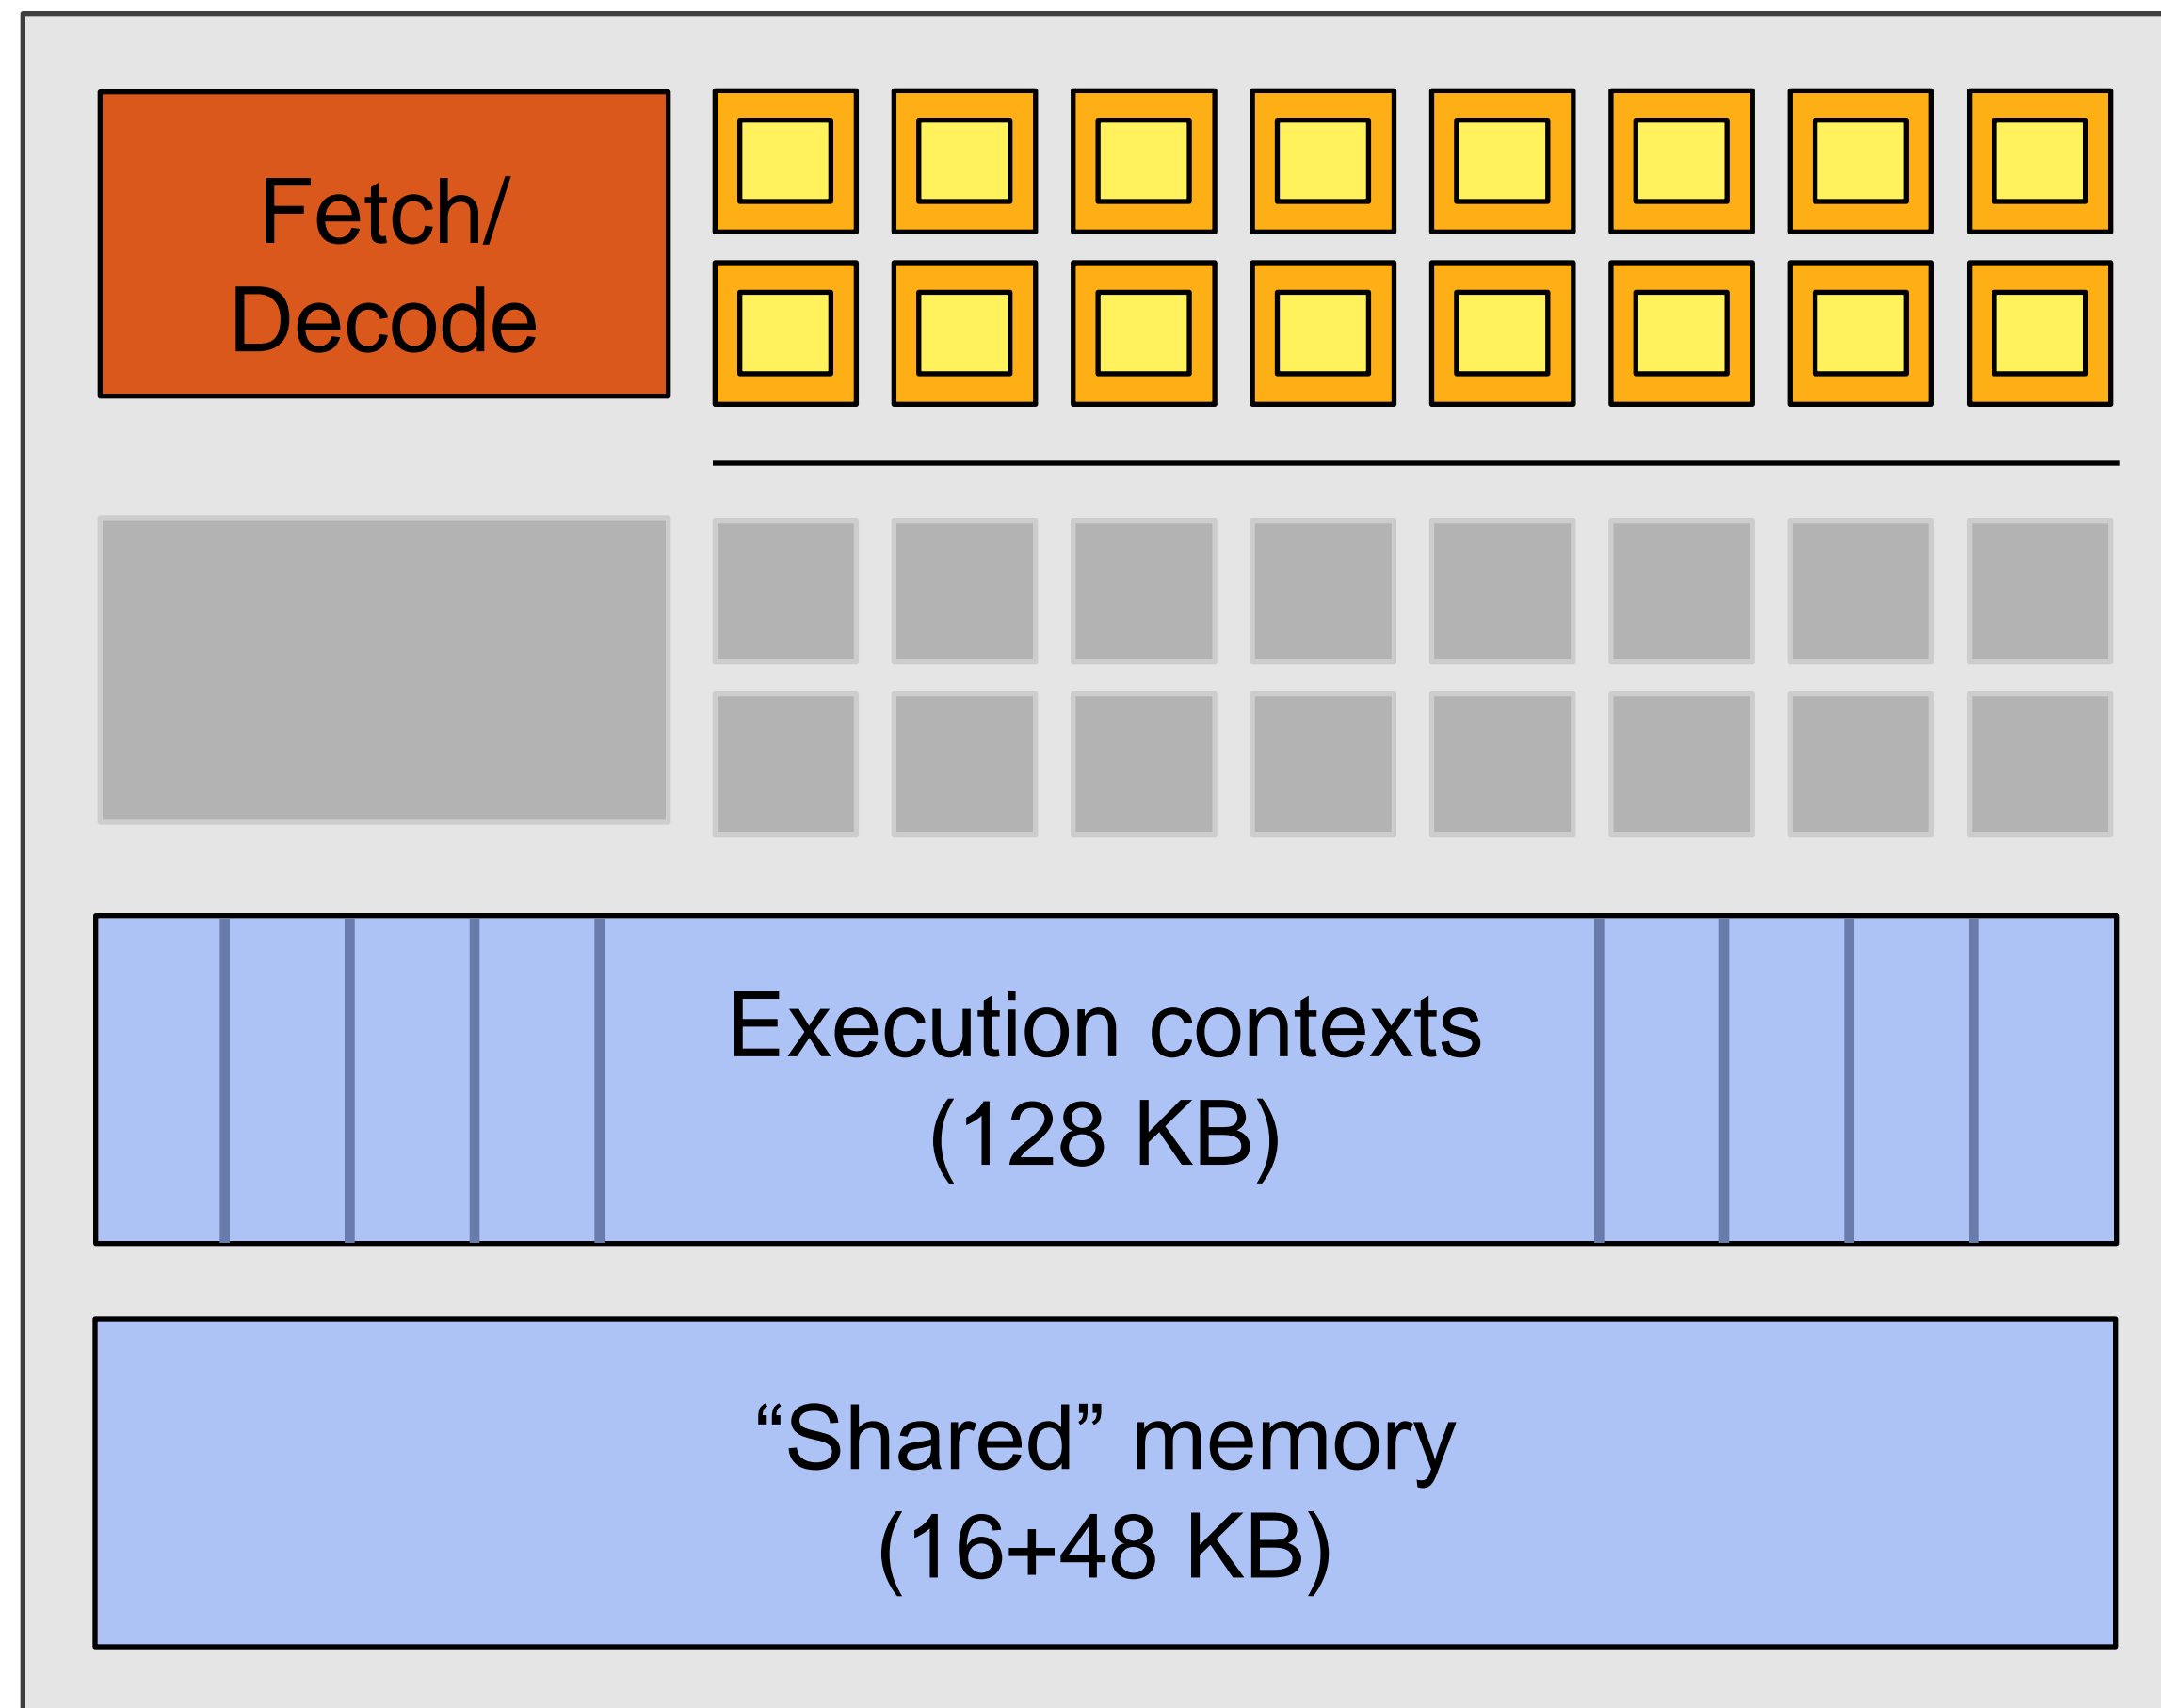
- The following slides describe “a reasonable way to think” about the architecture of commercial GPUs
- Many factors play a role in actual chip performance

# NVIDIA GeForce GTX 580 (Fermi)

- NVIDIA-speak:
  - 512 stream processors (“CUDA cores”)
  - “SIMT execution”
- Generic speak:
  - 16 cores
  - 2 groups of 16 SIMD functional units per core





# NVIDIA GeForce GTX 580 “core”



= SIMD function unit,  
control shared across 16 units  
(1 MUL-ADD per clock)

- Groups of 32 [fragments/vertices/CUDA threads] share an instruction stream
- Up to 48 groups are simultaneously interleaved
- Up to 1536 individual contexts can be stored


# NVIDIA GeForce GTX 580 “core”



= SIMD function unit,  
control shared across 16 units  
(1 MUL-ADD per clock)

- The core contains 32 functional units
- Two groups are selected each clock (decode, fetch, and execute two instruction streams in parallel)


# NVIDIA GeForce GTX 580 “SM”



 = **CUDA core**  
(1 MUL-ADD per clock)

- The **SM** contains 32 **CUDA cores**
- Two **warps** are selected each clock (decode, fetch, and execute two **warps** in parallel)
- Up to 48 warps are interleaved, totaling 1536 **CUDA threads**

# NVIDIA GeForce GTX 580



There are 16 of these things on the GTX 480:

That's 24,500 fragments!  
Or 24,000 CUDA threads!





# AMD Radeon HD 6970 (Cayman)

- AMD-speak:
  - 1536 stream processors
- Generic speak:
  - 24 cores
  - 16 “beefy” SIMD functional units per core
  - 4 multiply-adds per functional unit (VLIW processing)




# ATI Radeon HD 6970 “core”

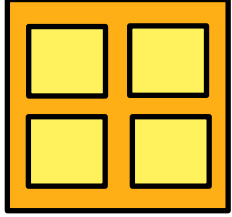


= SIMD function unit,  
control shared across 16 units  
(Up to 4 MUL-ADDs per clock)

- Groups of 64 [fragments/vertices/etc.] share instruction stream
- Four clocks to execute an instruction for all fragments in a group


# ATI Radeon HD 6970 “SIMD-Engine”



 = **Stream Processor**,  
control shared across 16 units  
(Up to 4 MUL-ADDs per clock)

- Groups of 64 [fragments/vertices/etc.] are in a “wavefront”
- Four clocks to execute an instruction for an entire “wavefront”

# ATI Radeon HD 6970

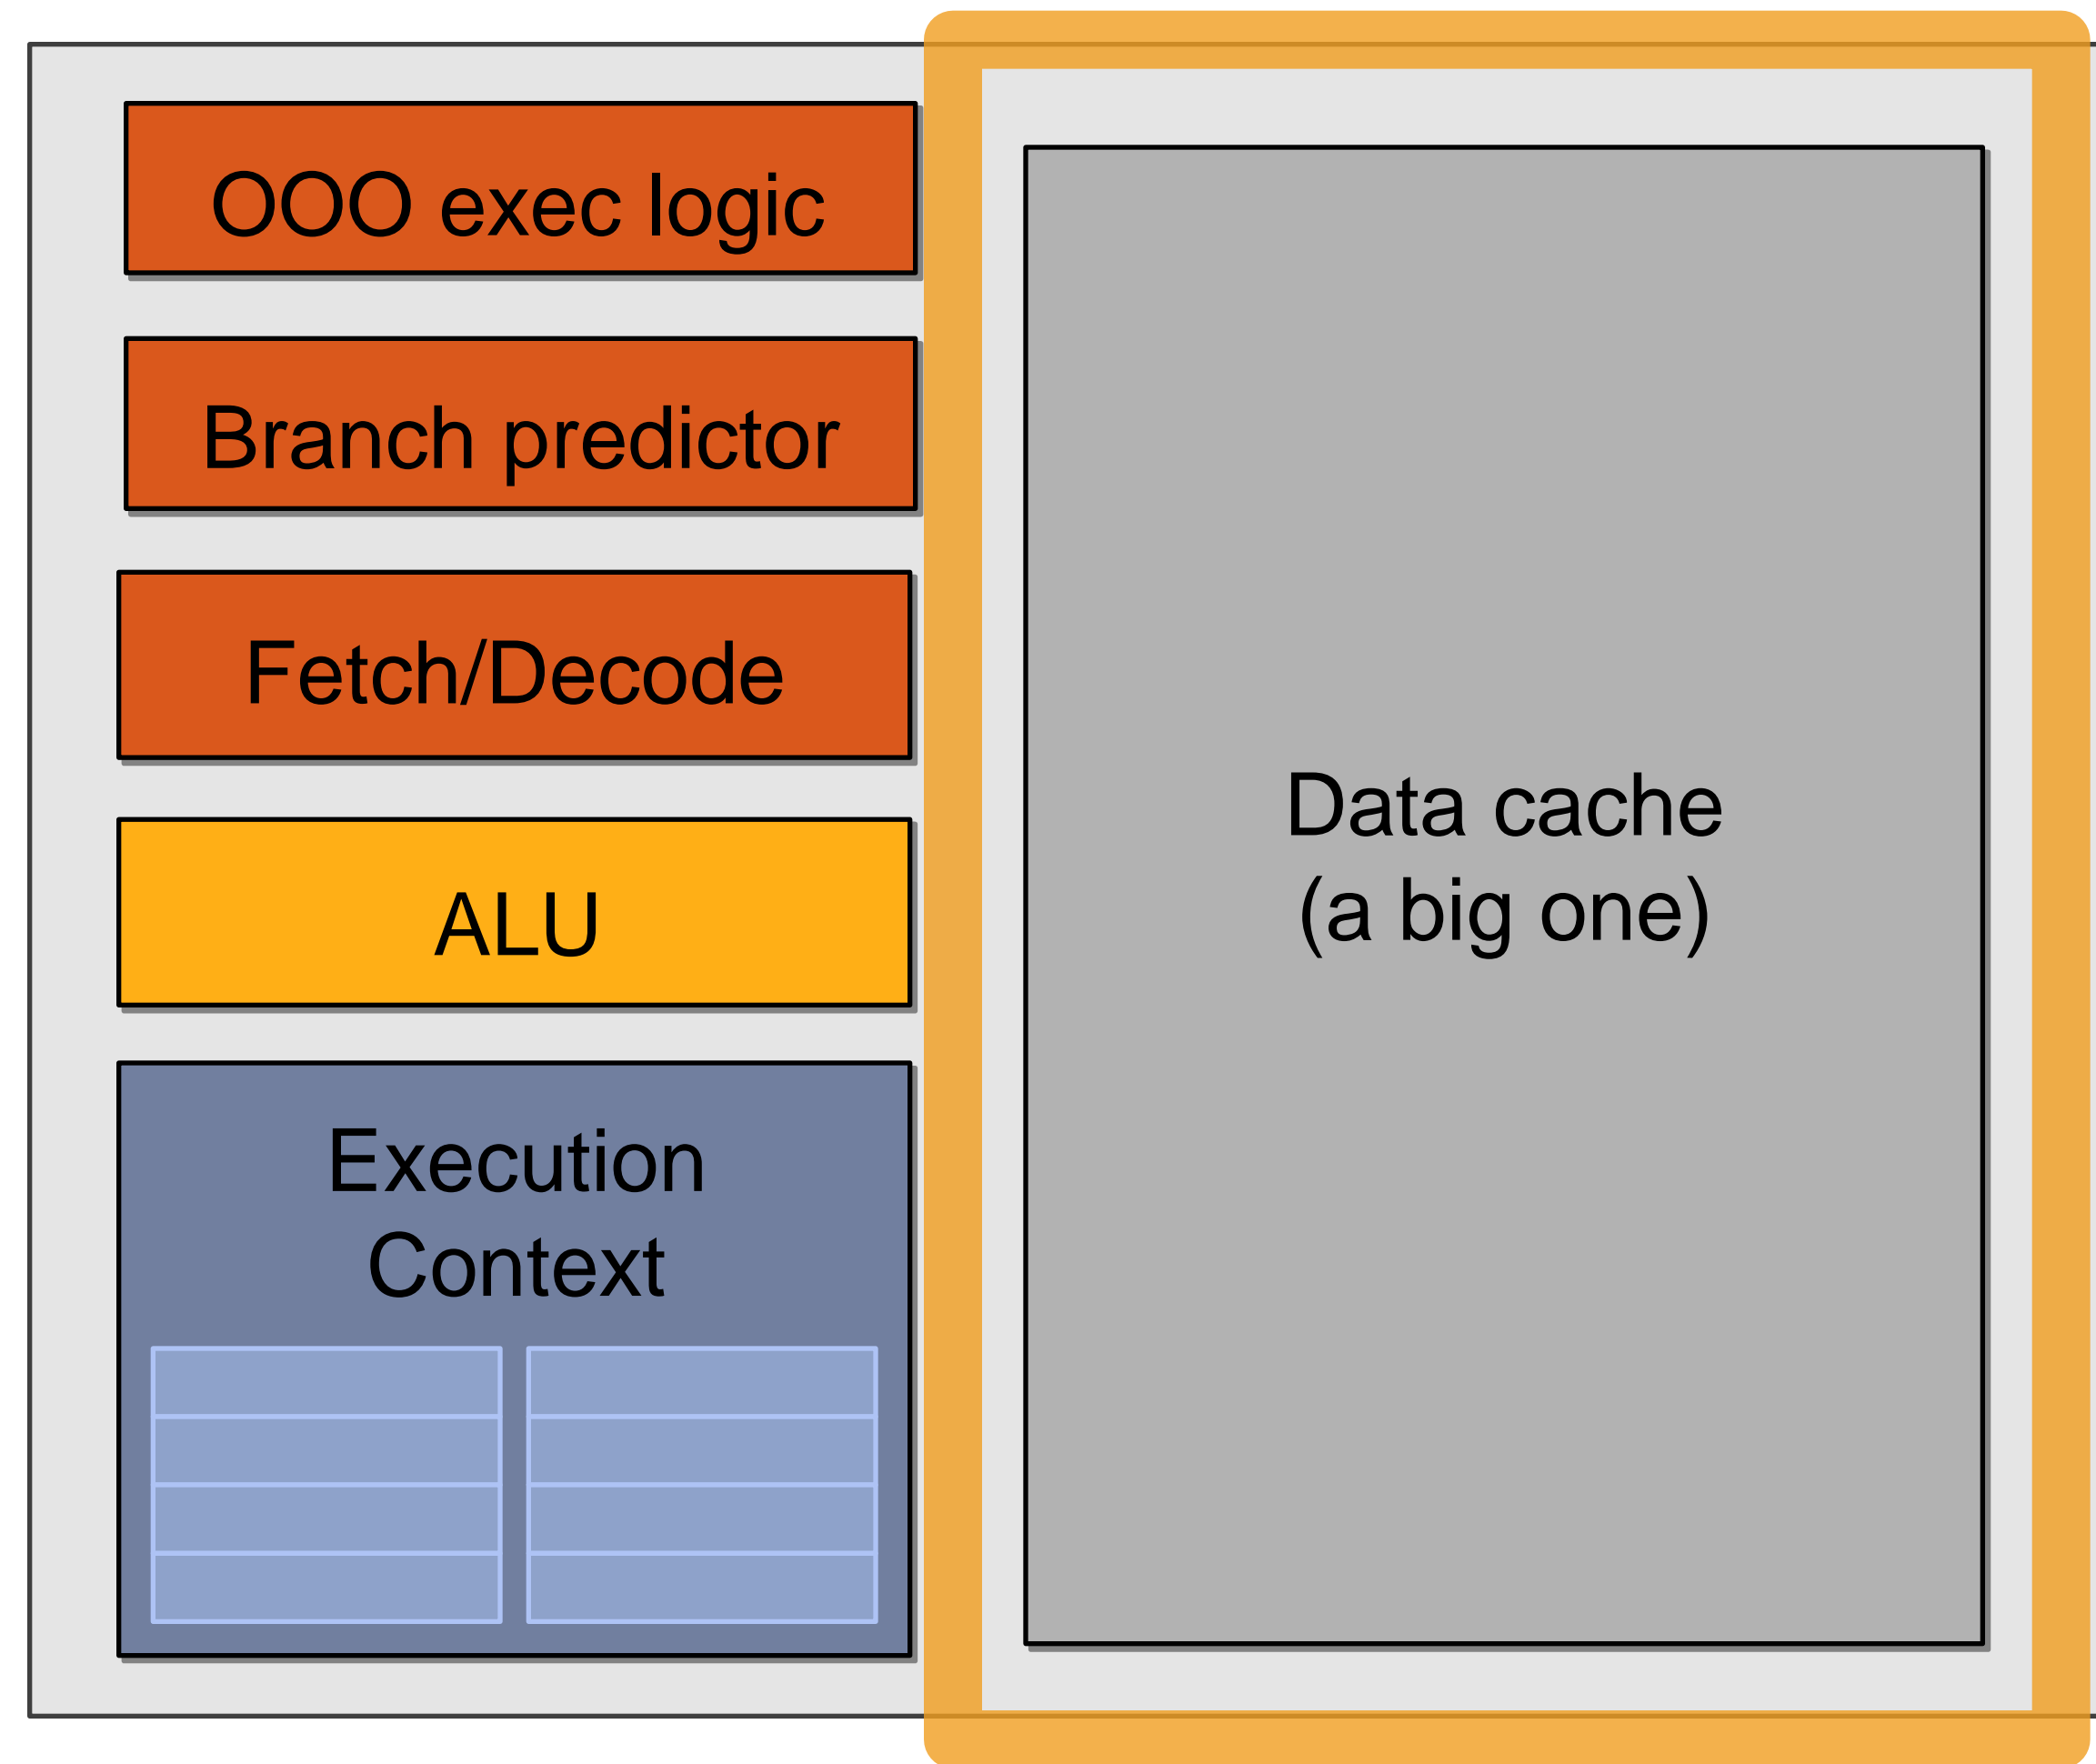


There are 24 of these “cores” on the 6970: that’s about 32,000 fragments!  
(there is a global limitation of 496 wavefronts)

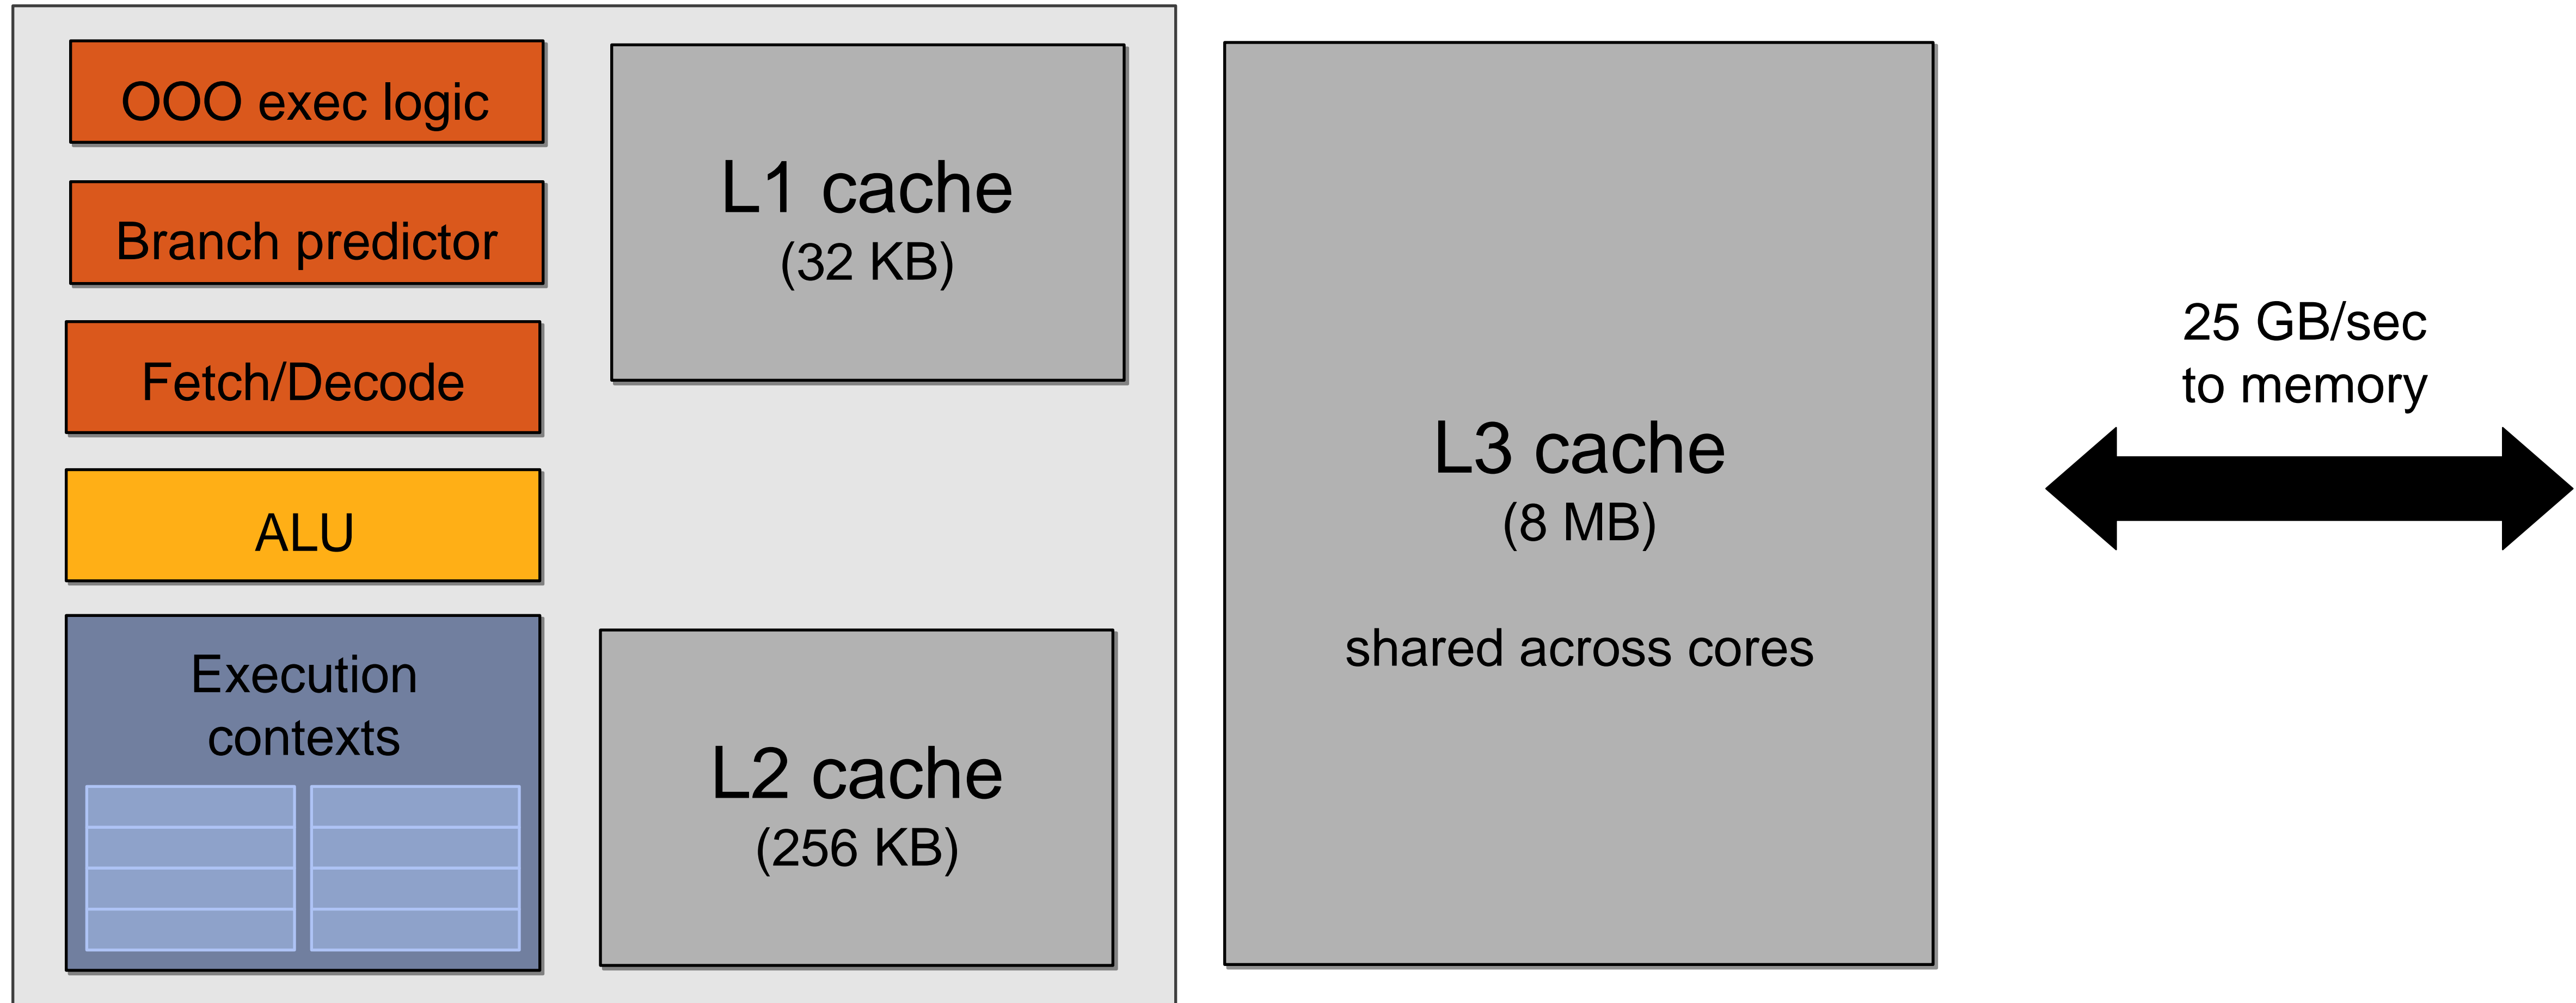
The talk thus far: processing data

## Part 3: moving data to processors

# Recall: “CPU-style” core




# “CPU-style” memory hierarchy



CPU cores run efficiently when data is resident in cache  
(caches reduce latency, provide high bandwidth)

# Throughput core (GPU-style)



More ALUs, no large traditional cache hierarchy:  
Need high-bandwidth connection to memory



# Bandwidth is a critical resource

- A high-end GPU (e.g. Radeon HD 6970) has...
  - Over **twenty times** (2.7 TFLOPS) the compute performance of quad-core CPU
  - No large cache hierarchy to absorb memory requests
- GPU memory system is designed for throughput
  - Wide bus (150 GB/sec)
  - Repack/reorder/interleave memory requests to maximize use of memory bus
  - Still, this is only **six times** the bandwidth available to CPU

# Bandwidth thought experiment

Task: element-wise multiply two long vectors A and B


1. Load input A[i]

2. Load input B[i]

3. Load input C[i]

4. Compute  $A[i] \times B[i] + C[i]$

5. Store result into D[i]



Four memory operations (16 bytes) for every MUL-ADD

Radeon HD 6970 can do 1536 MUL-ADDS per clock

Need ~20 TB/sec of bandwidth to keep functional units busy

Less than 1% efficiency... but 6x faster than CPU!

# Bandwidth limited!

If processors request data at too high a rate,  
the memory system cannot keep up.

No amount of latency hiding helps this.


Overcoming bandwidth limits are a common challenge  
for GPU-compute application developers.

# Reducing bandwidth requirements

- Request data less often (instead, do more math)
  - “arithmetic intensity”
- Fetch data from memory less often (share/reuse data across fragments)
  - on-chip communication or storage

# Reducing bandwidth requirements

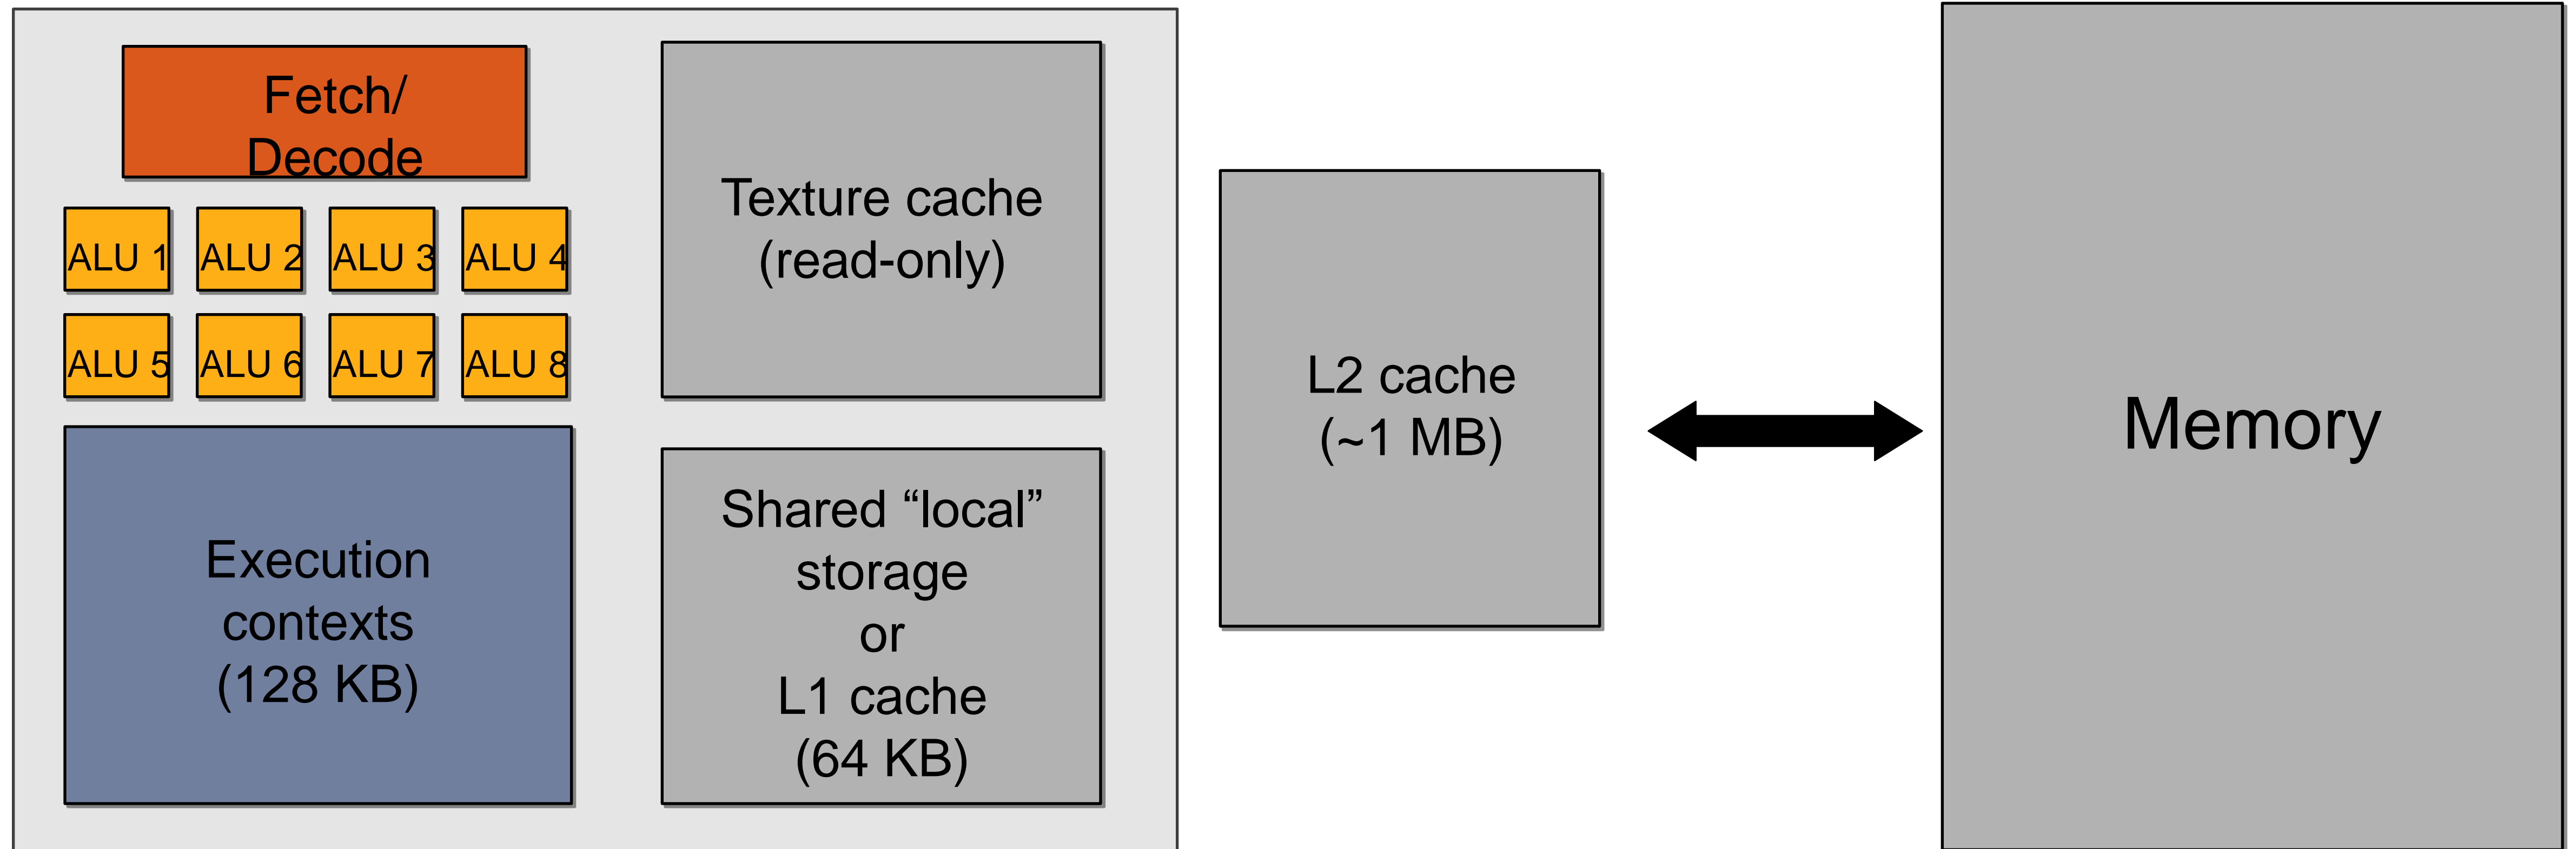
- Two examples of on-chip storage
  - Texture caches
  - OpenCL “local memory” (CUDA shared memory)



Texture caches:

Capture reuse across fragments, not temporal reuse within a single shader program

# Modern GPU memory hierarchy



On-chip storage takes load off memory system.  
Many developers calling for more cache-like storage  
(particularly GPU-compute applications)

# Don't forget about offload cost...


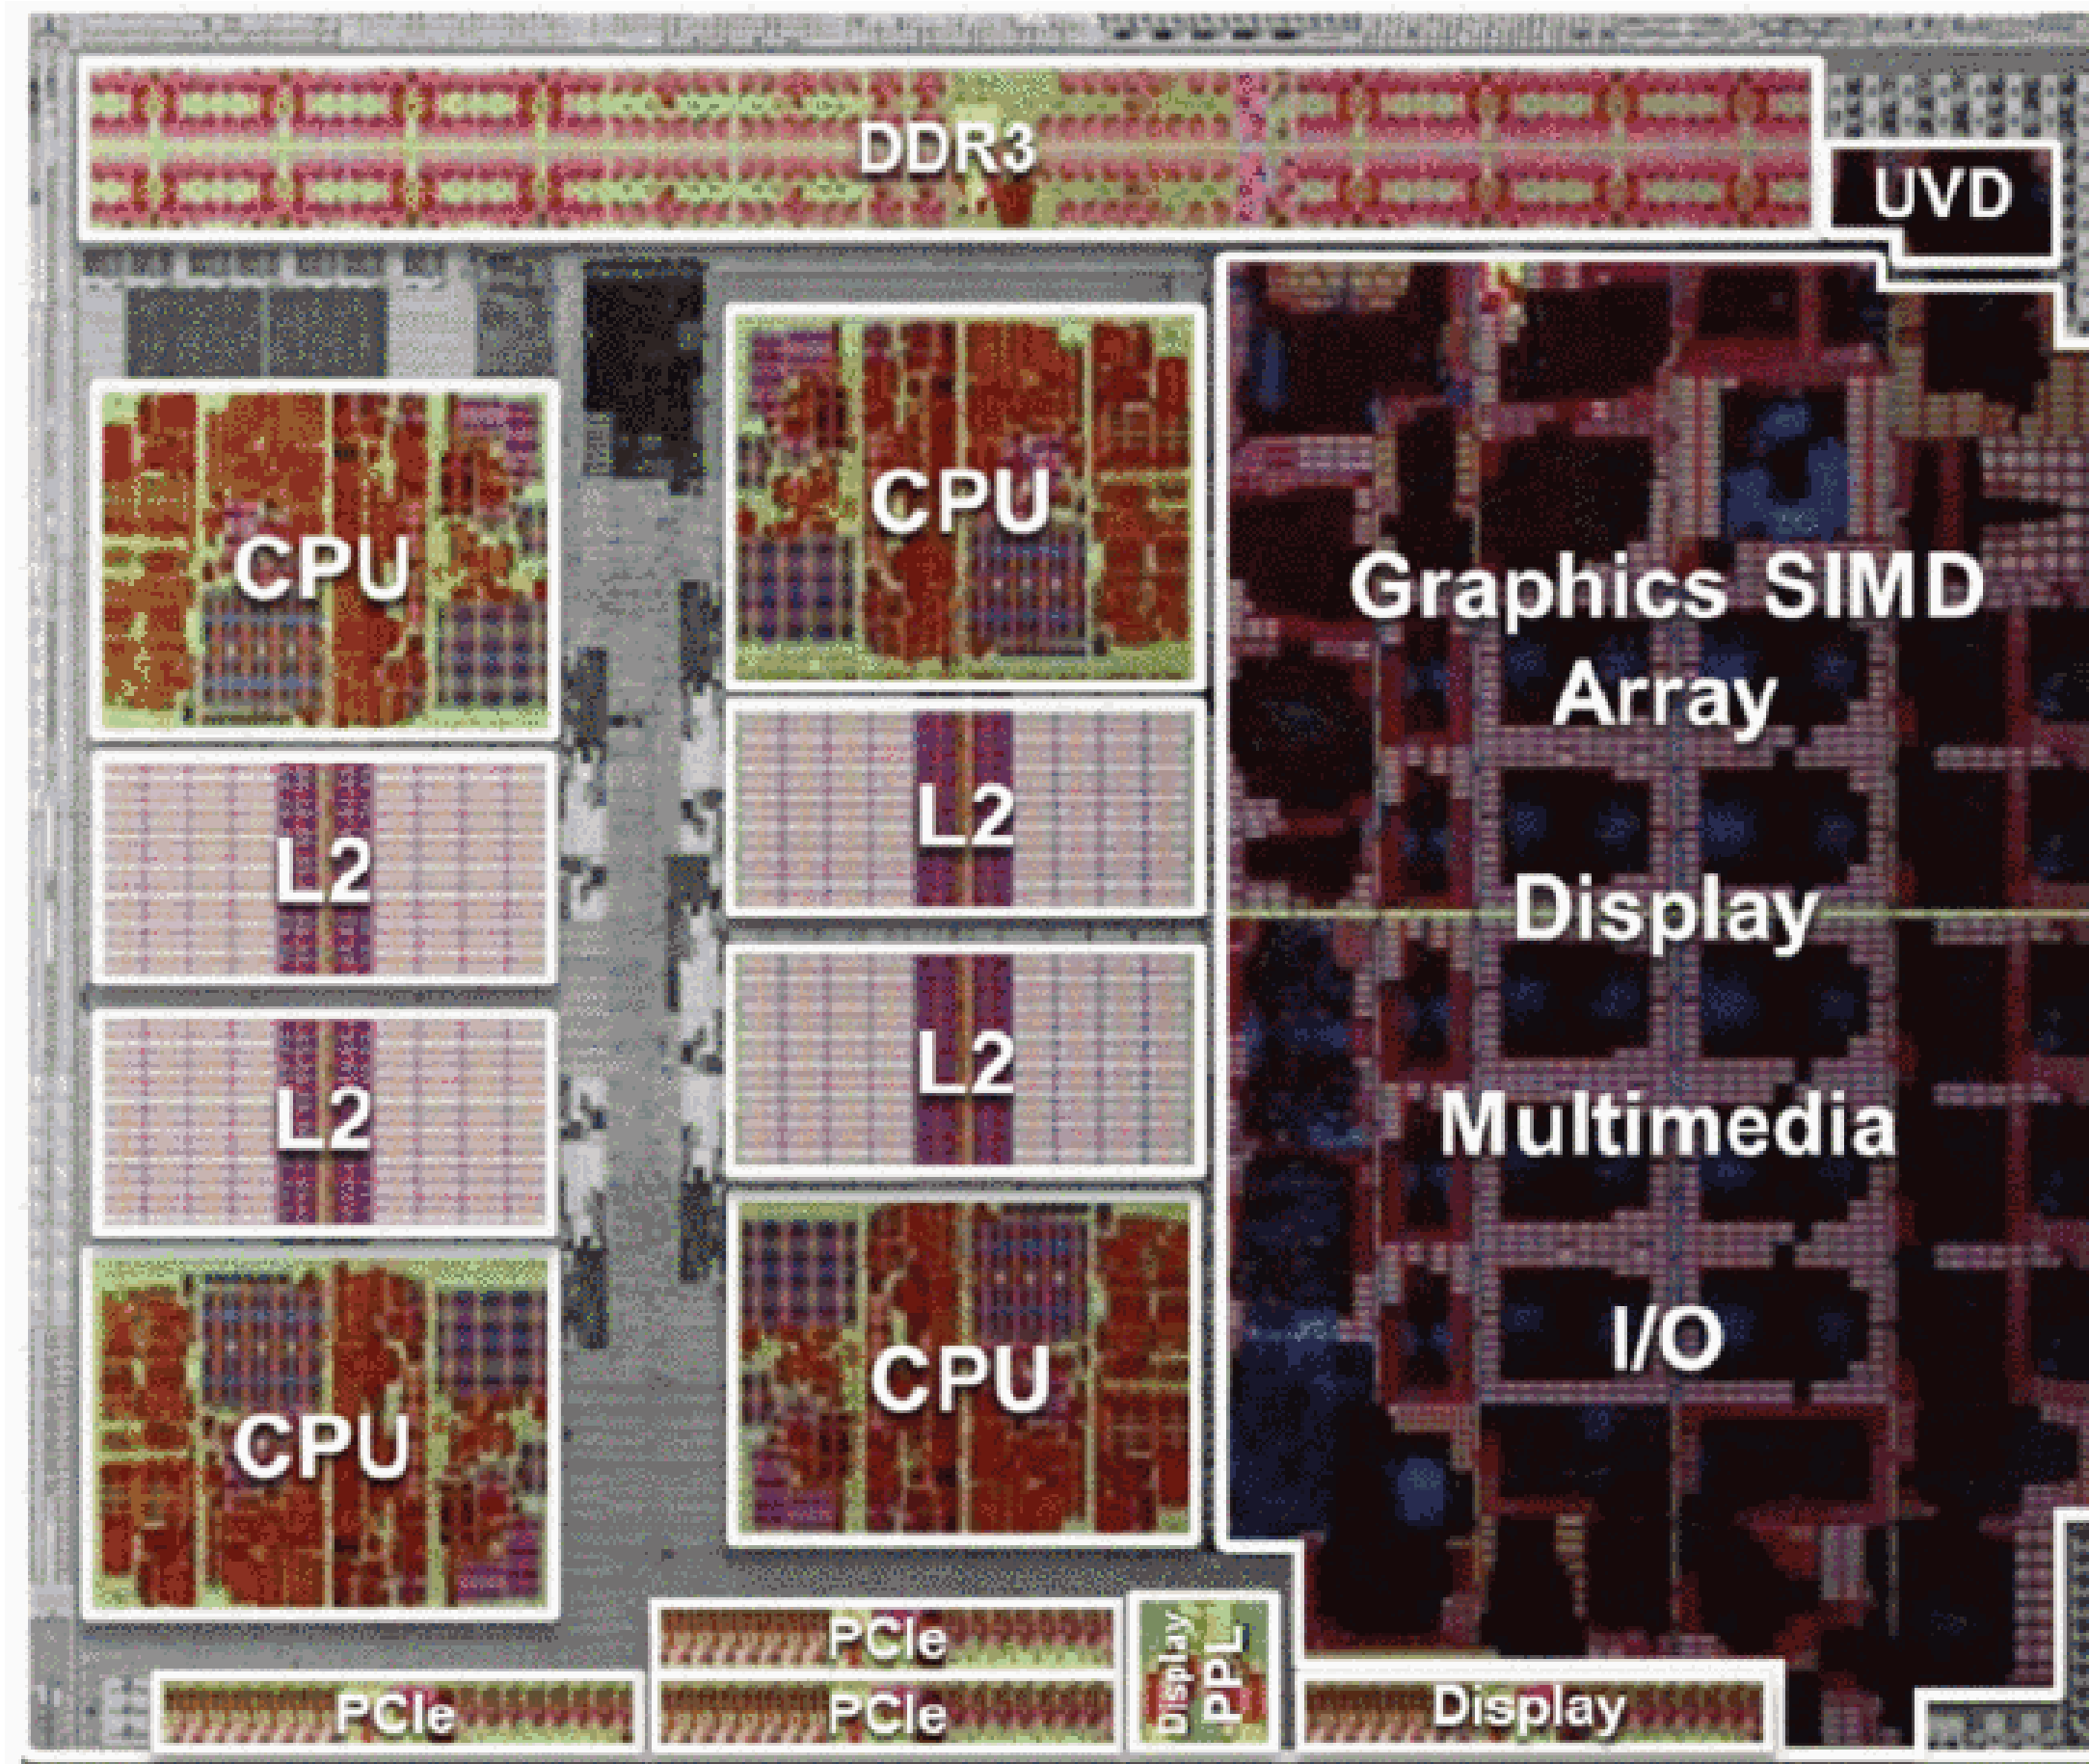
- PCIe bandwidth/latency
  - 8GB/s each direction in practice
  - Attempt to pipeline/multi-buffer uploads and downloads
- Dispatch latency
  - $O(10)$  usec to dispatch from CPU to GPU
  - This means offload cost if  $O(10M)$  instructions

# Heterogeneous cores to the rescue ?

- Tighter integration of CPU and GPU style cores
  - Reduce offload cost
  - Reduce memory copies/transfers
  - Power management
- Industry shifting rapidly in this direction
  - AMD Fusion APUs
  - Intel SandyBridge
  - ...
  - NVIDIA Tegra 2
  - Apple A4 and A5
  - QUALCOMM Snapdragon
  - TI OMAP
  - ...

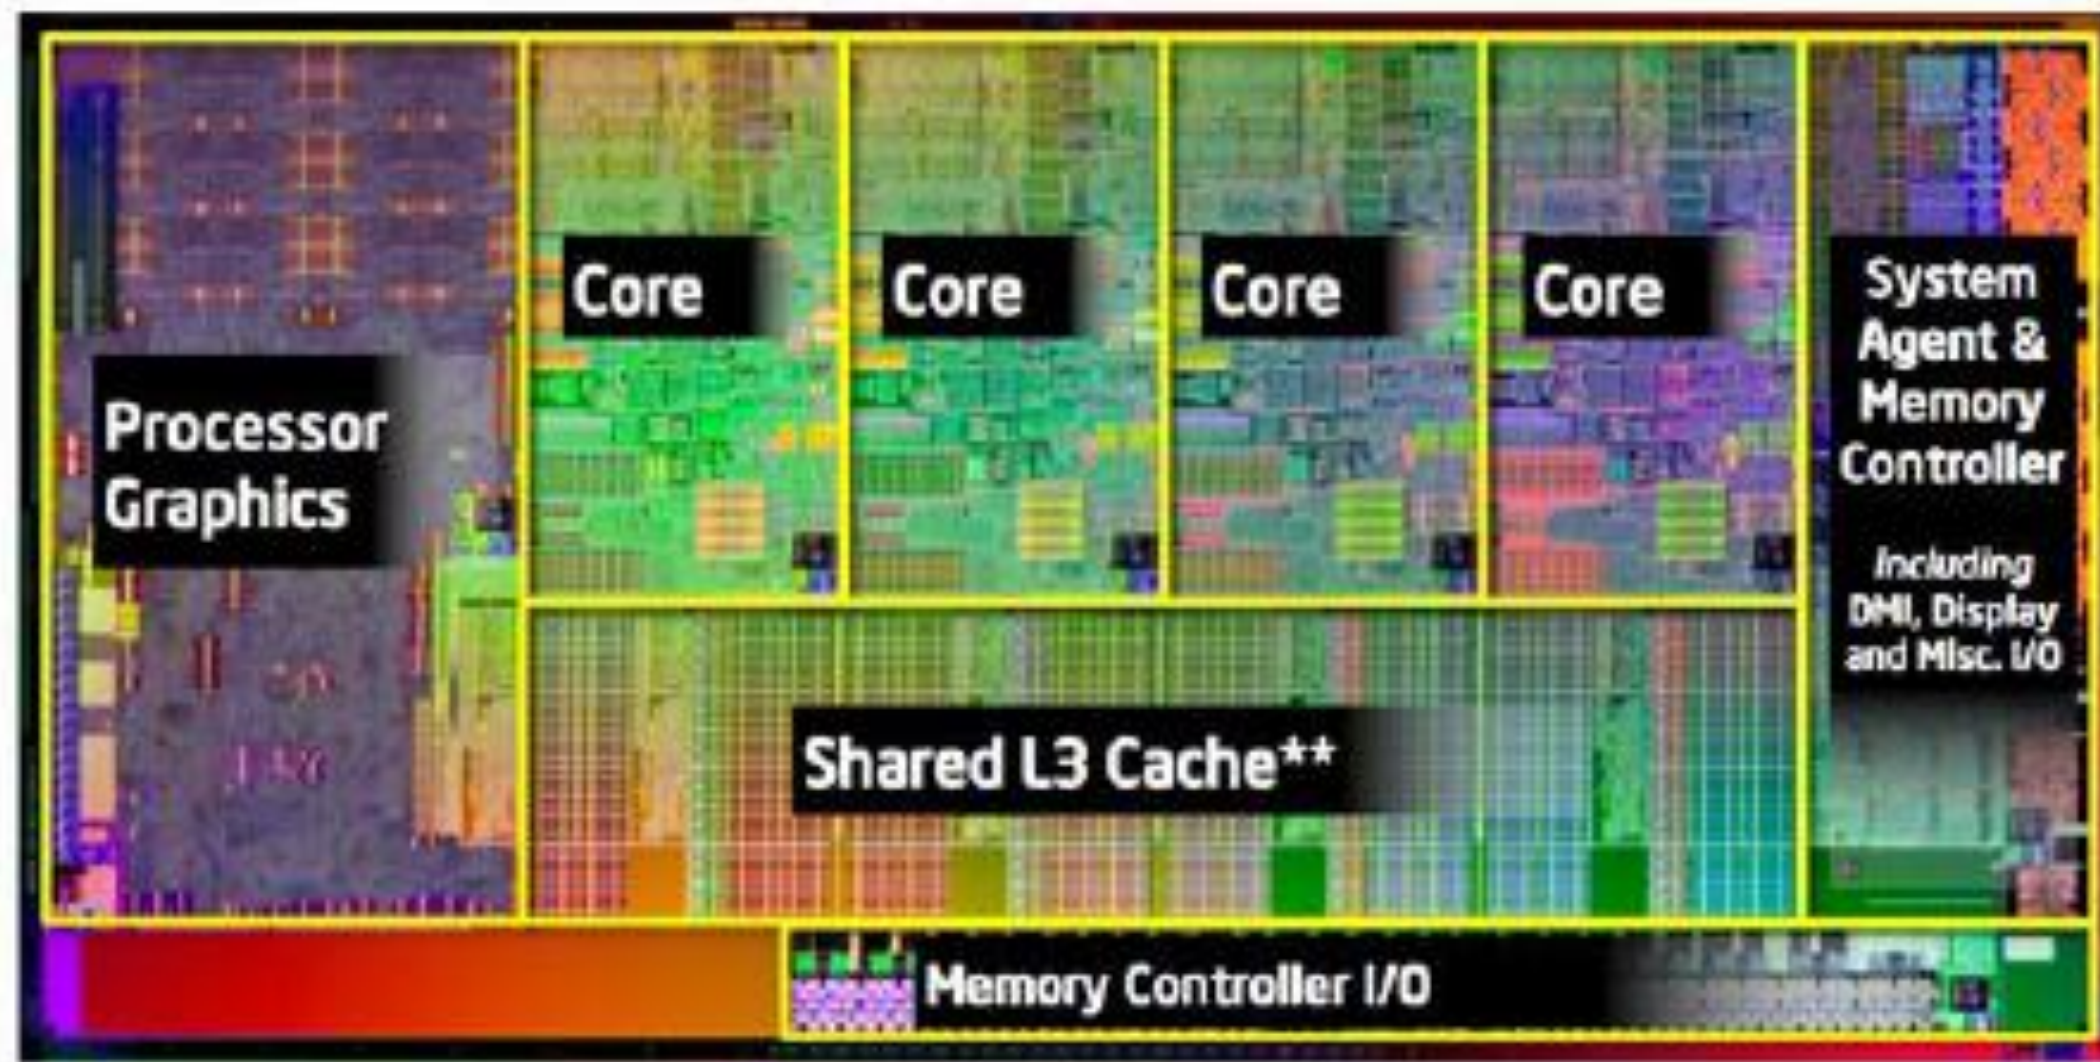


# AMD A-Series APU ("Llano")

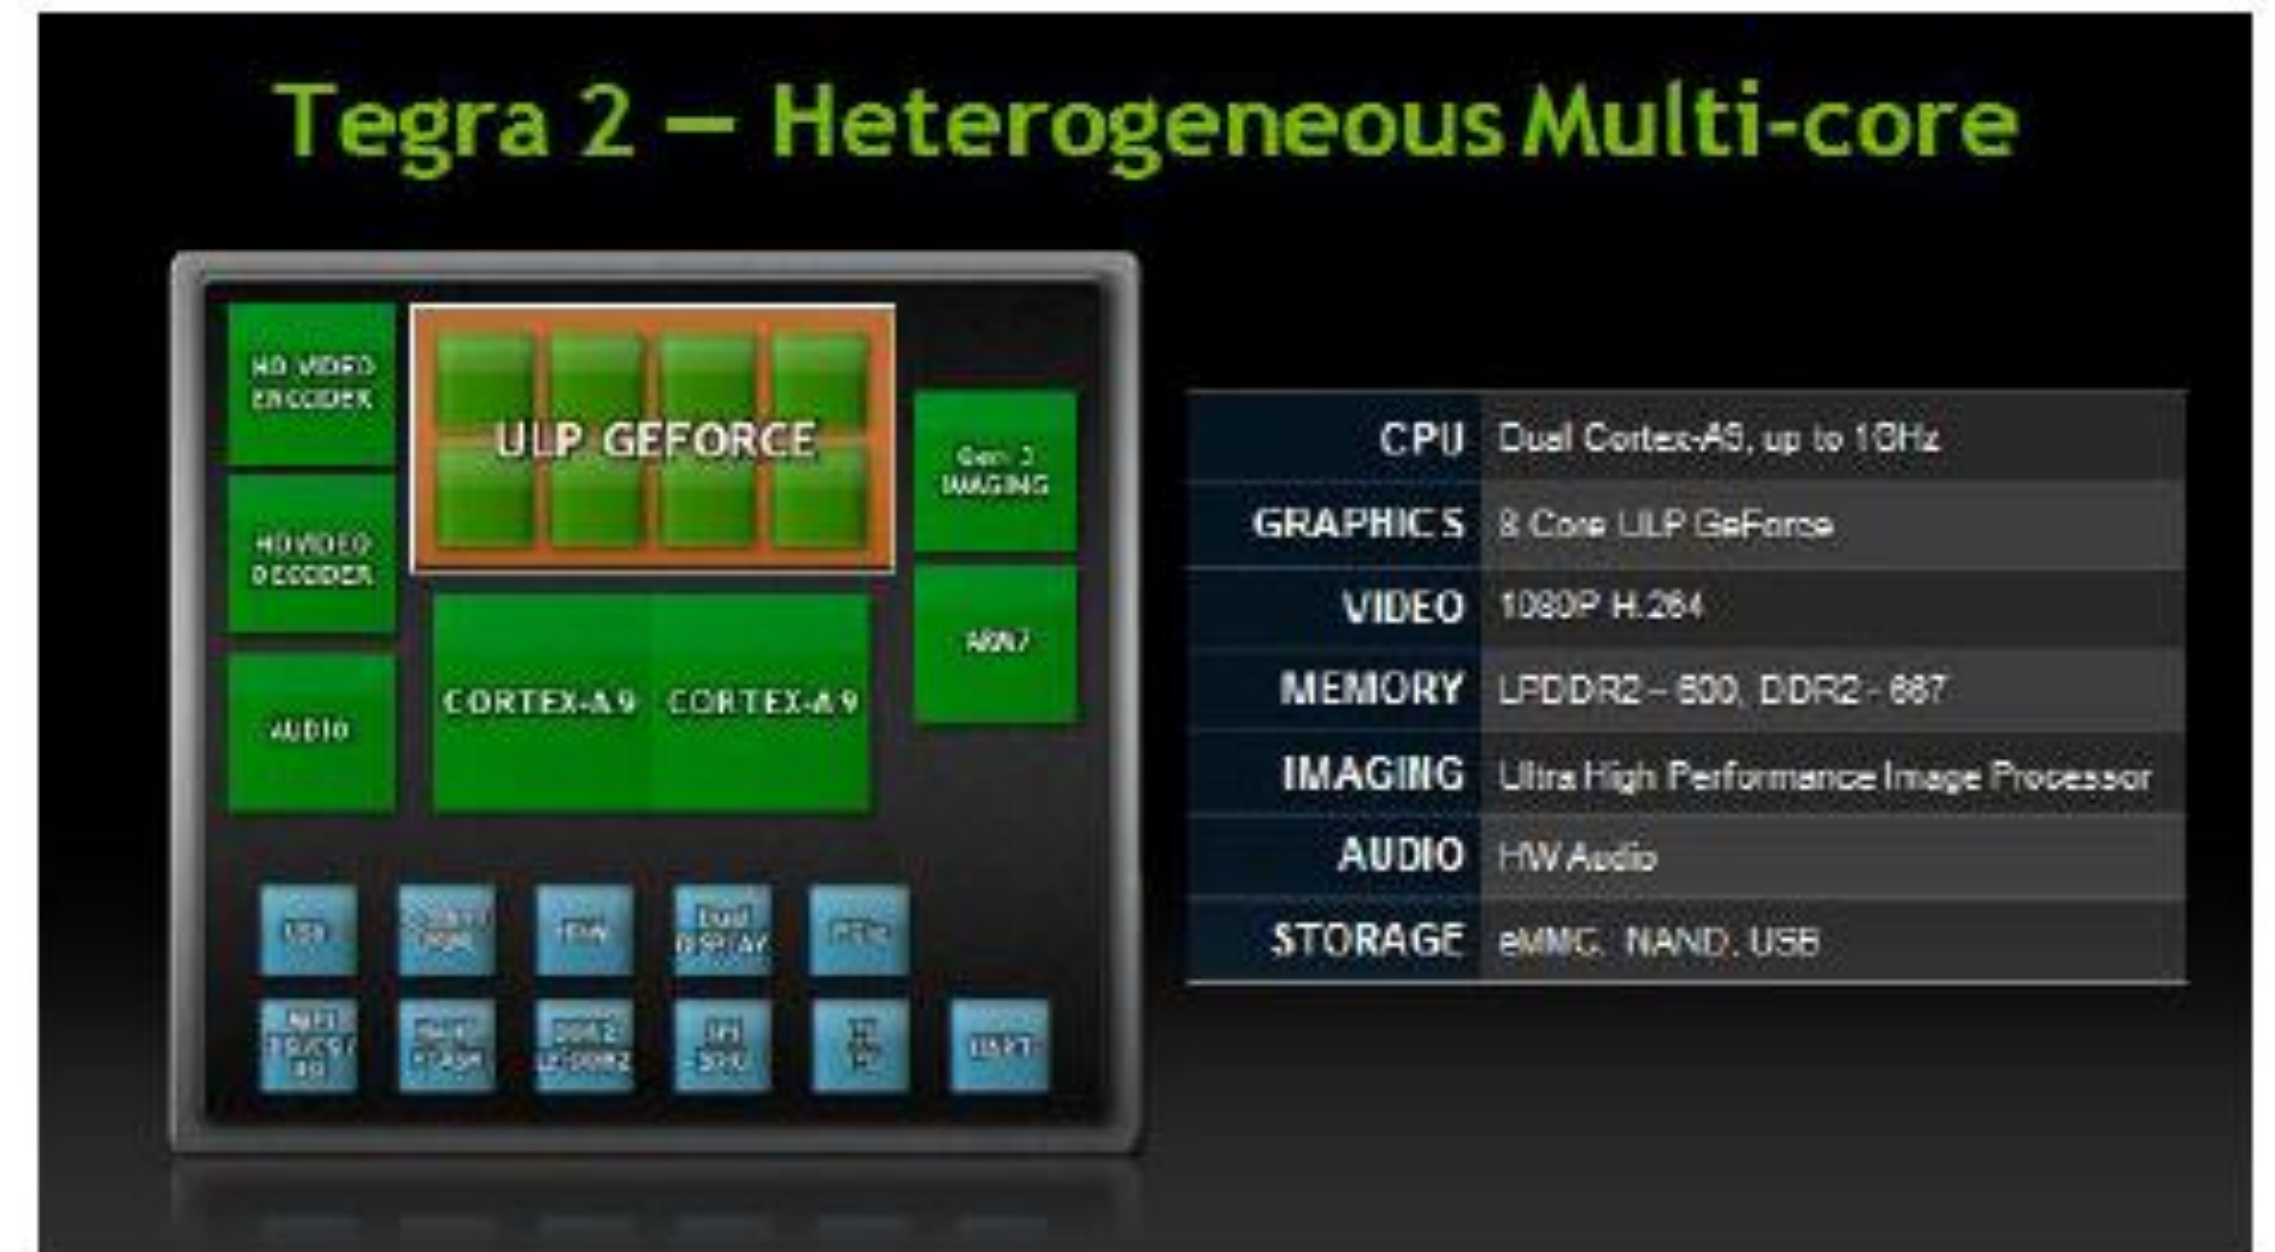




Others – GPU is not compute ready, yet



# Intel SandyBridge



# NVIDIA Tegra 2